

CS:APP Chapter 4

Computer Architecture

Pipelined Implementation

Randal E. Bryant
adapted by Jason Fritts

<http://csapp.cs.cmu.edu>

Overview

General Principles of Pipelining

- Goal
- Difficulties

Creating a Pipelined Y86 Processor

- Rearranging SEQ to create pipelined datapath, PIPE
- Inserting pipeline registers
- Problems with data and control hazards

Fundamentals of Pipelining

Real-World Pipelines: Car Washes

Sequential



Parallel



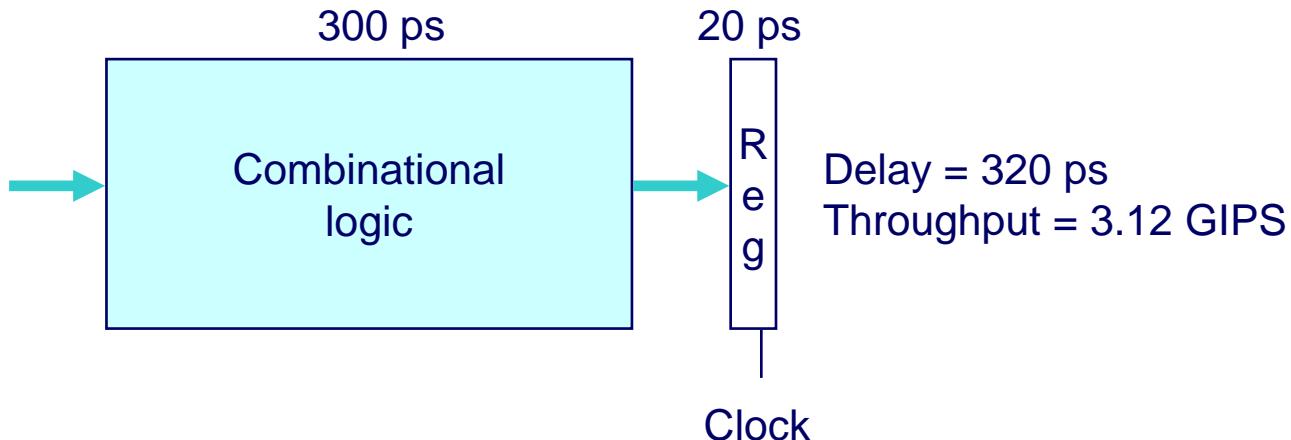
Pipelined



Idea

- Divide process into independent stages
- Move objects through stages in sequence
- At any given times, multiple objects being processed

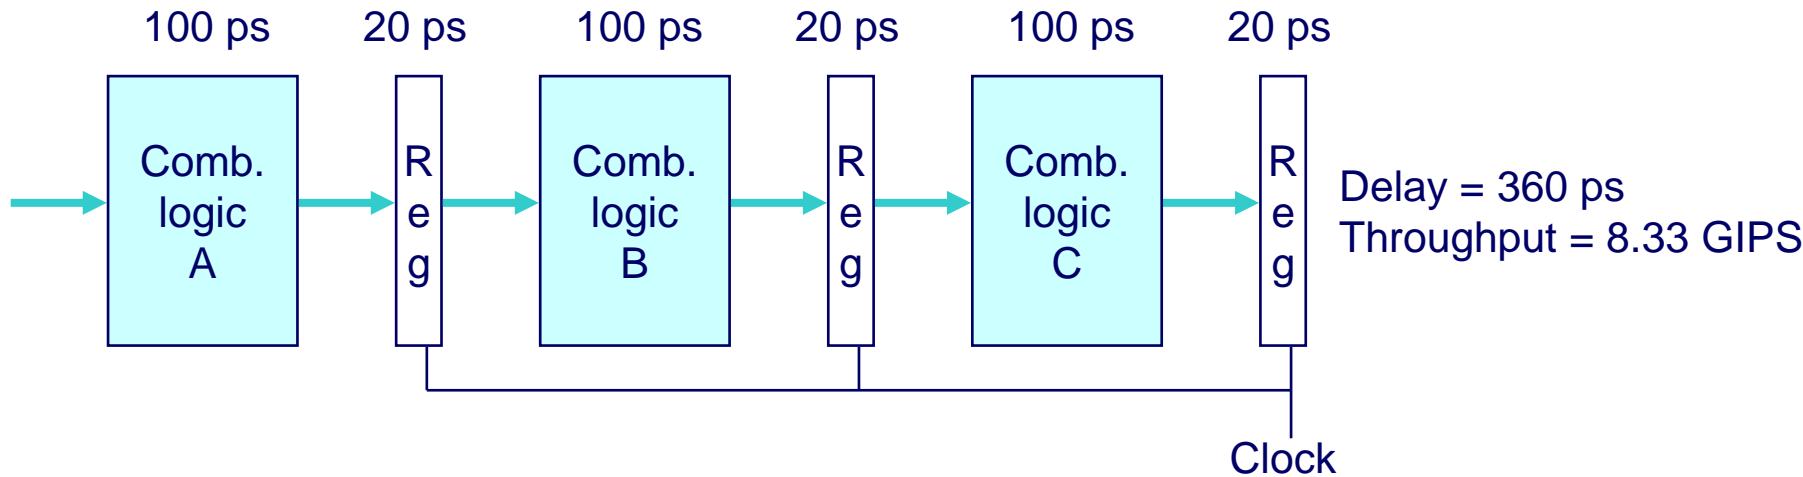
Computational Example



System

- Computation requires total of 300 picoseconds
- Additional 20 picoseconds to save result in register
- Must have clock cycle of at least 320 ps

3-Way Pipelined Version



System

- Divide combinational logic into 3 blocks of 100 ps each
- Can begin new operation as soon as previous one passes through stage A.
 - Begin new operation every 120 ps
- Overall latency increases
 - 360 ps from start to finish

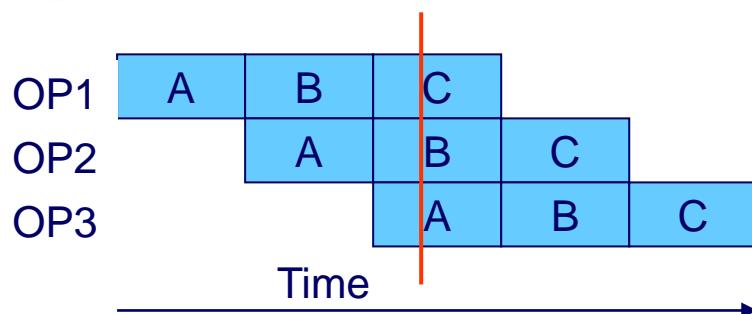
Pipeline Diagrams

Unpipelined



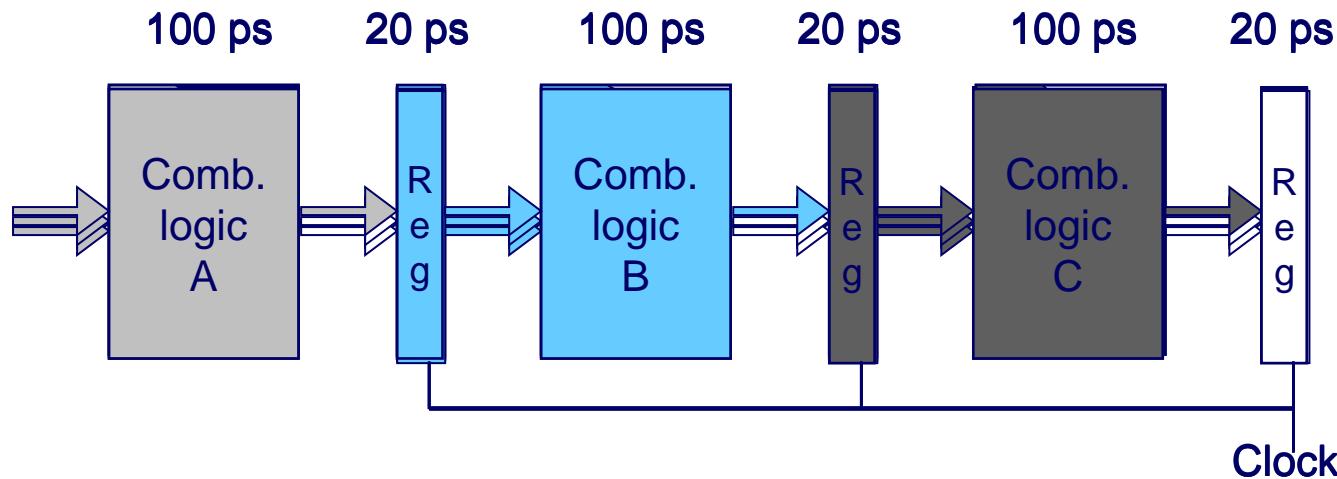
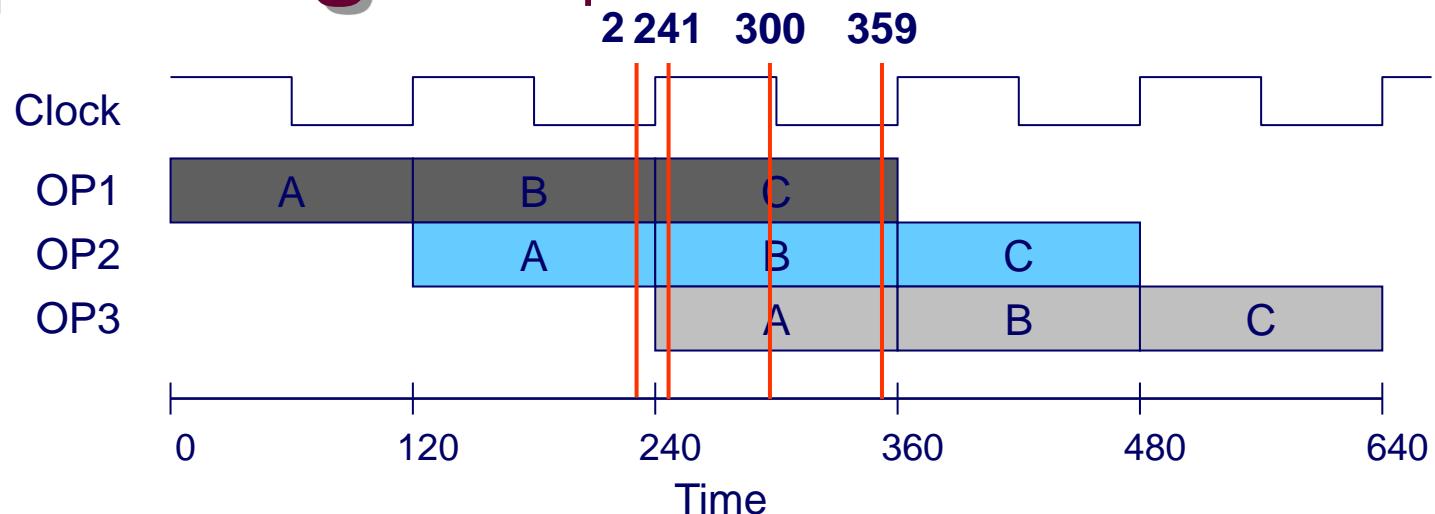
- Cannot start new operation until previous one completes

3-Way Pipelined

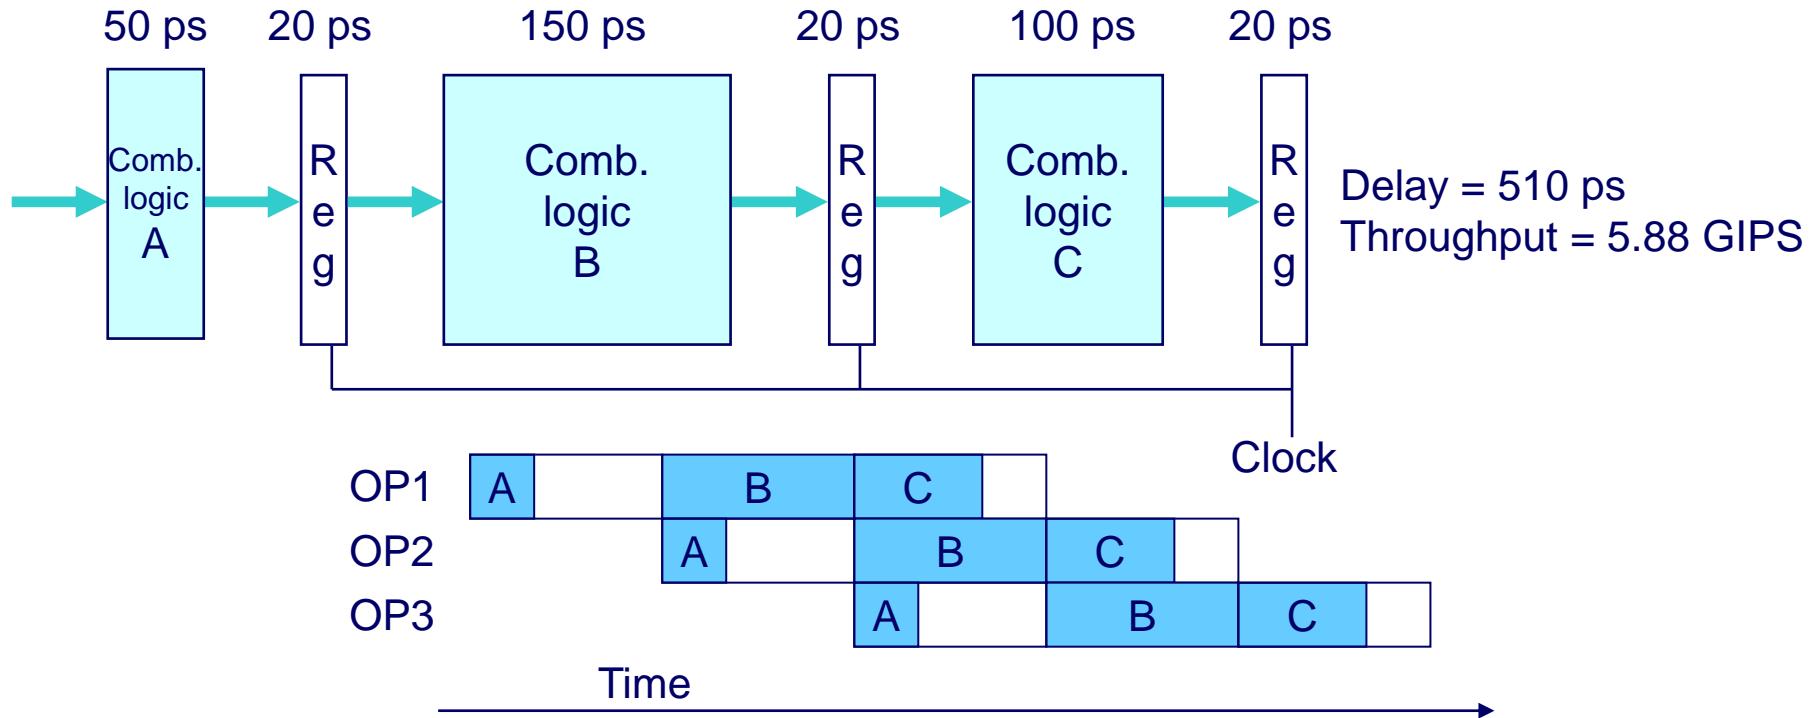


- Up to 3 operations in process simultaneously

Operating a Pipeline



Limitations: Nonuniform Delays



- Throughput limited by slowest stage
- Other stages sit idle for much of the time
- Challenging to partition system into balanced stages

Sample Circuit Delays & Pipelining

<i>instruction memory</i>	220ps
<i>decode</i>	70ps
<i>register fetch</i>	120ps
<i>ALU</i>	180ps
<i>data memory</i>	260ps
<i>register writeback</i>	120ps

20ps delay for
hardware register at
end of cycle

Single-cycle processor:

- Clock cycle = $220 + 70 + 120 + 180 + 260 + 120 + 20 = 990\text{ps}$
- Clock freq = $1 / 990\text{ps} = 1 / 990 * 10^{-12} = 1.01 \text{ GHz}$

Combine and/or split stages for pipelining

- Need to balance time per stage since clock freq determined by slowest time
- Must maintain original order of stages, so can't combine non-neighboring stages (e.g. can't combine *decode* & *data mem*)

Sample Circuit Delays & Pipelining

<i>instruction memory</i>	220ps	
<i>decode</i>	70ps	
<i>register fetch</i>	120ps	
<i>ALU</i>	180ps	
<i>data memory</i>	260ps	
<i>register writeback</i>	120ps	
		20ps delay added for hardware register at end of each cycle

3-stage pipeline:

- Best combination for minimizing clock cycle time:

- 1st stage – *instr mem & decode*: 220 + 70 + 20 = 310ps
- 2nd stage – *reg fetch & ALU*: 120 + 180 + 20 = 320ps
- 3rd stage – *data mem & reg WB*: 260 + 120 + 20 = 400ps

- Slowest stage is 400ps, so clock cycle time is 400ps

- Clock freq = 1 / 400ps = 1 / 400*10⁻¹² = 2.5 GHz

Sample Circuit Delays & Pipelining

<i>instruction memory</i>	220ps	
<i>decode</i>	70ps	
<i>register fetch</i>	120ps	
<i>ALU</i>	180ps	
<i>data memory</i>	260ps	
<i>register writeback</i>	120ps	
		20ps delay added for hardware register at end of each cycle

5-stage pipeline:

- Best combination for minimizing clock cycle time:

● 1 st stage – <i>instr mem</i> :	220 + 20ps	= 240ps
● 2 nd stage – <i>decode & reg fetch</i> :	70 + 120 + 20ps	= 210ps
● 3 rd stage – <i>ALU</i> :	180 + 20ps	= 200ps
● 4 th stage – <i>data mem</i> :	260 + 20ps	= 280ps
● 5 th stage – <i>reg WB</i> :	120 + 20ps	= 140ps

- Slowest stage is 280ps, so clock cycle time is 280ps
- Clock freq = 1 / 280ps = 1 / 280*10⁻¹² = 3.57 GHz

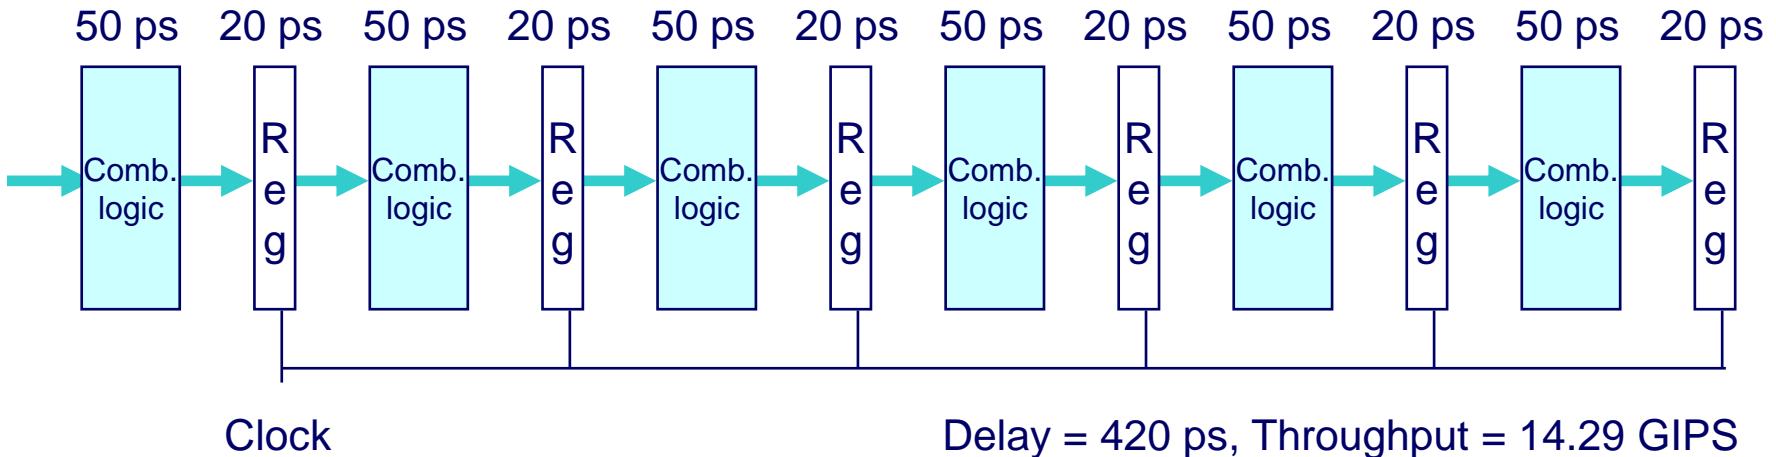
Sample Circuit Delays & Pipelining

<i>instruction memory</i>	220ps
<i>decode</i>	70ps
<i>register fetch</i>	120ps
<i>ALU</i>	180ps
<i>data memory</i>	260ps
<i>register writeback</i>	120ps
	20ps delay added for hardware register at end of each cycle

9-stage pipeline:

- Assuming can split stages evenly into halves, thirds, or quarters
 - not a valid assumption, but useful for simplifying problem
- Best combination for minimizing clock cycle time:
 - Each circuit is its own stage, with 20ps added delay for reg
 - Split *instr mem* circuit into two stages, each 110+20ps
 - Split *data mem* circuit into two stages, each 130+20ps
 - Split ALU circuit into two stages, each 90+20ps
- Slowest stage is 150ps, so clock cycle time is 150ps
- Clock freq = $1 / 150\text{ps} = 1 / 150 \cdot 10^{-12} = 6.67 \text{ GHz}$

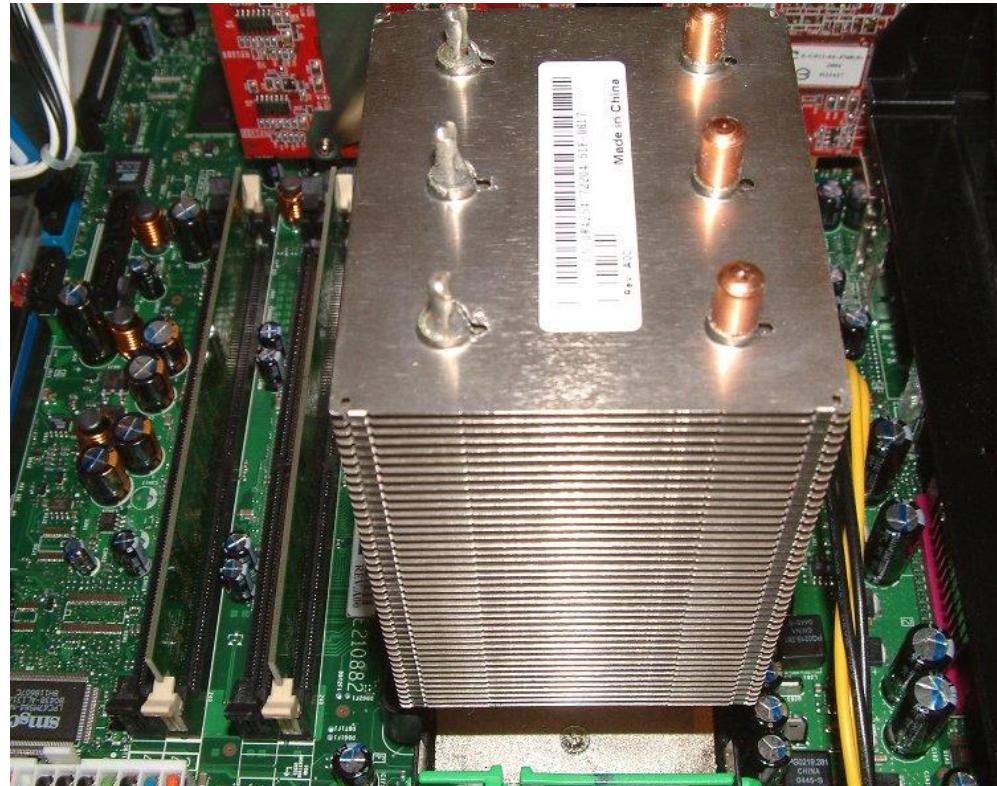
Limitations: Register Overhead



- As try to deepen pipeline, overhead of loading registers becomes more significant
- Percentage of clock cycle spent loading register:
 - 1-stage pipeline: 6.25%
 - 3-stage pipeline: 16.67%
 - 6-stage pipeline: 28.57%
- High speeds of modern processor designs obtained through very deep pipelining

In Practice

- i386 – 3 stage pipeline
- i486 – 5 stages
- Pentium 3 – 11 stages
- Pentium 4 (willamette)
 - 20 stages
- Pentium 4 (prescott)
 - 31 stages!
 - Up to 3.8GHz
 - Severe heat problems
 - Long pipeline actually hurt some application's performance
 - 115 Watts dissipation



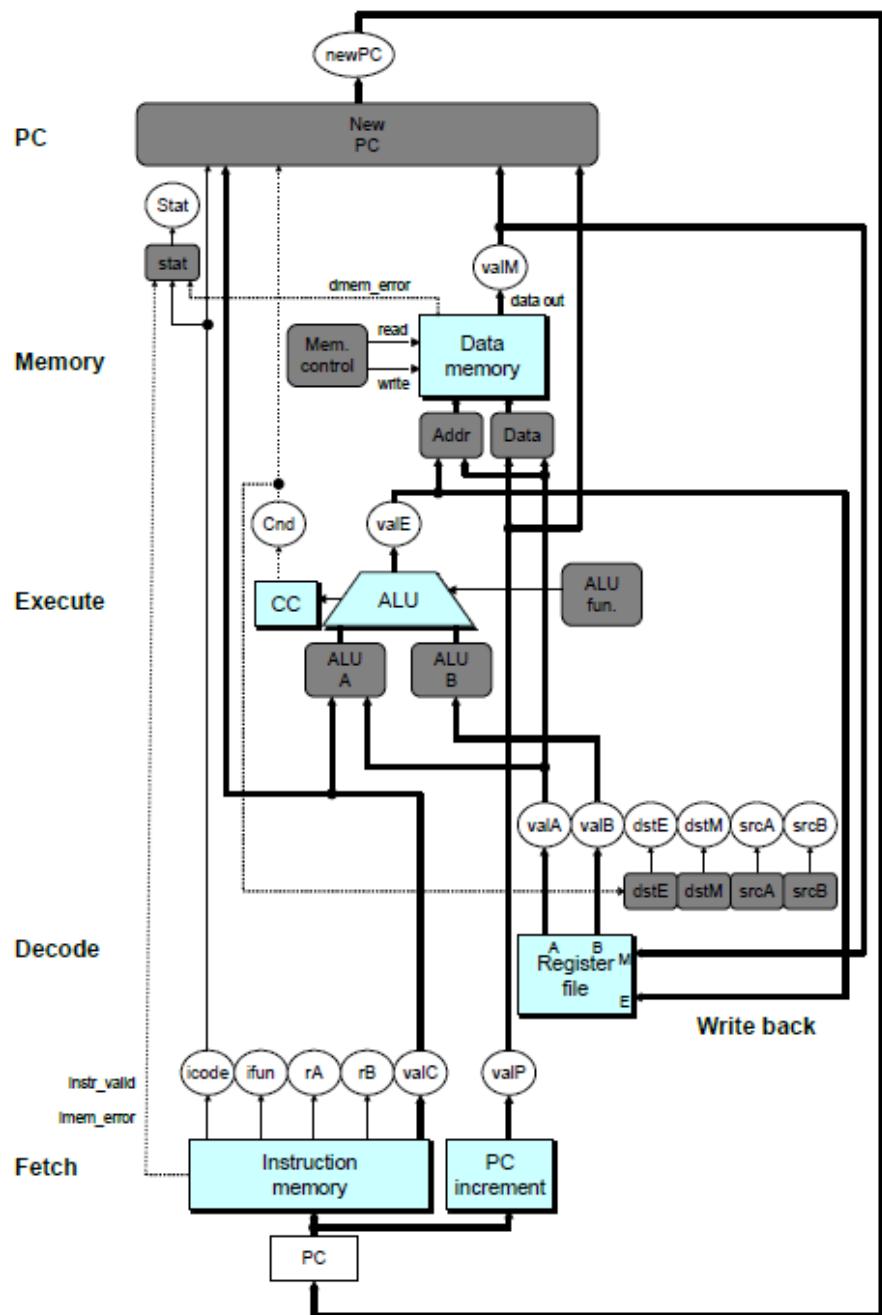
Converting SEQ to PIPE, a pipelined datapath

SEQ Hardware

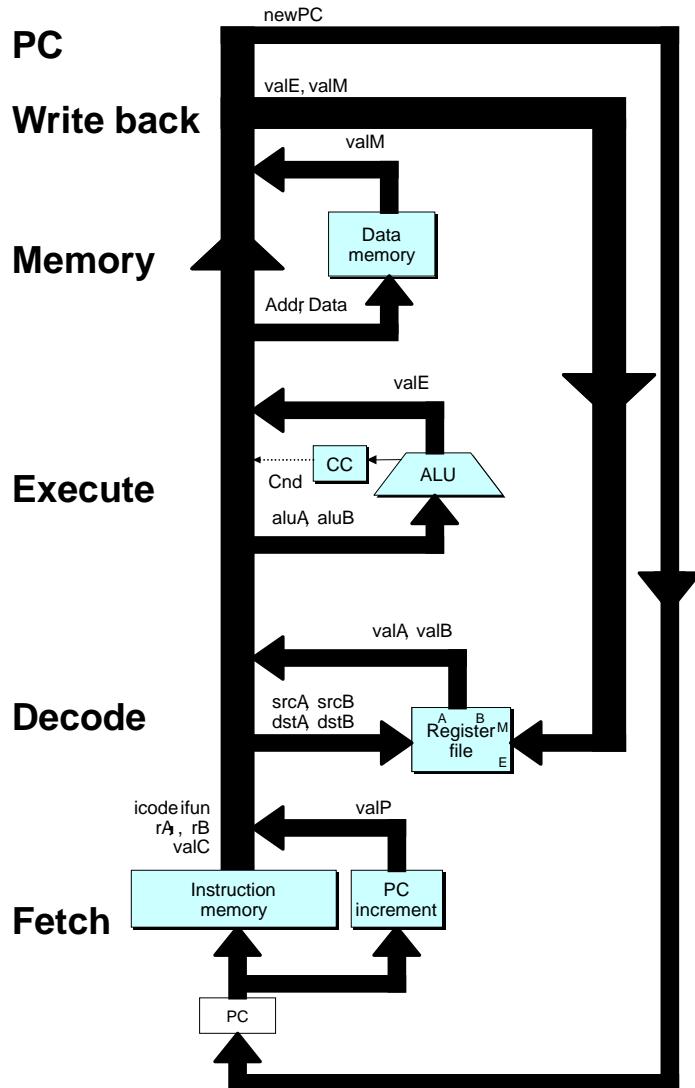
- Stages occur in sequence
- One operation in process at a time

To convert to pipelined datapath, start by adding registers between stages, resulting in 5 pipeline stages:

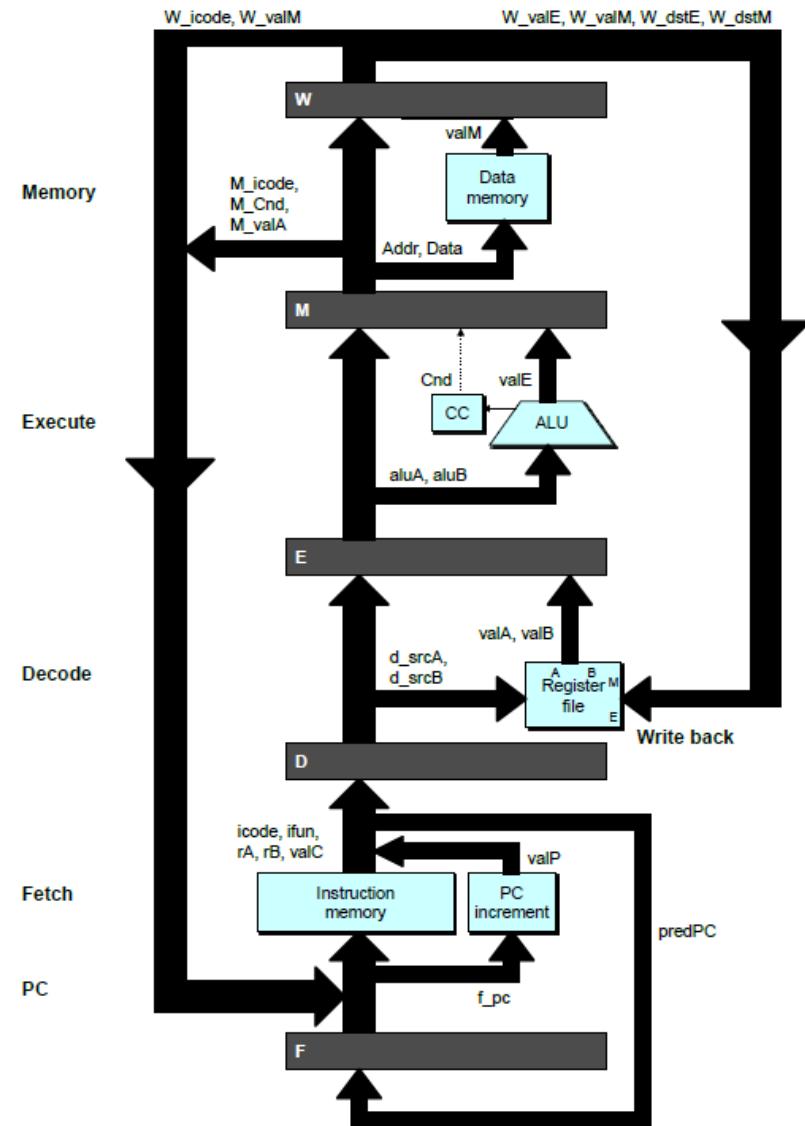
- Fetch
- Decode
- Execute
- Memory
- Writeback



Converting to pipelined datapath



Add pipeline registers between stages



Problem: Fetching a new instruction each cycle

Two problems

- PC generated in last stage of SEQ datapath
- PC sometimes not available until end of Execute or Memory stage

PC needs to be computed early

- In order to fetch a new instruction every cycle, PC generation must be moved to first stage of datapath
- Solve first problem by moving PC generation from end of SEQ to beginning of SEQ

Use prediction to select PC early

- Solve second problem by predicting next instruction from current instruction
- If prediction is wrong, squash (kill) predicted instructions

SEQ+ Hardware

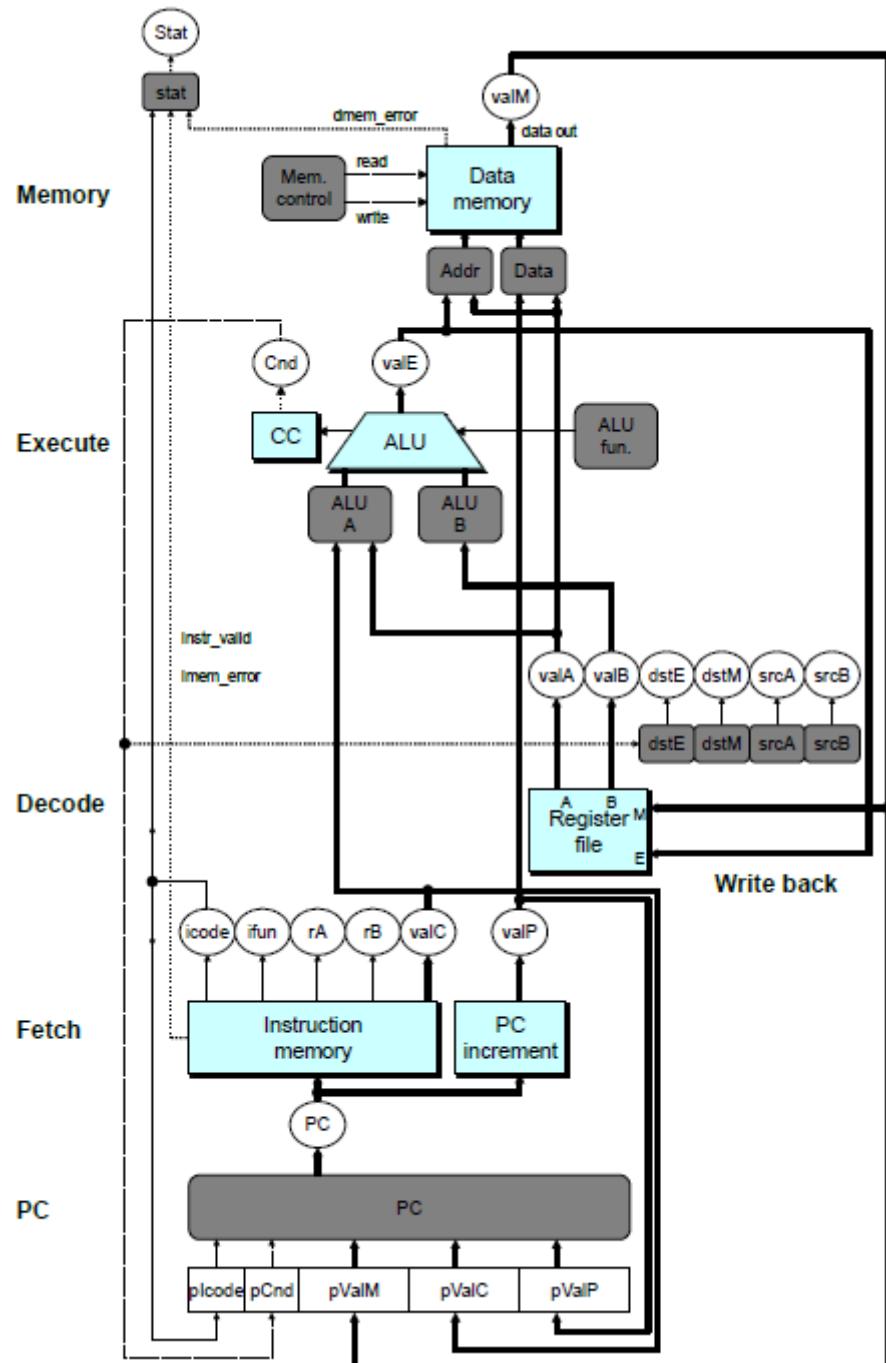
- Still sequential implementation
- Reorder PC stage to put at beginning

PC Stage

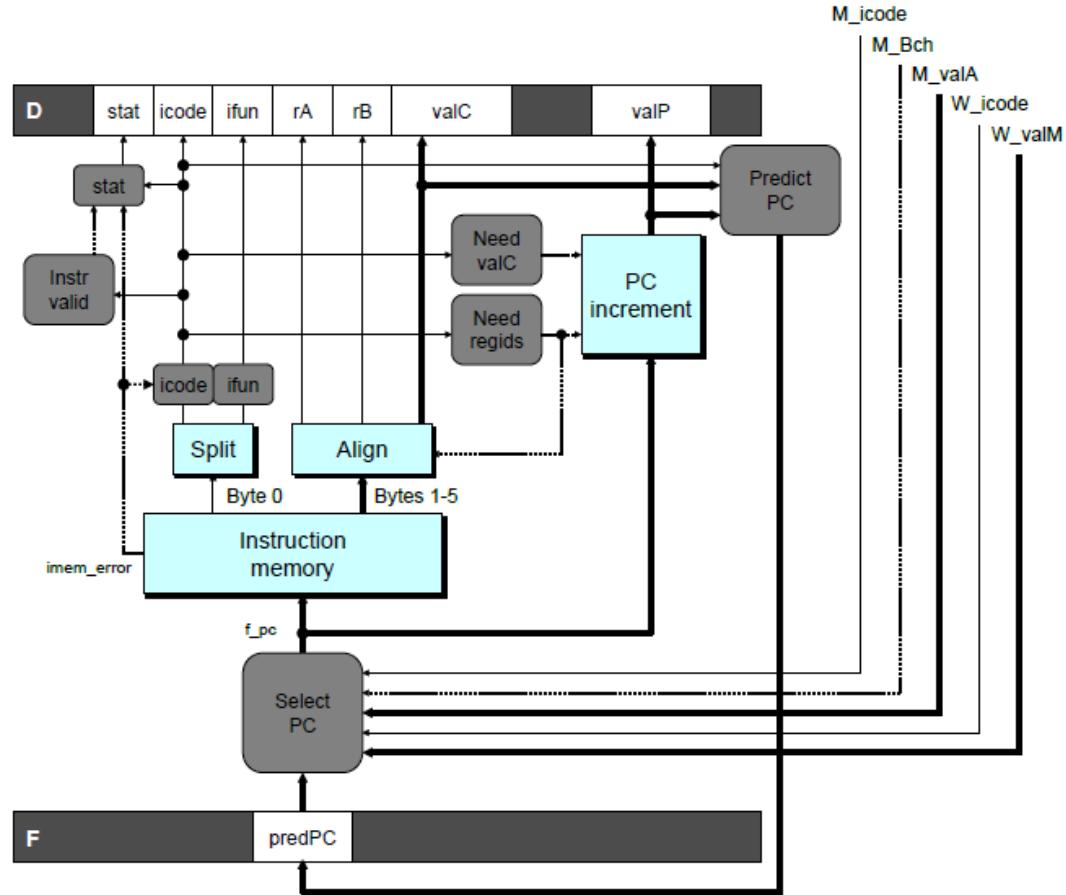
- Task is to select PC for current instruction
- Based on results computed by previous instruction

Processor State

- PC is no longer stored in register
- But, can determine PC based on other stored information



Predicting the PC



- Start fetch of new instruction after current has been fetched
 - Not enough time to fully determine next instruction
- Attempt to predict which instruction will be next
 - Recover if prediction was incorrect

Our Prediction Strategy

Predict next instruction from current instruction

Instructions that Don't Transfer Control

- Predict next PC to be valP
- Always reliable

Call and Unconditional Jumps

- Predict next PC to be valC (destination)
- Always reliable

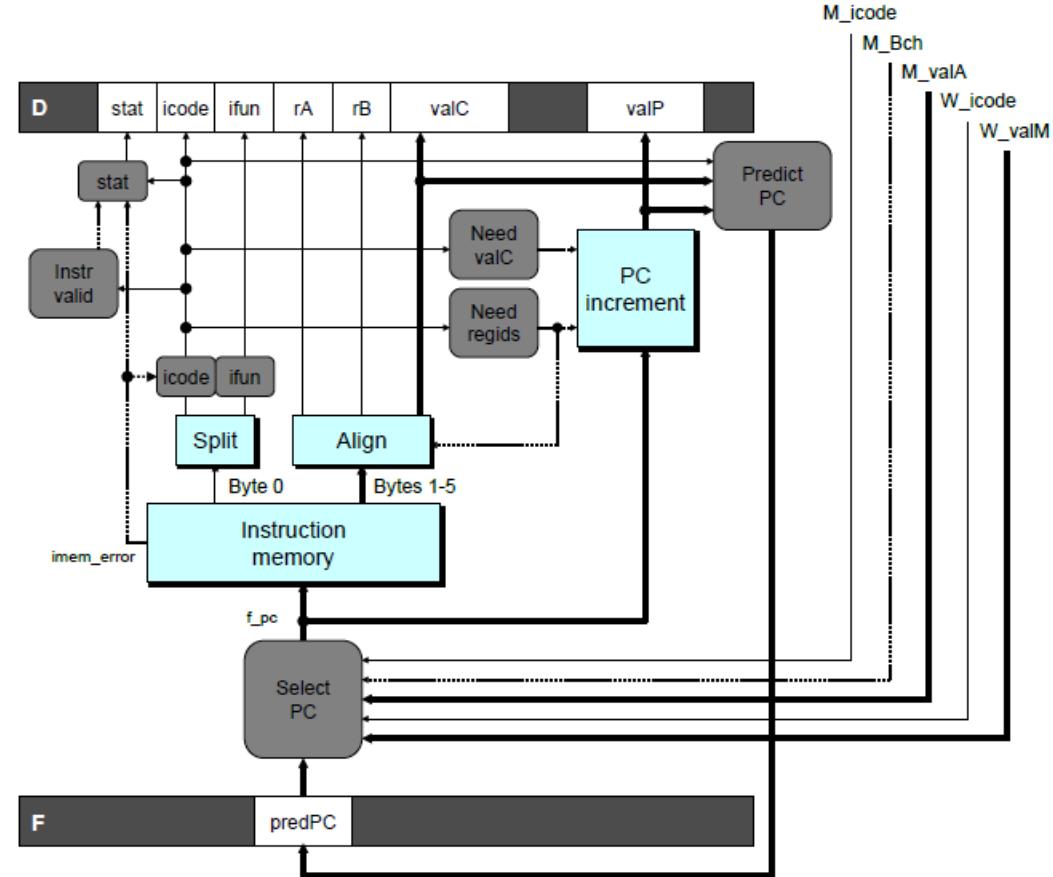
Conditional Jumps

- Predict next PC to be valC (destination)
- Only correct if branch is taken
 - Typically right 60% of time

Return Instruction

- Don't predict, just stall

Recovering from PC Misprediction



Mispredicted Jump

- Will see branch condition flag once instruction reaches memory stage
- Can get fall-through PC from valA (value M_valA)

Return Instruction

- Will get return PC when ret reaches write-back stage (W_valM)

Pipeline Stages

Fetch

- Select current PC
- Read instruction
- Compute incremented PC

Decode

- Read program registers

Execute

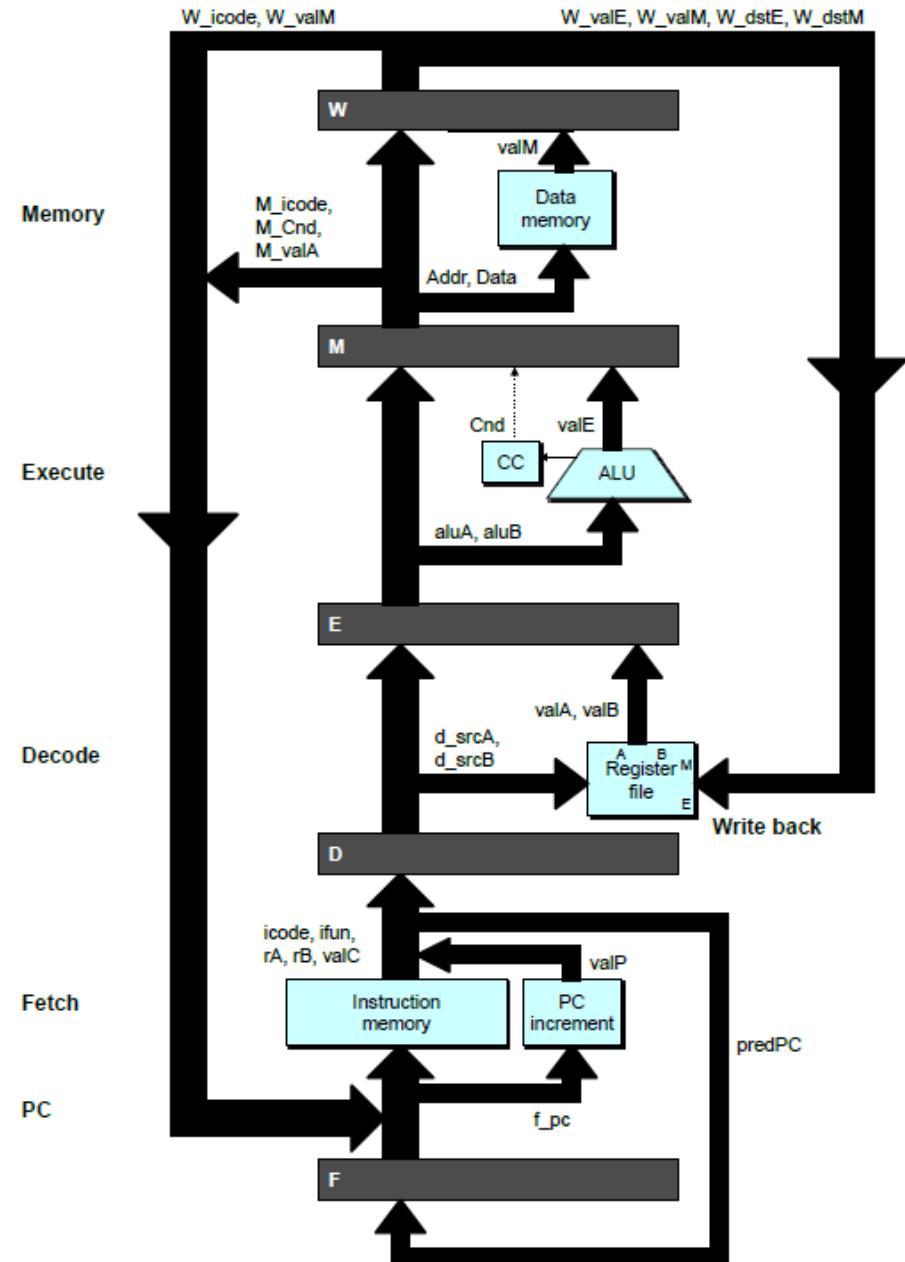
- Operate ALU

Memory

- Read or write data memory

Write Back

- Update register file

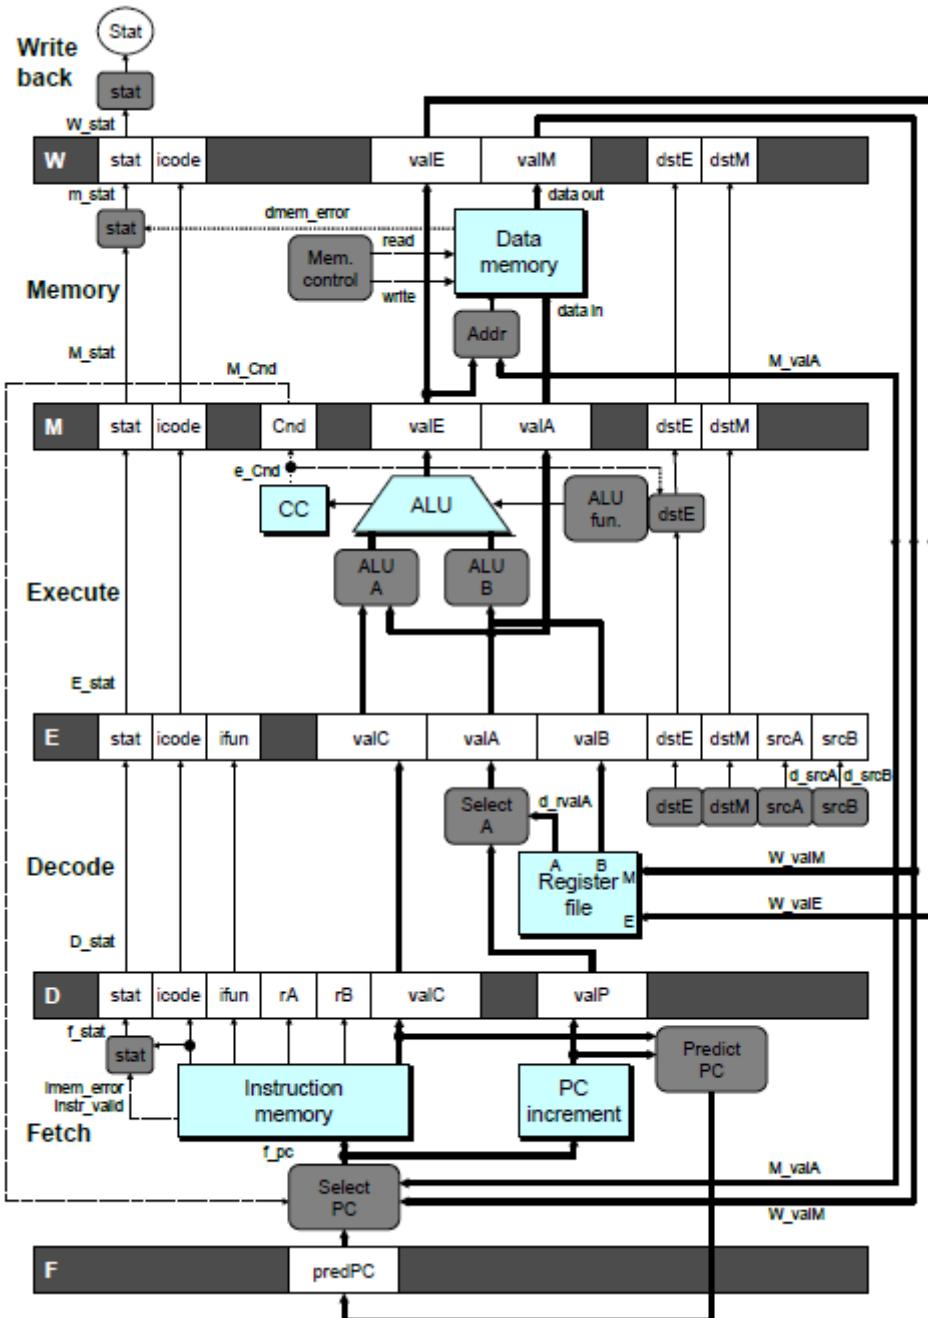


PIPE- Hardware

- Pipeline registers hold intermediate values from instruction execution

Forward (Upward) Paths

- Values passed from one stage to next
- Cannot jump past stages
 - e.g., valC passes through decode



Feedback Paths

Important for distinguishing dependencies between pipeline stages

Predicted PC

- Guess value of next PC

Branch information

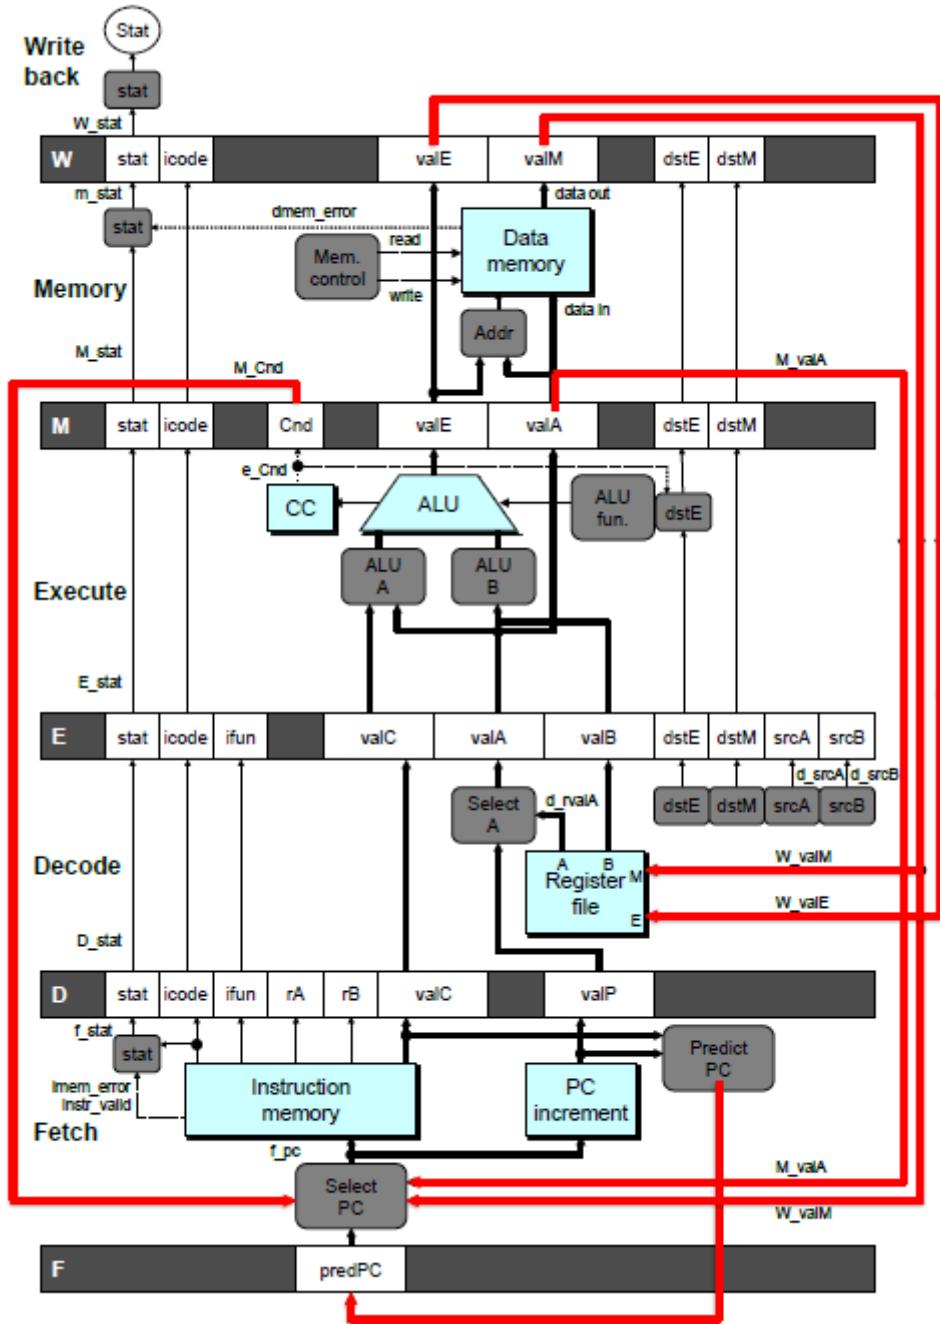
- Jump taken/not-taken
- Fall-through or target address

Return point

- Read from memory

Register updates

- To register file write ports



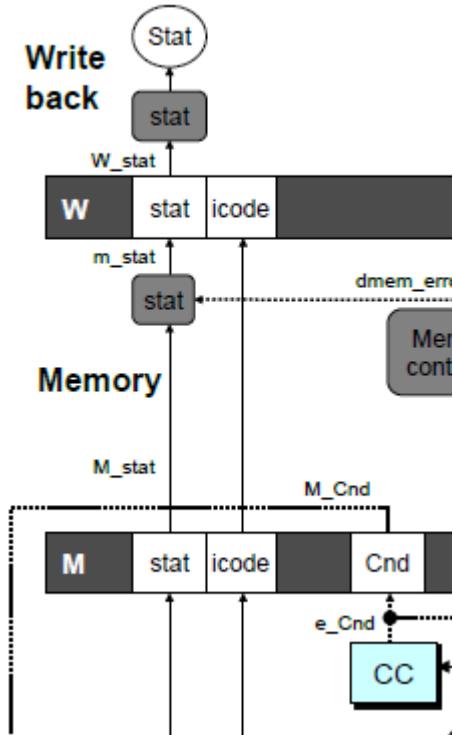
Signal Naming Conventions

S_Field

- Value of Field held in stage S pipeline register

s_Field

- Value of Field computed in stage S



Dealing with Dependencies between Instructions

Hazards

Hazards

- Problems caused by dependencies between separate instructions in the pipeline

Data Hazards

- Instruction having register R as source follows shortly after instruction having register R as destination
- Common condition, don't want to slow down pipeline

Control Hazards

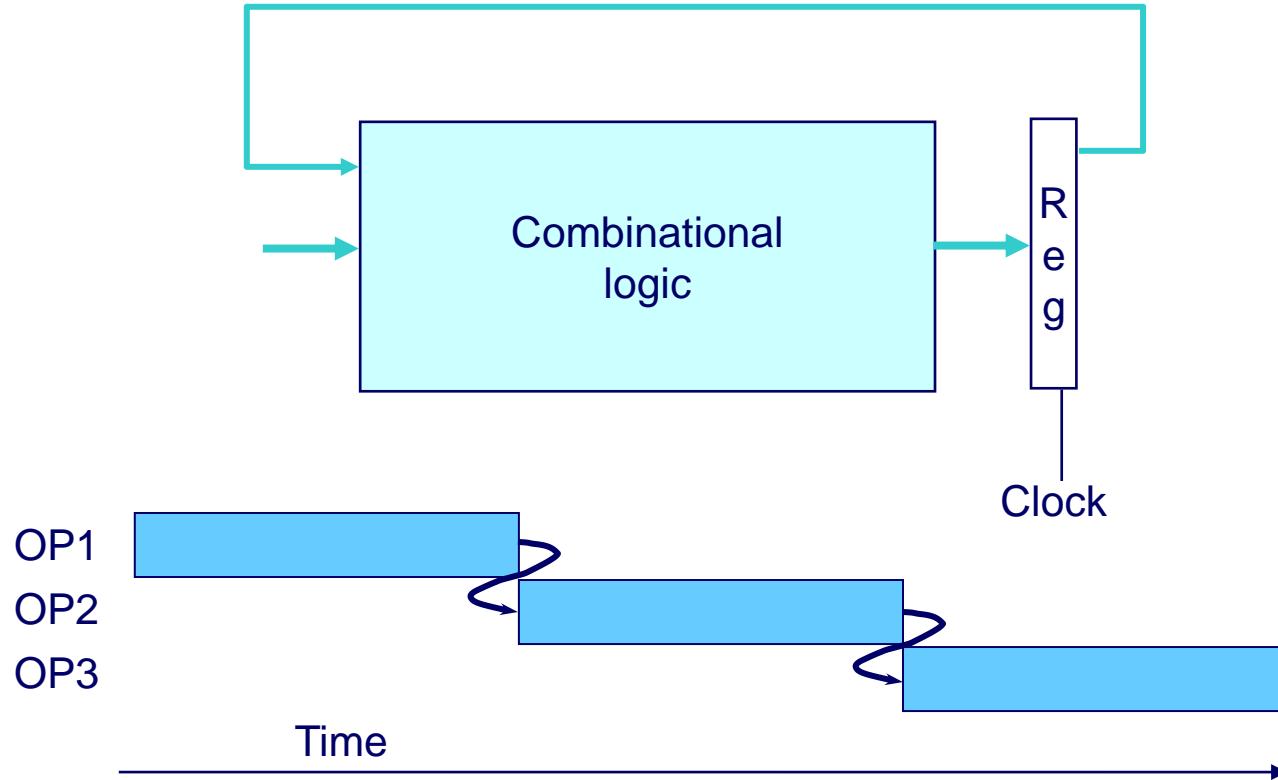
- Mispredict conditional branch
 - Our design predicts all branches as being taken
 - Naïve pipeline executes two extra instructions
- Getting return address for `ret` instruction
 - Naïve pipeline executes three extra instructions

Dealing with Dependencies between Instructions

Data Hazards

Data Dependencies

- not a problem in SEQ

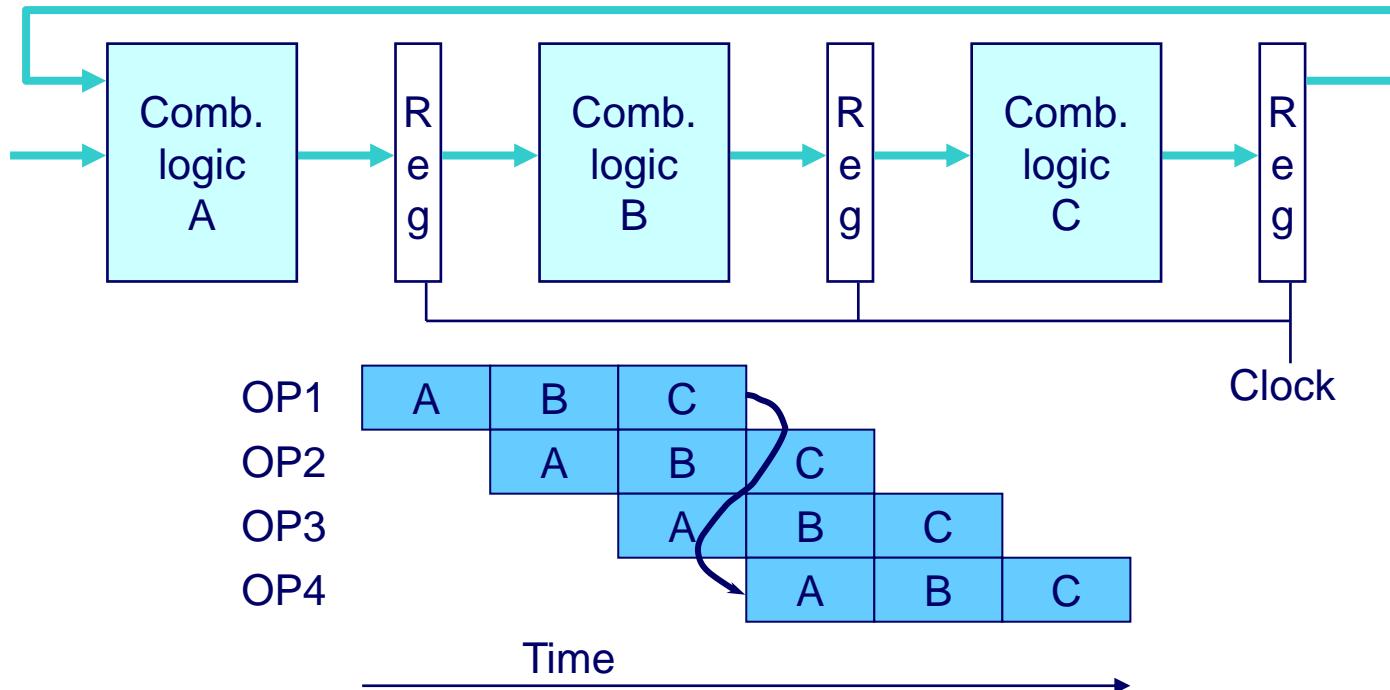


System

- Each operation depends on result from preceding one

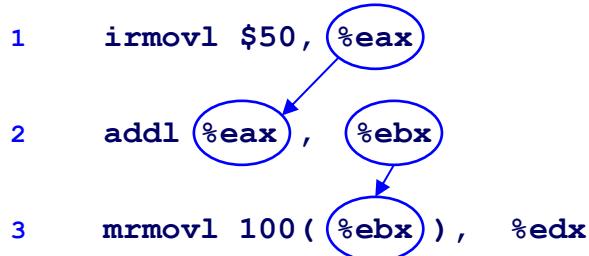
Data Hazards

- the problems caused by data dependences in pipelined datapaths



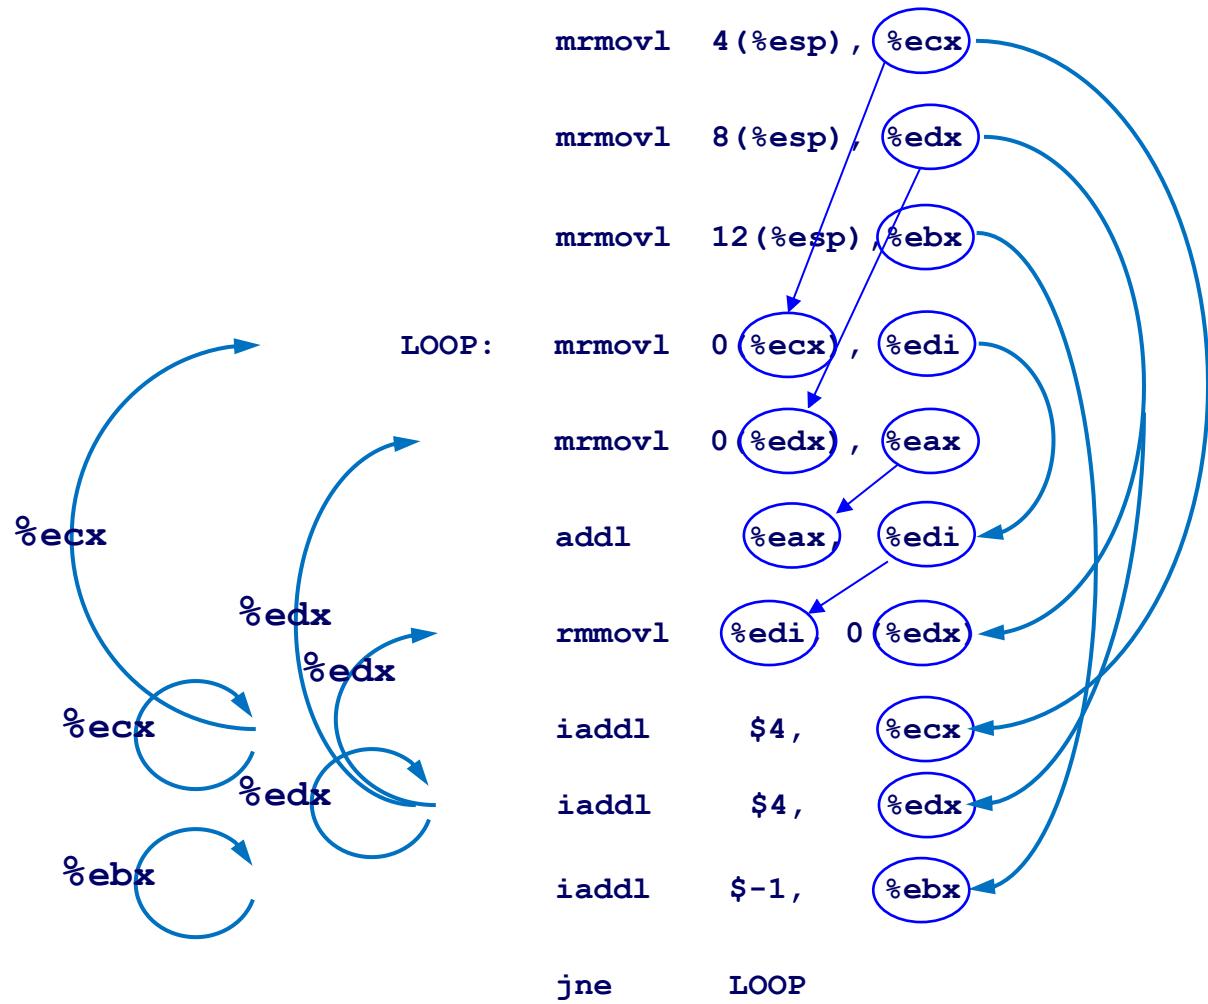
- Result does not feed back around in time for next operation
- Pipelining has changed behavior of system

Data Dependencies between Instructions



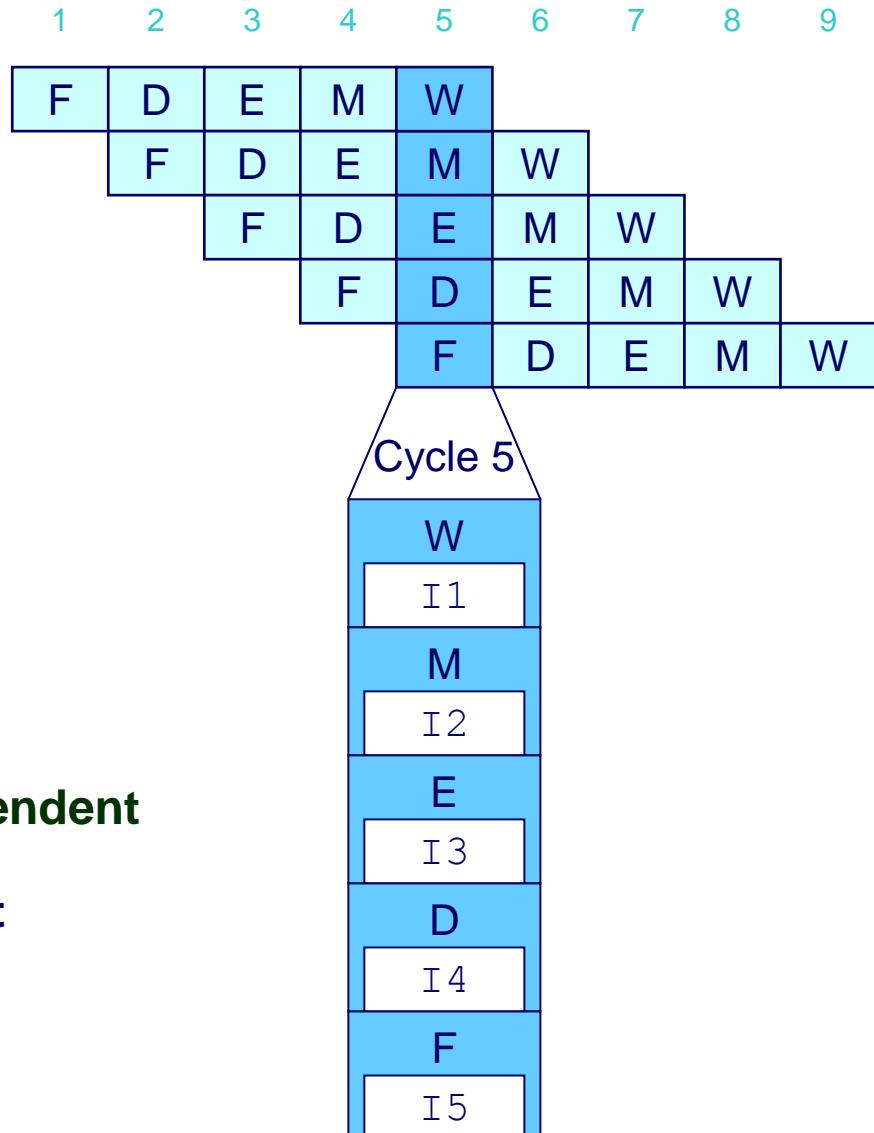
- Result from one instruction used as operand for another
 - Read-after-write (RAW) dependency
 - Dependency is between writeback stage of earlier instruction and decode stage of later instruction
- Very common in actual programs
- Must make sure our pipeline handles these properly
 - Get correct results
 - Minimize performance impact

Data Dependencies – Loop-Carried Dependencies



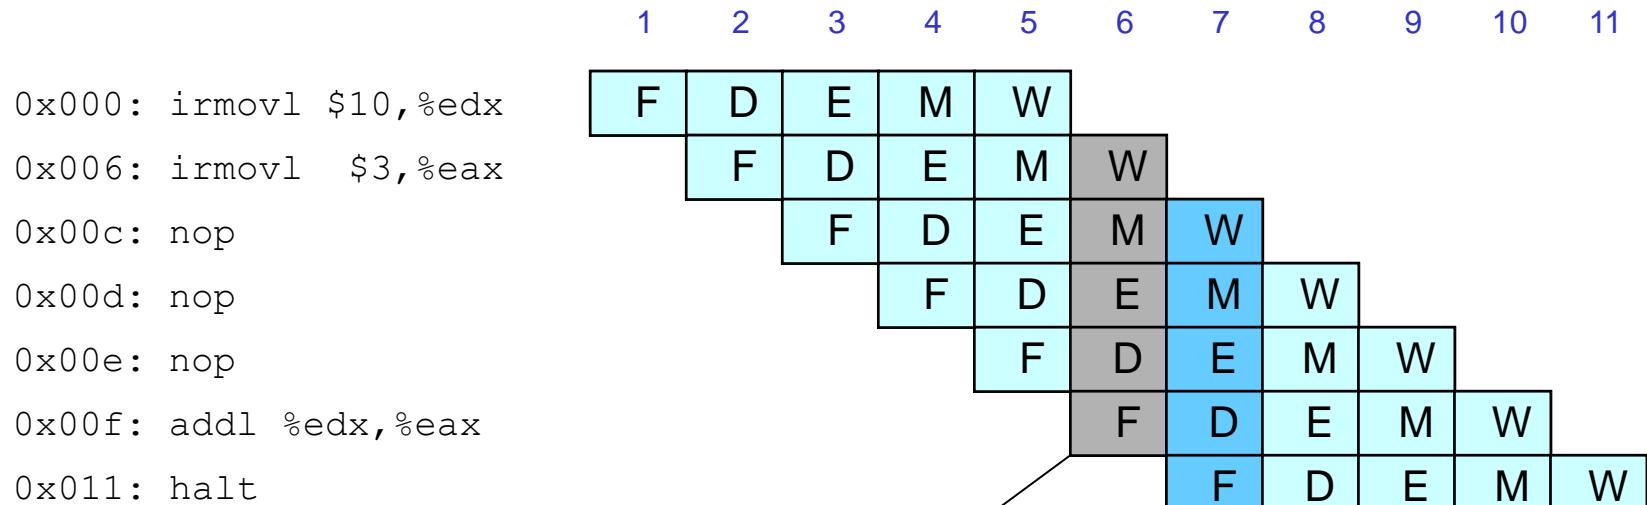
Pipeline Demonstration

irmovl	\$1,%eax	#I1
irmovl	\$2,%ecx	#I2
irmovl	\$3,%edx	#I3
irmovl	\$4,%ebx	#I4
halt		#I5



All the instructions are independent
of each other
- No dependencies exist

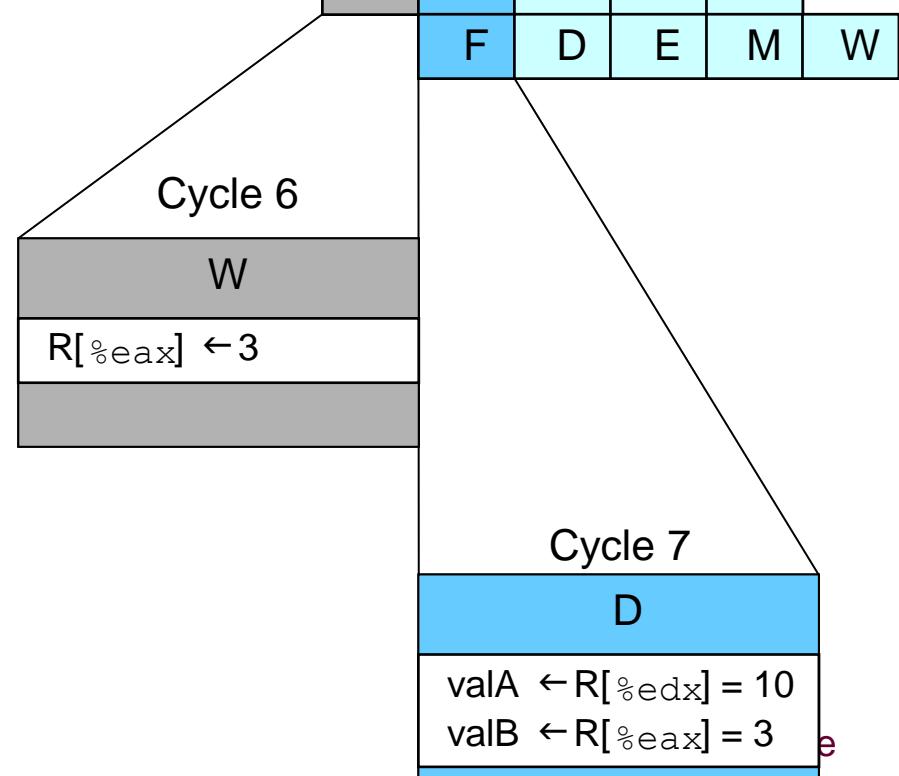
Data Dependencies: 3 Nop's



The addl instruction depends on the first two instructions

- addl depends upon %edx from the 1st instr
- addl depends upon %eax from the 2nd instr

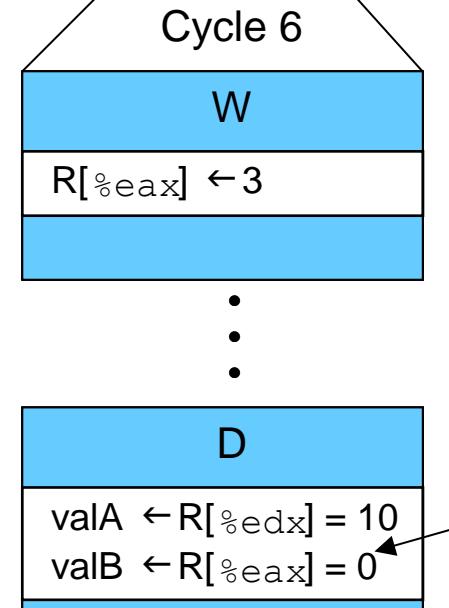
addl must wait 3 cycles after the 2nd instruction, so that it doesn't fetch the two registers before they've been written to the register file



Data Dependencies: 2 Nop's

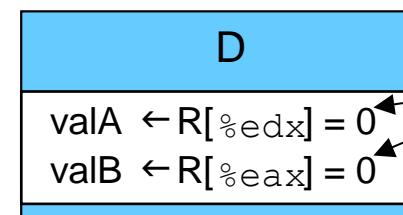
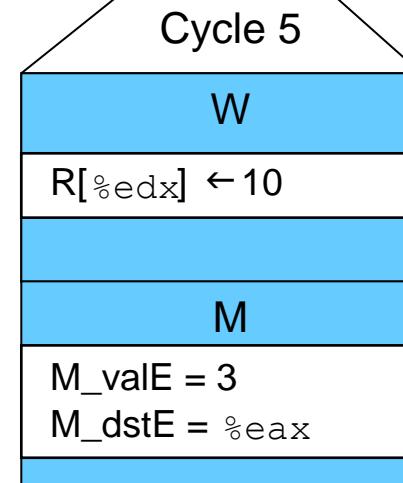
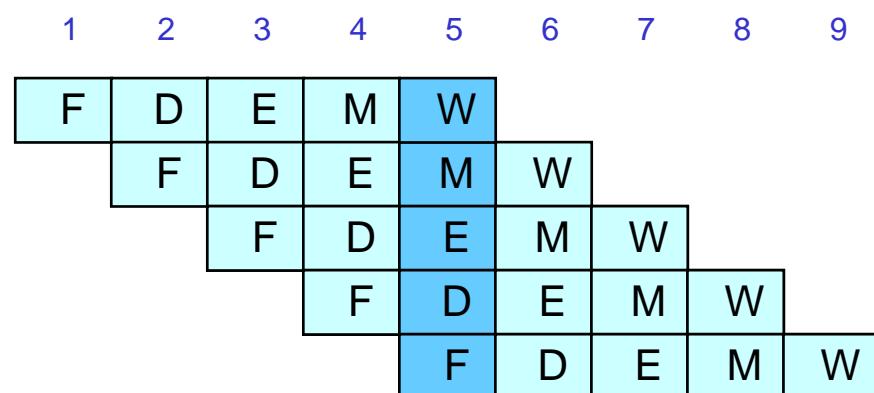
	1	2	3	4	5	6	7	8	9	10
0x000: irmovl \$10,%edx	F	D	E	M	W					
0x006: irmovl \$3,%eax		F	D	E	M	W				
0x00c: nop			F	D	E	M	W			
0x00d: nop				F	D	E	M	W		
0x00e: addl %edx,%eax					F	D	E	M	W	
0x010: halt						F	D	E	M	W

If addl executes one cycle earlier,
it gets the wrong value for %eax



Data Dependencies: 1 Nop

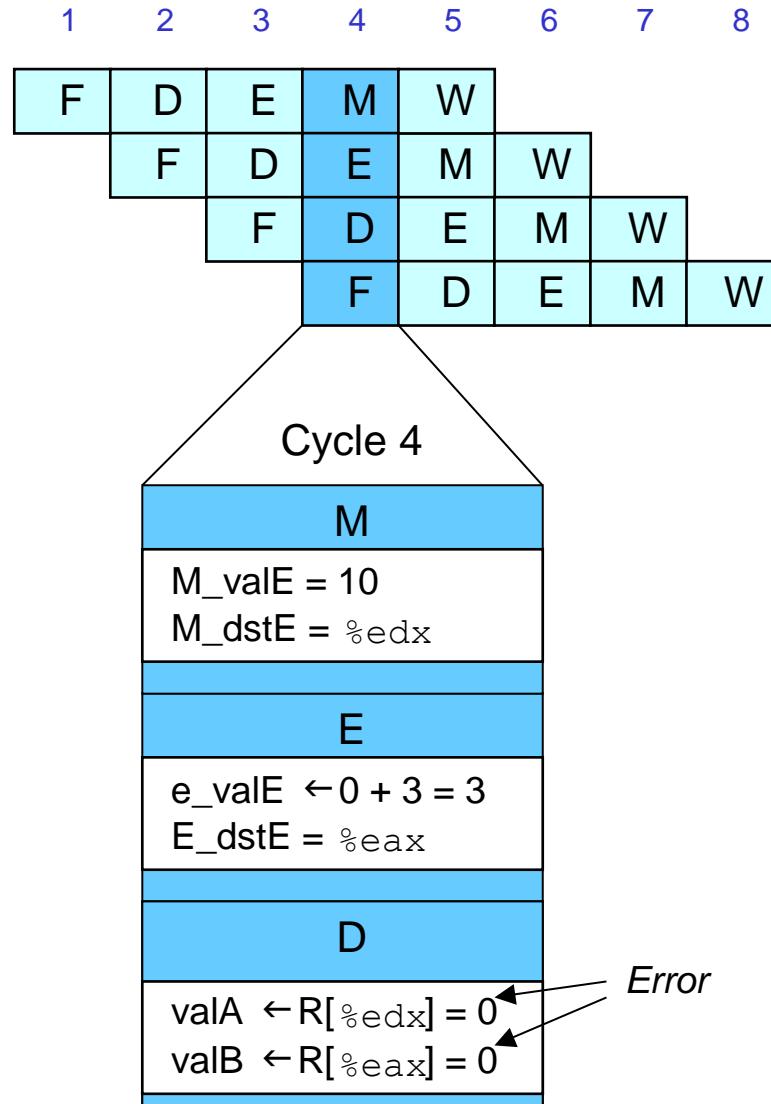
	1	2	3	4	5	6	7	8	9
0x000: irmovl \$10,%edx	F	D	E	M	W				
0x006: irmovl \$3,%eax	F	D	E	M	W				
0x00c: nop	F	D	E	M	W				
0x00d: addl %edx,%eax	F	D	E	M	W				
0x00f: halt	F	D	E	M	W				



If addl executes two cycles earlier,
it gets the wrong value for both
%eax and %edx

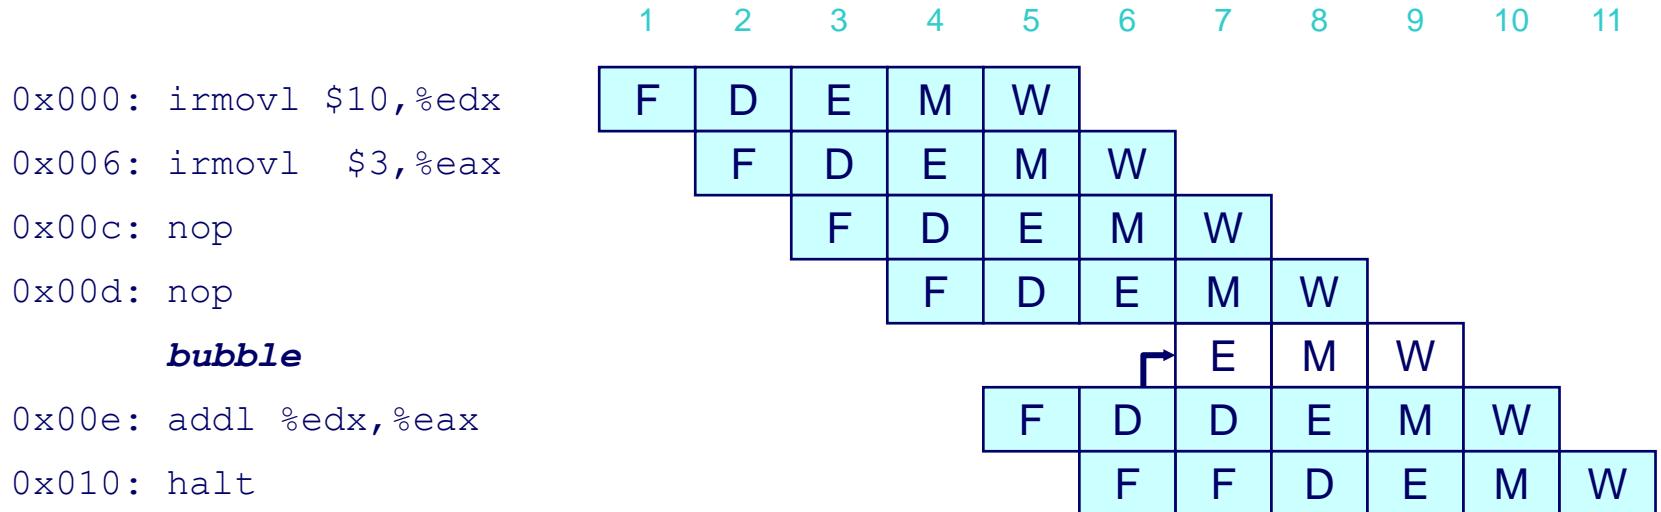
Data Dependencies: No Nop

```
0x000: irmovl $10,%edx  
0x006: irmovl $3,%eax  
0x00c: addl %edx,%eax  
0x00e: halt
```



Like the prior case, if addl executes three cycles earlier, it gets the wrong value for both %eax and %edx

Stalling for Data Dependencies



- If instruction follows too closely after one that writes register, slow it down
- Hold instruction in decode
- Dynamically inject nop into execute stage

Stall Condition

Source Registers

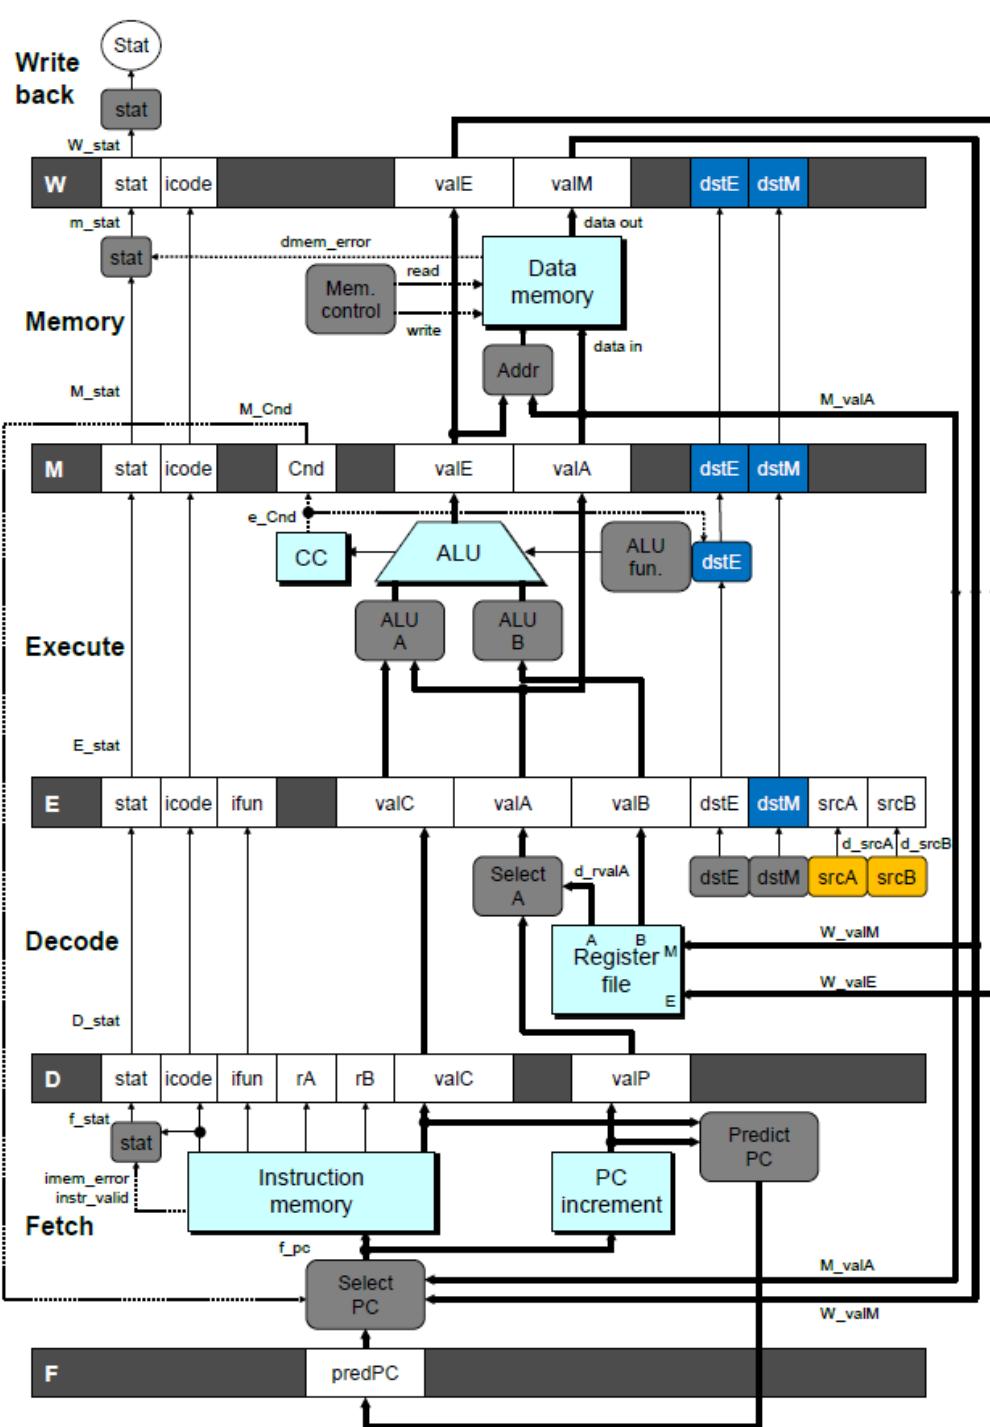
- srcA and srcB of current instruction in decode stage

Destination Registers

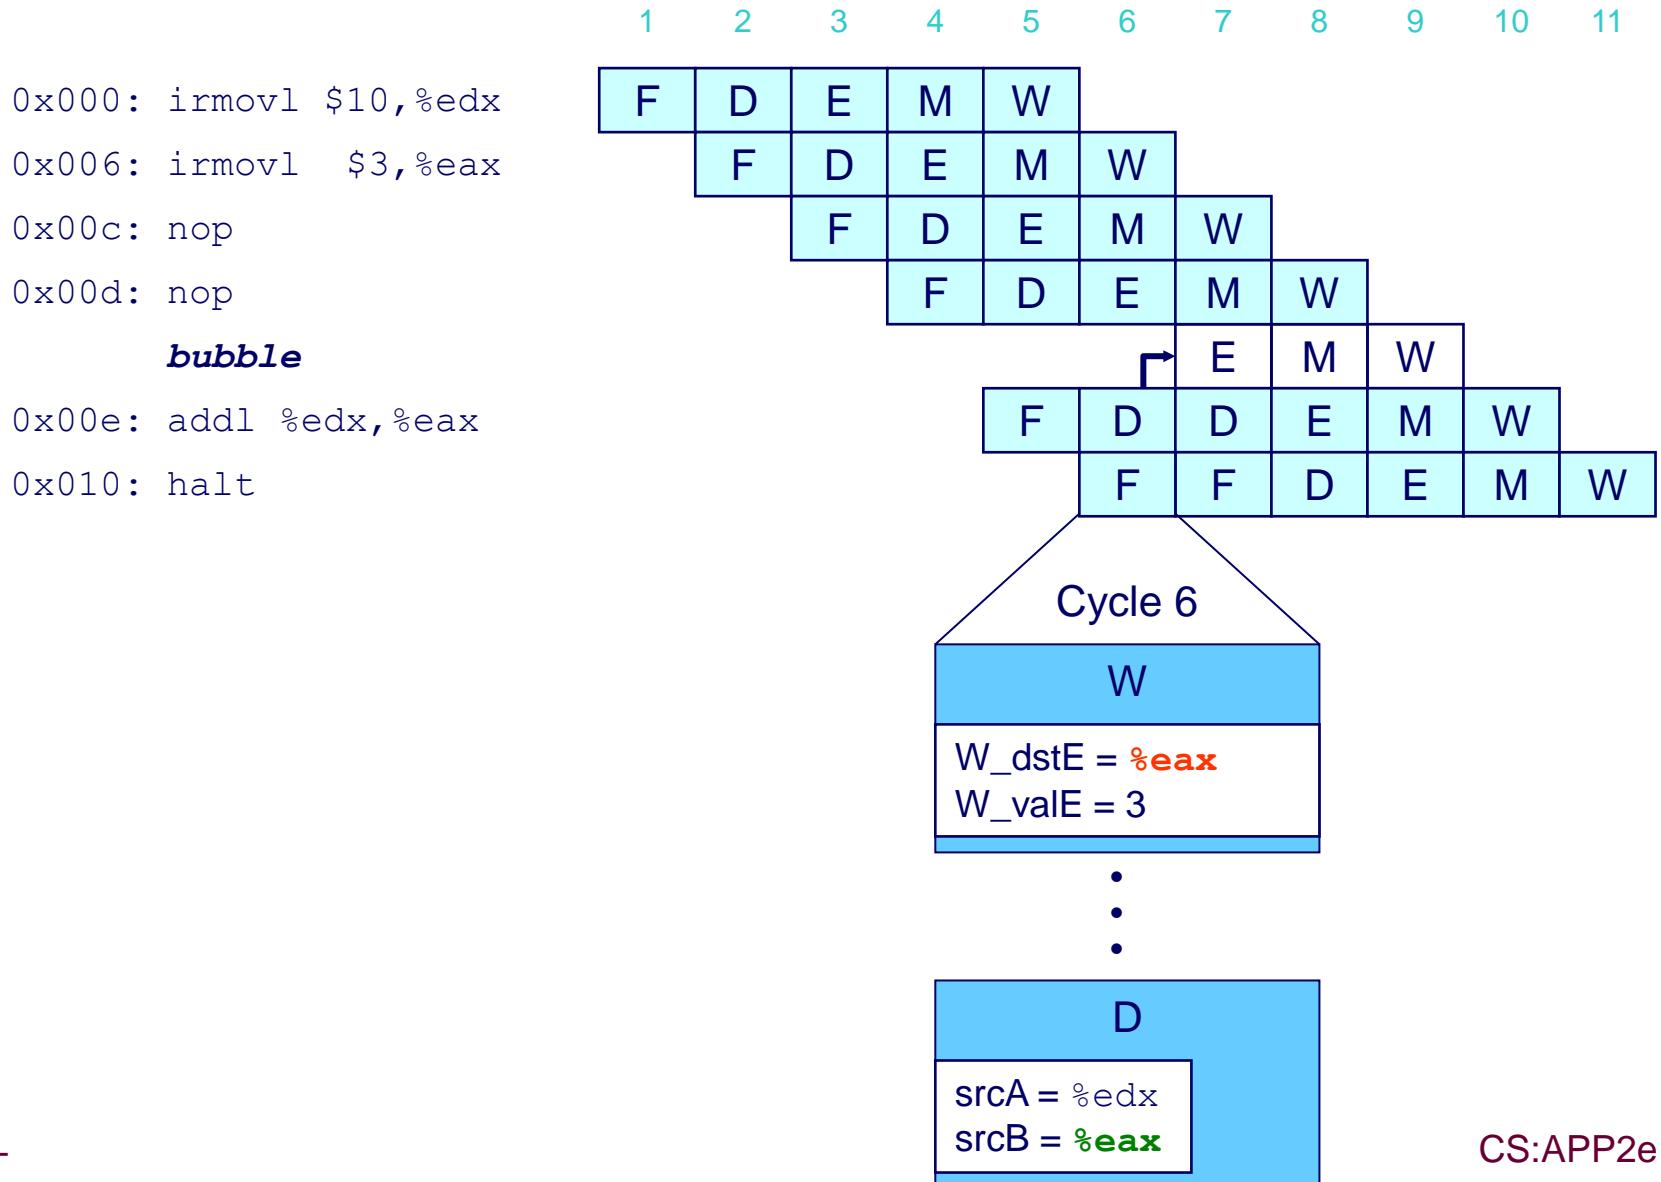
- dstE and dstM fields
- Instructions in execute, memory, and write-back stages

Special Case

- Don't stall for register ID 15 (0xF)
 - Indicates absence of register operand
- Don't stall for failed conditional move



Detecting Stall Condition

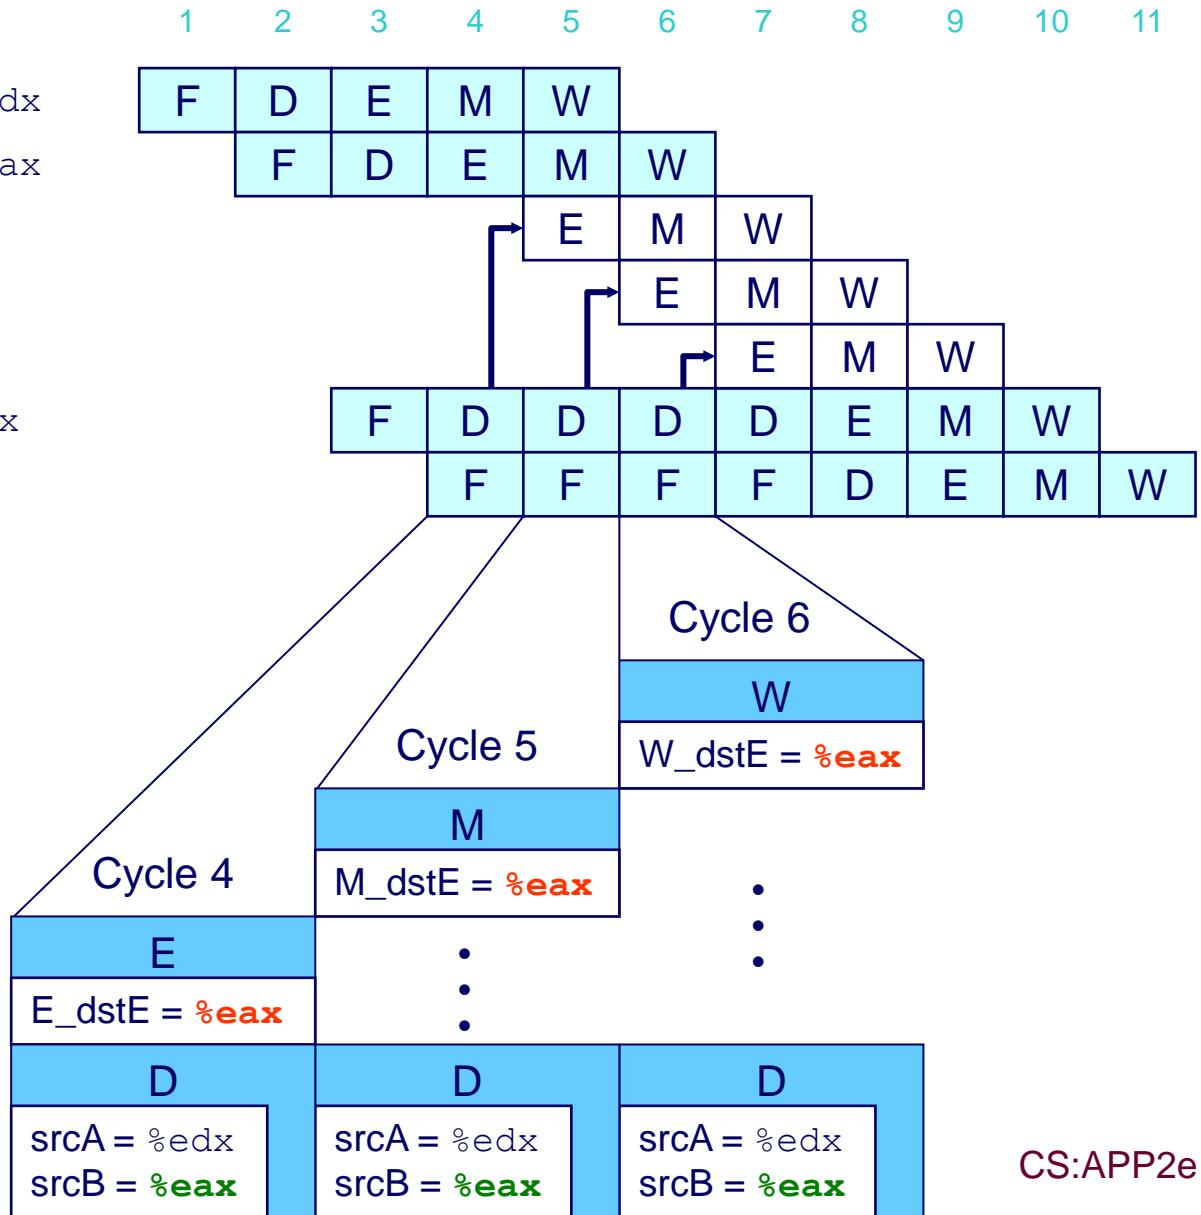


Stalling X3

```

0x000: irmovl $10,%edx
0x006: irmovl $3,%eax
bubble
bubble
bubble
0x00c: addl %edx,%eax
0x00e: halt

```



What Happens When Stalling?

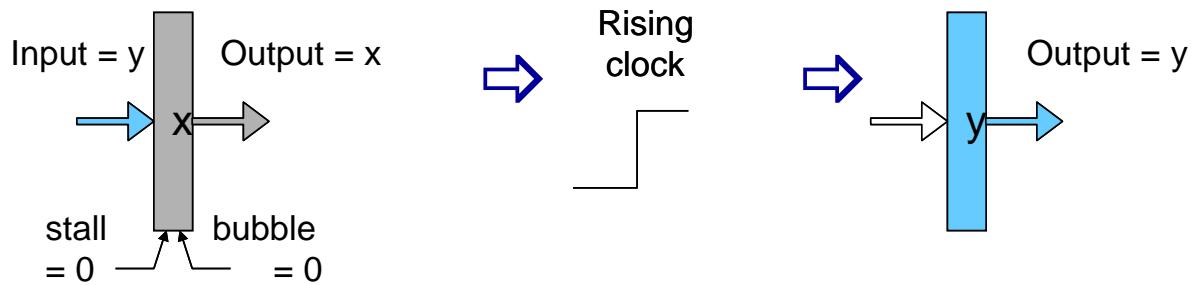
```
0x000: irmovl $10,%edx  
0x006: irmovl $3,%eax  
0x00c: addl %edx,%eax  
0x00e: halt
```

Cycle 8	
Write Back	<i>bubble</i>
Memory	<i>bubble</i>
Execute	0x00c: addl %edx,%eax
Decode	0x00e: halt
Fetch	

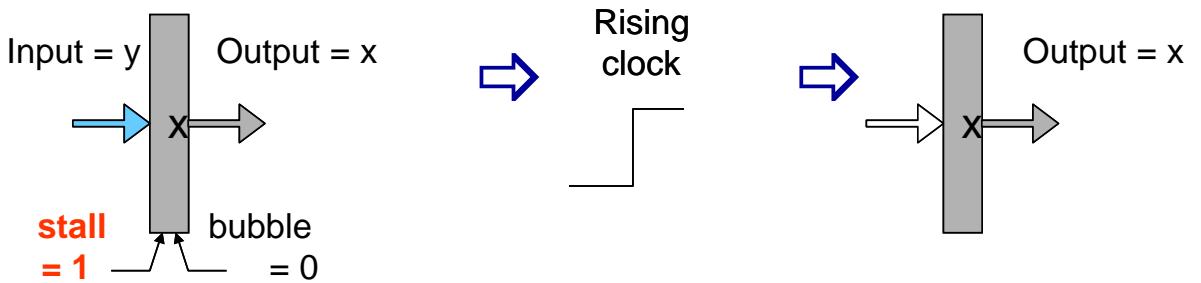
- Stalling instruction held back in decode stage
- Following instruction stays in fetch stage
- Bubbles injected into execute stage
 - Like dynamically generated nop's
 - Move through later stages

Pipeline Register Modes

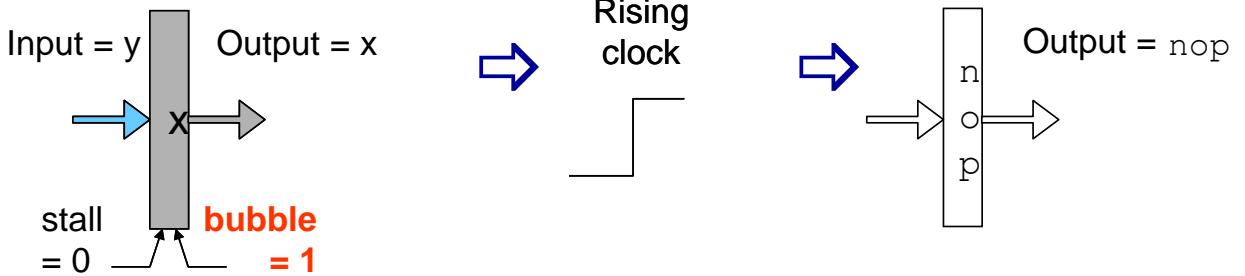
Normal



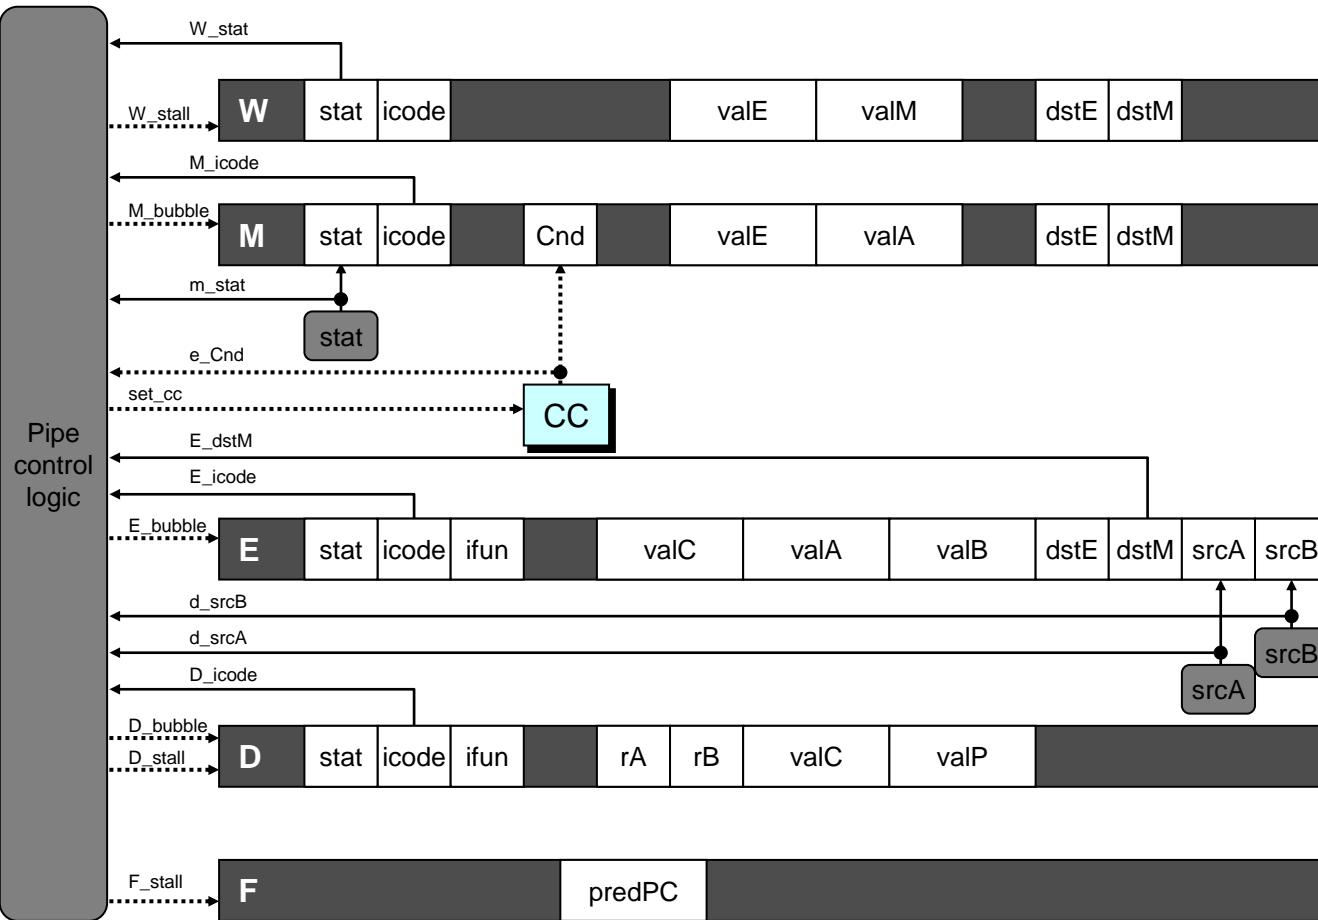
Stall



Bubble



Implementing Stalling



Pipeline Control

- Combinational logic detects stall condition
- Sets mode signals for how pipeline registers should update

Data Forwarding

Naïve Pipeline

- Register isn't written until completion of write-back stage
- Source operands read from register file in decode stage
 - Needs to be in register file at start of stage

Observation

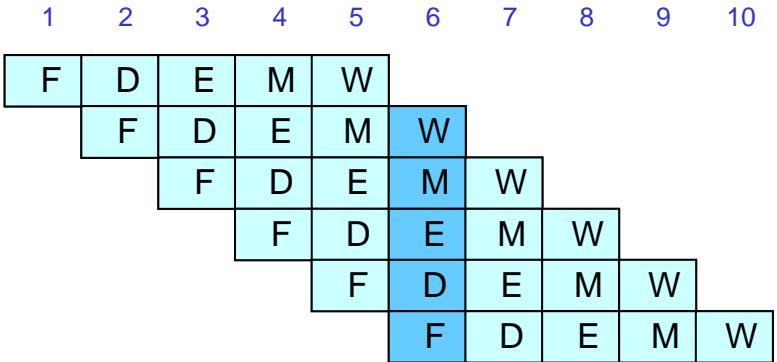
- Value generated in execute or memory stage

Trick

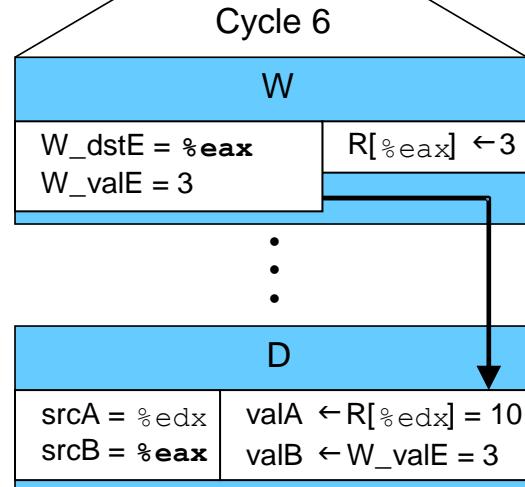
- Pass value directly from generating instruction to decode stage
- Needs to be available at end of decode stage

Data Forwarding Example

```
0x000: irmovl $10,%edx  
0x006: irmovl $3,%eax  
0x00c: nop  
0x00d: nop  
0x00e: addl %edx,%eax  
0x010: halt
```



- **irmovl in write-back stage**
- **Destination value in W pipeline register**
- **Forward as valB for decode stage**



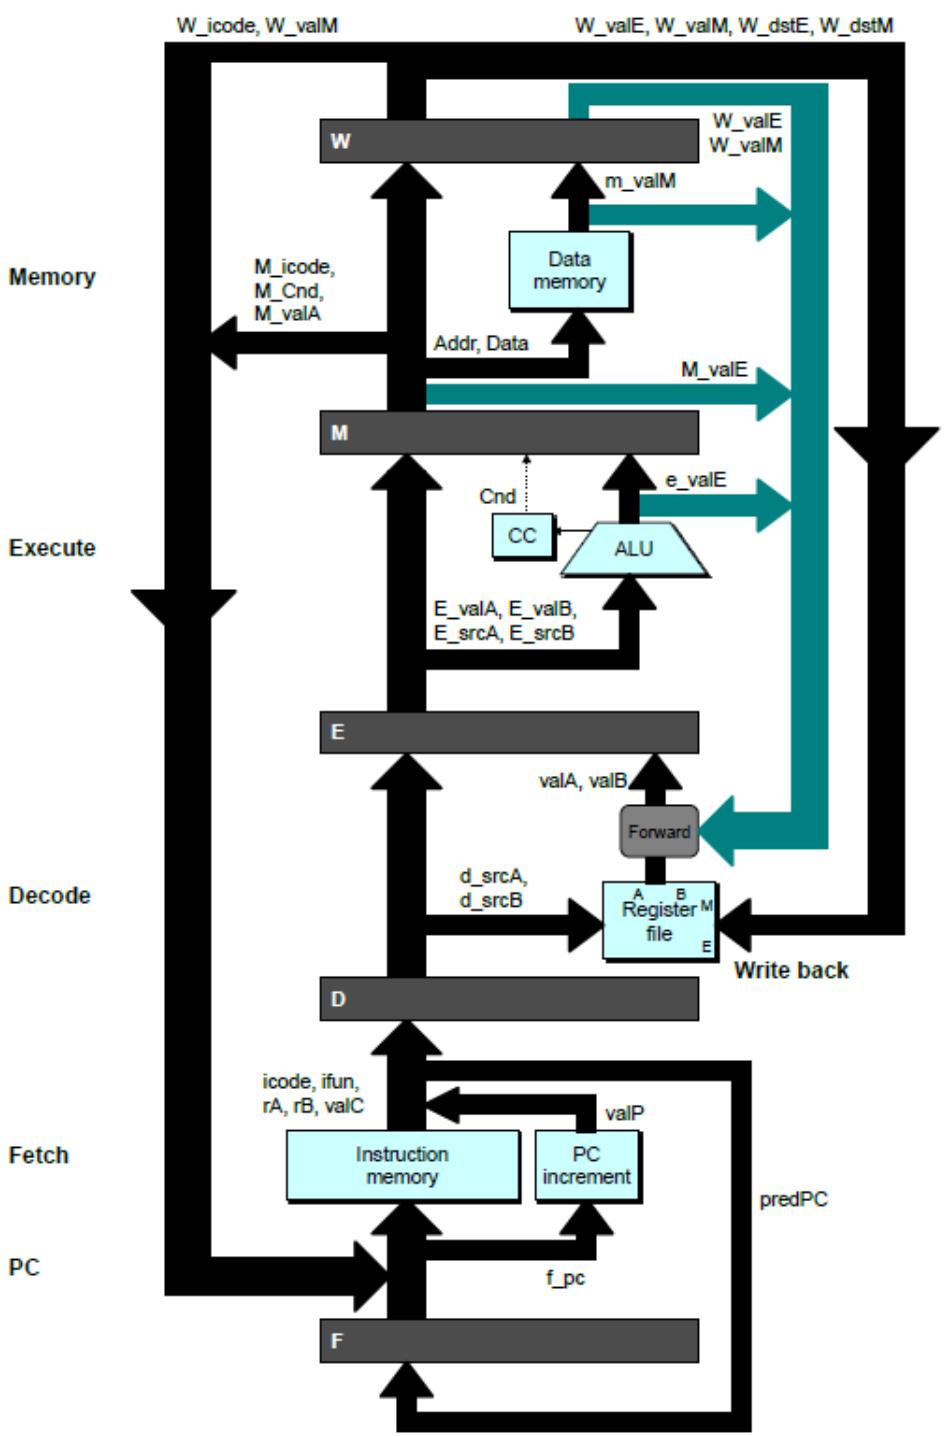
Bypass Paths

Decode Stage

- Forwarding logic selects valA and valB
- Normally from register file
- Forwarding: get valA or valB from later pipeline stage

Forwarding Sources

- Execute: valE
- Memory: valE, valM
- Write back: valE, valM



Data Forwarding Example #2

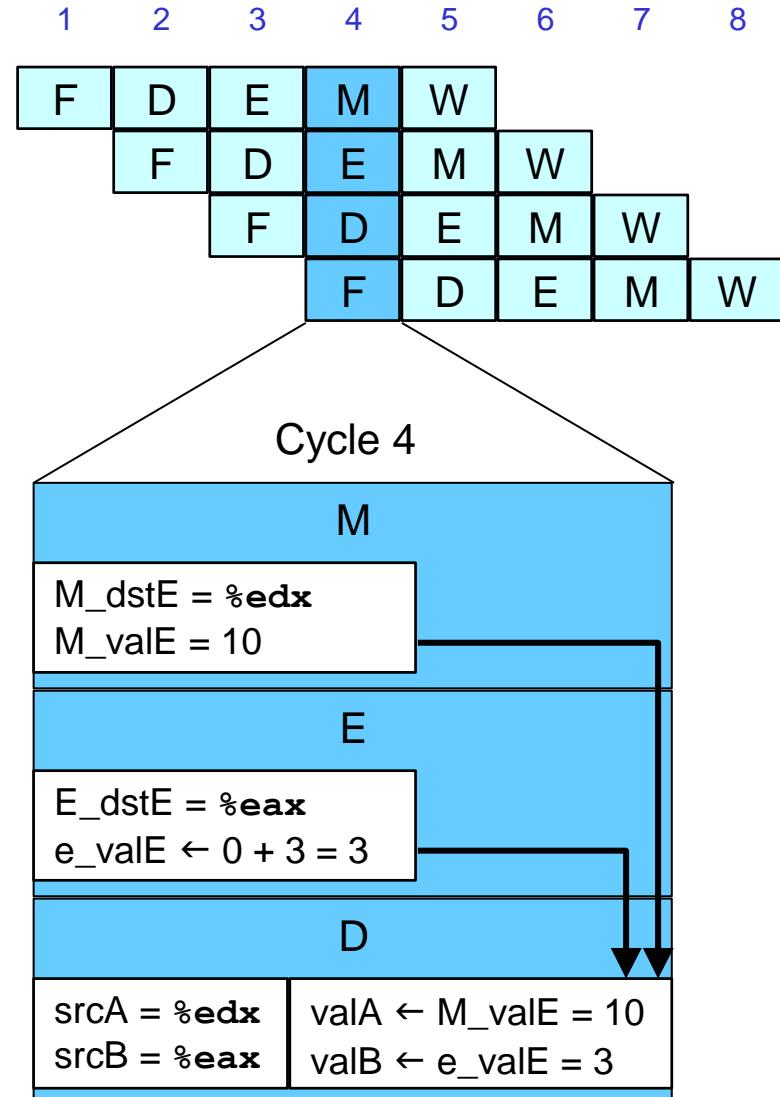
```
# demo-h0.ys
0x000: irmovl $10,%edx
0x006: irmovl $3,%eax
0x00c: addl %edx,%eax
0x00e: halt
```

Register %edx

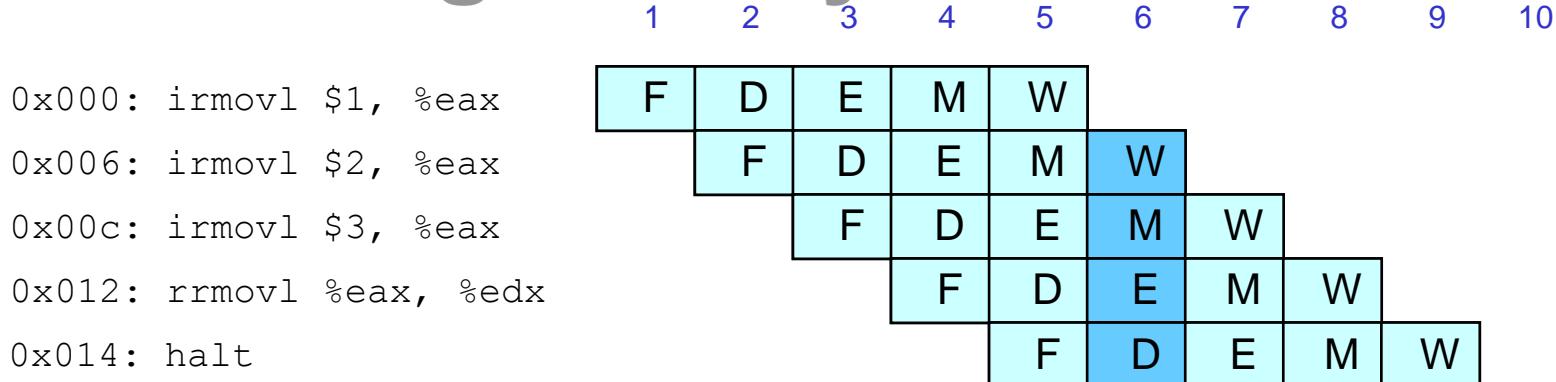
- Generated by ALU during previous cycle
- Forward from memory as valA

Register %eax

- Value just generated by ALU
- Forward from execute as valB

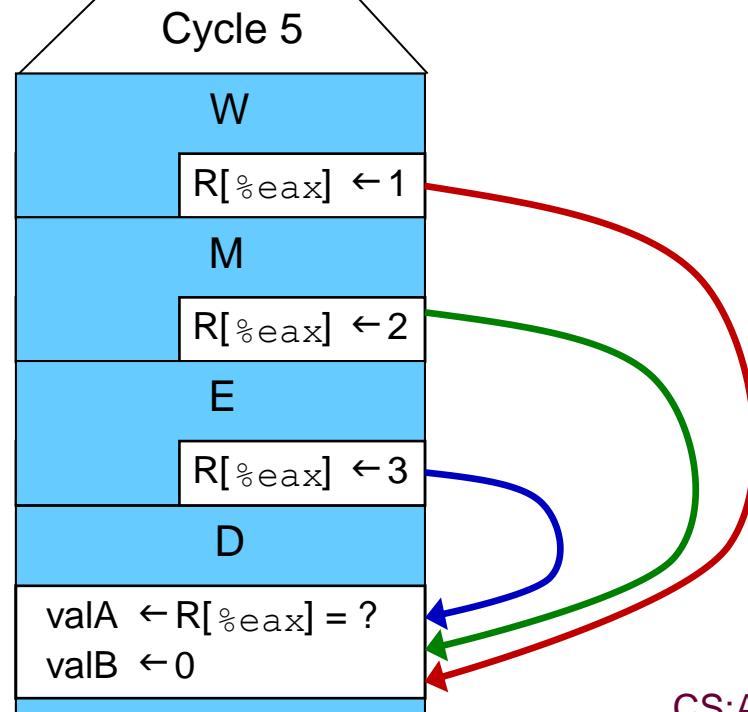


Forwarding Priority

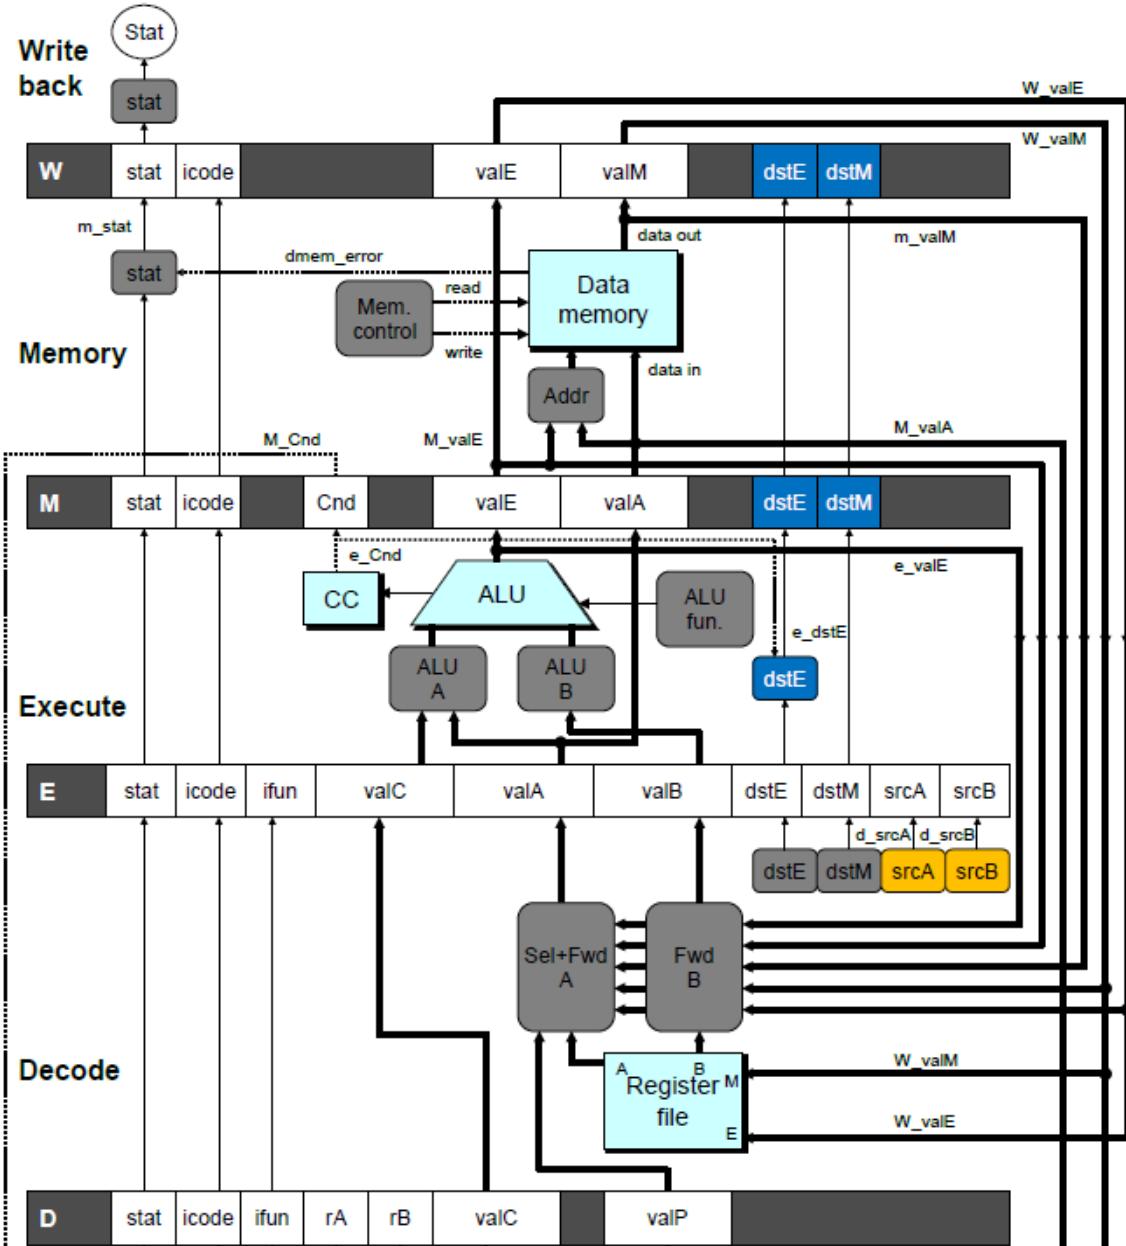


Multiple Forwarding Choices

- Which one should have priority
- Match serial semantics
- Use matching value from earliest pipeline stage



Implementing Forwarding

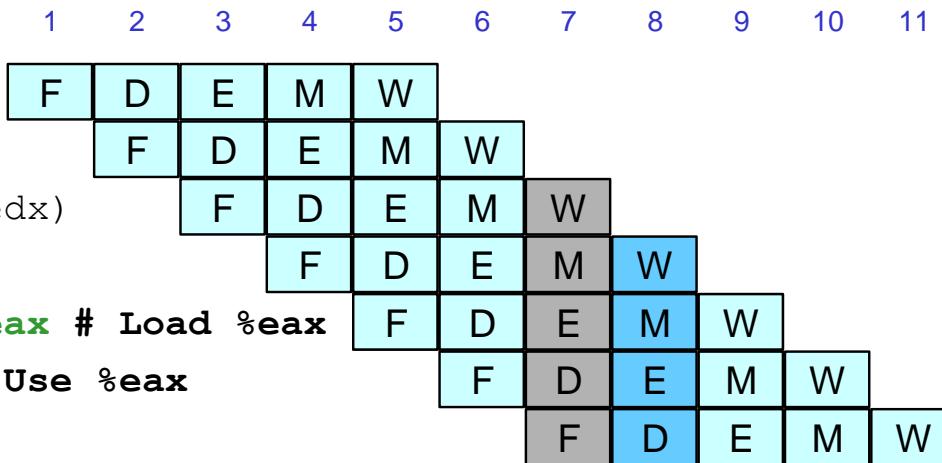


- Add additional feedback paths from E, M, and W pipeline registers into decode stage
- Create logic blocks to select from multiple sources for valA and valB in decode stage

Limitation of Forwarding

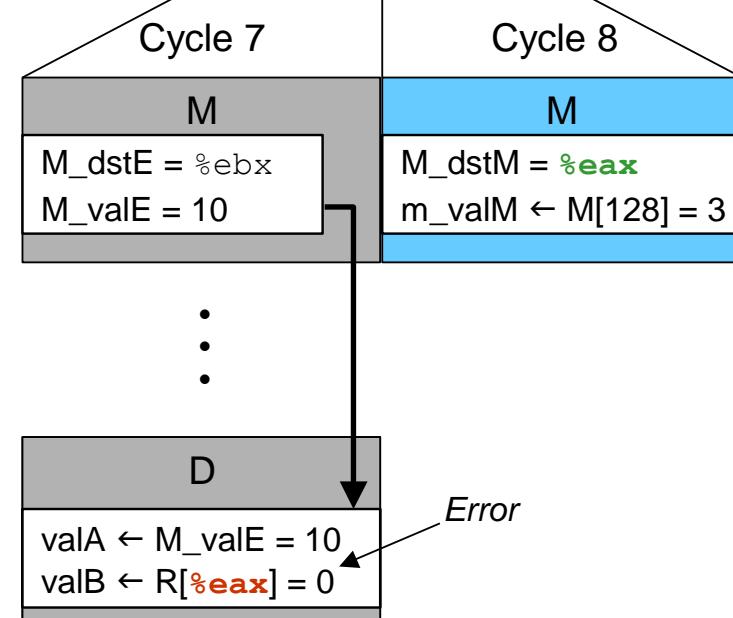
```
# demo-luh.ys
```

```
0x000: irmovl $128,%edx  
0x006: irmovl $3,%ecx  
0x00c: rmmovl %ecx, 0(%edx)  
0x012: irmovl $10,%ebx  
0x018: mrmovl 0(%edx),%eax # Load %eax  
0x01e: addl %ebx,%eax # Use %eax  
0x020: halt
```



Load-use dependency

- Value needed by end of decode stage in cycle 7
- Value read from memory in memory stage of cycle 8

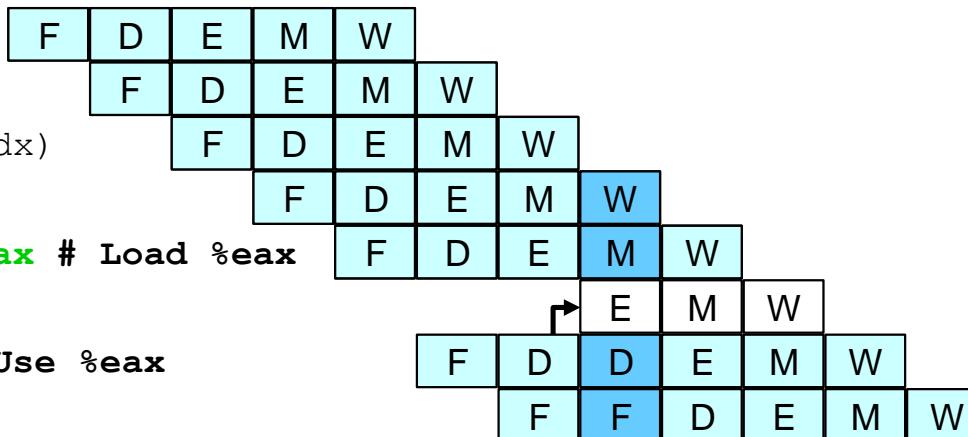


Avoiding Load/Use Hazard

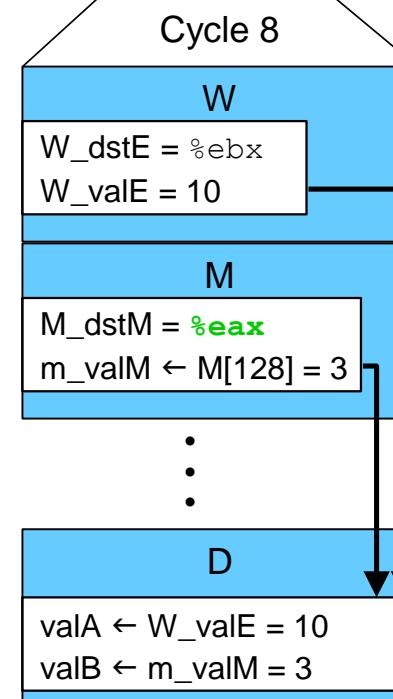
demo-luh.y8

```
0x000: irmovl $128,%edx  
0x006: irmovl $3,%ecx  
0x00c: rmmovl %ecx, 0(%edx)  
0x012: irmovl $10,%ebx  
0x018: mrmovl 0(%edx),%eax # Load %eax  
bubble  
0x01e: addl %ebx,%eax # Use %eax  
0x020: halt
```

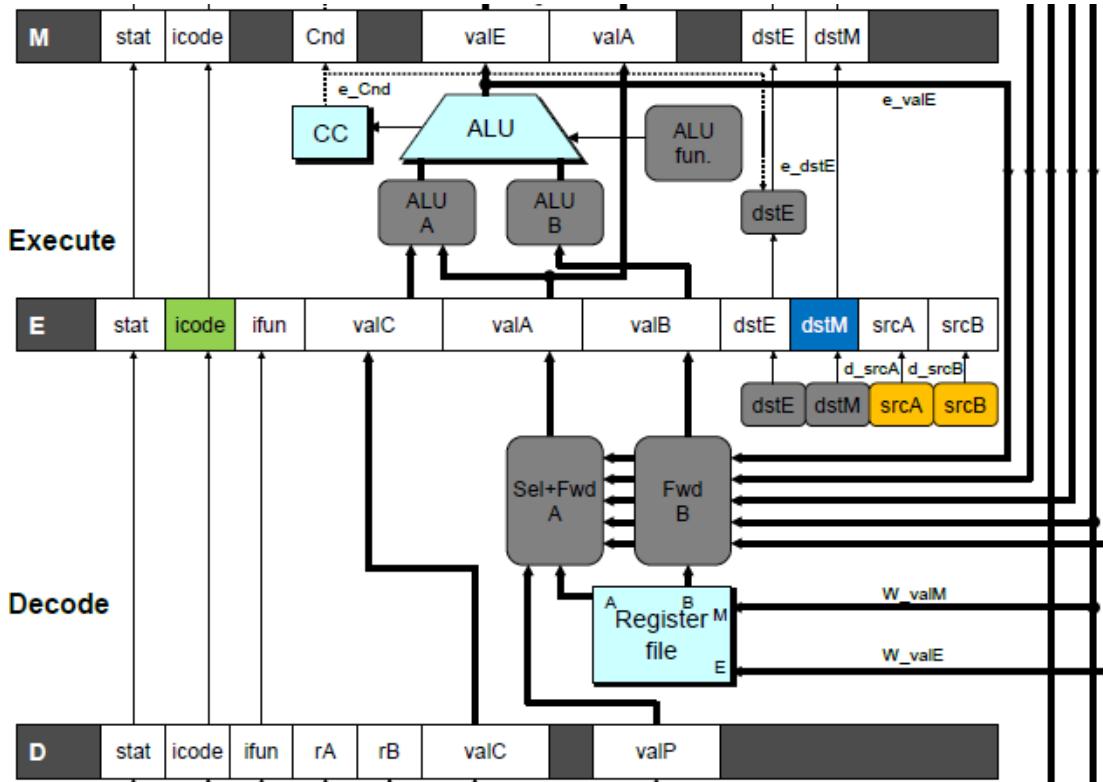
1 2 3 4 5 6 7 8 9 10 11 12



- Stall using instruction for one cycle
- Can then pick up loaded value by forwarding from memory stage

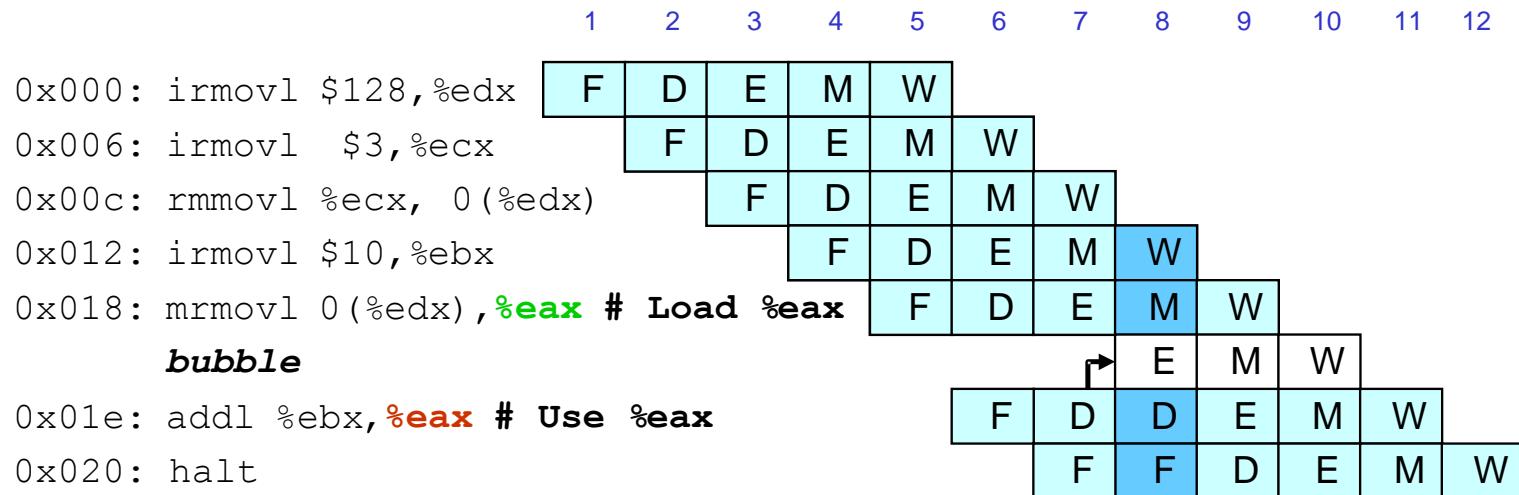


Detecting Load/Use Hazard



Condition	Trigger
Load/Use Hazard	E_icode in { MRMOVL, POPL } && E_dstM in { d_srcA, d_srcB }

Control for Load/Use Hazard



- Stall instructions in fetch and decode stages
- Inject bubble into execute stage

Condition	F	D	E	M	W
Load/Use Hazard	stall	stall	bubble	normal	normal

Dealing with Dependencies between Instructions

Control Hazards

Branch Misprediction Example

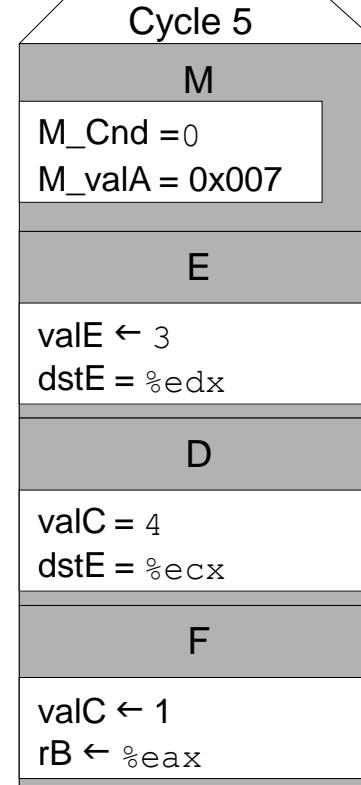
```
0x000: xorl %eax,%eax
0x002: jne t          # Not taken
0x007: irmovl $1, %eax # Fall through
0x00d: nop
0x00e: nop
0x00f: nop
0x010: halt
0x011: t: irmovl $3, %edx # Target (Should not execute)
0x017: irmovl $4, %ecx # Should not execute
0x01d: irmovl $5, %edx # Should not execute
```

- Should only execute first 7 instructions

Branch Misprediction Trace

	1	2	3	4	5	6	7	8	9	
0x000:	xorl %eax, %eax	F	D	E	M	W				
0x002:	jne t # Not taken		F	D	E	M	W			
0x011:	t: irmovl \$3, %edx # Target			F	D	E	M	W		
0x017:	irmovl \$4, %ecx # Target+1				F	D	E	M	W	
0x007:	irmovl \$1, %eax # Fall Through					F	D	E	M	W

- Incorrectly execute two instructions at branch target



Handling Misprediction

```
# demo-j.ys
```

	1	2	3	4	5	6	7	8	9	10
0x000:	xorl %eax,%eax	F	D	E	M	W				
0x002:	jne target # Not taken	F	D	E	M	W				
0x011:	t: irmovl \$2,%edx # Target		F	D						
		bubble			E	M	W			
0x017:	irmovl \$3,%ebx # Target+1		F							
		bubble			D	E	M	W		
0x007:	irmovl \$1,%eax # Fall through			F	D	E	M	W		
0x00d:	nop				F	D	E	M	W	

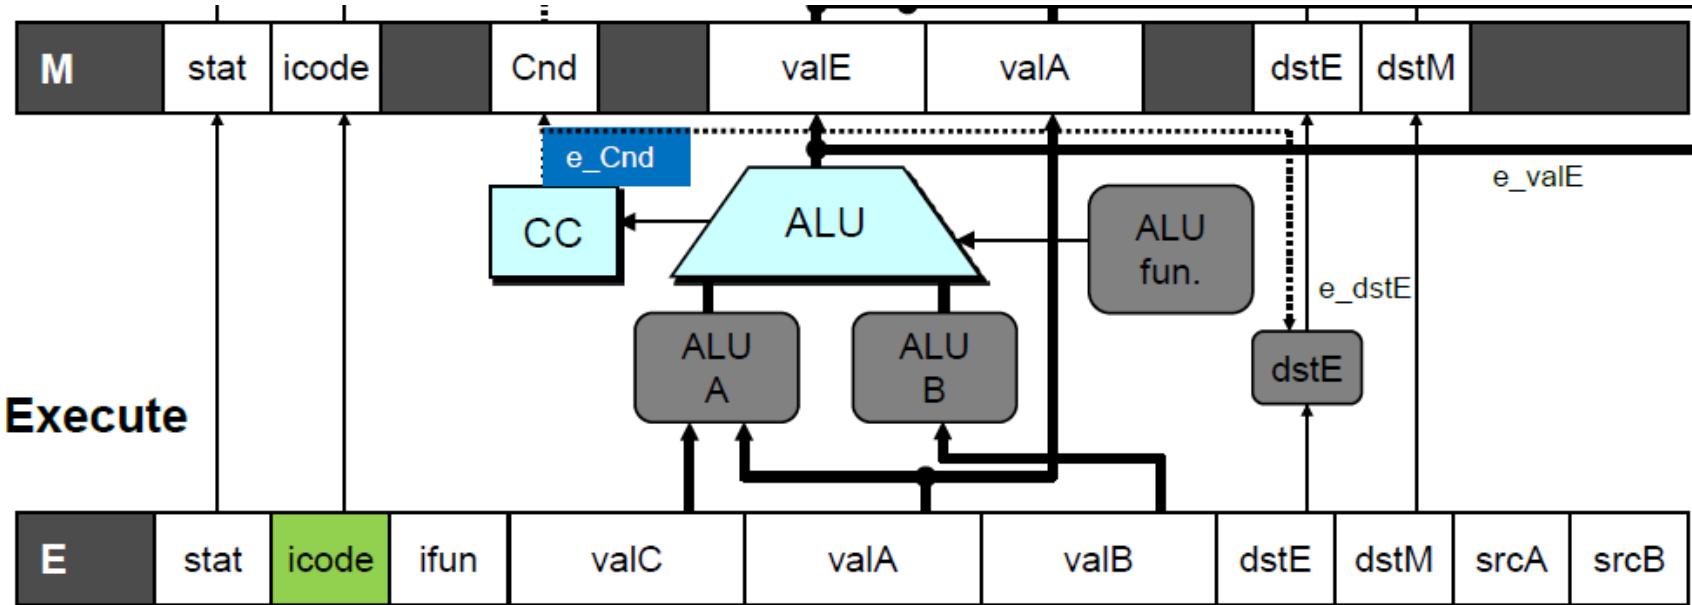
Predict branch as taken

- Fetch 2 instructions at target

Cancel when mispredicted

- Detect branch not-taken in execute stage
- On following cycle, replace instructions in execute and decode by bubbles
- No side effects have occurred yet

Detecting Mispredicted Branch

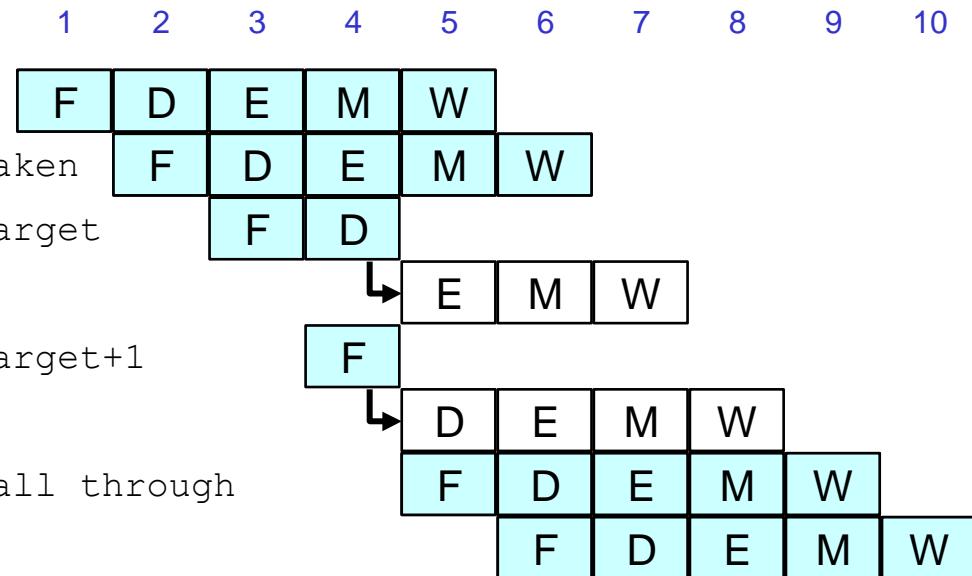


Condition	Trigger
Mispredicted Branch	$E_icode == JXX \& !e_Cnd$

Control for Misprediction

```
# demo-j.ys
```

```
0x000: xorl %eax,%eax  
0x002: jne target # Not taken  
0x011: t: irmovl $2,%edx # Target  
0x017: irmovl $3,%ebx # Target+1  
0x007: irmovl $1,%eax # Fall through  
0x00d: nop
```



Condition	F	D	E	M	W
Mispredicted Branch	normal	bubble	bubble	normal	normal

Return Example

demo-retb.ys

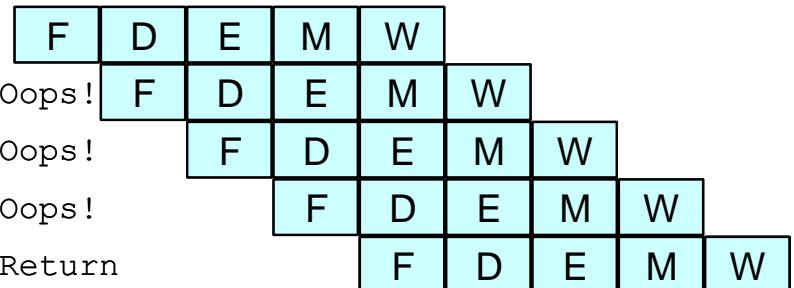
```
0x000:    irmovl Stack,%esp    # Initialize stack pointer
0x006:    call p              # Procedure call
0x00b:    irmovl $5,%esi      # Return point
0x011:    halt
0x020: .pos 0x20
0x020: p: irmovl $-1,%edi    # procedure
0x026:    ret
0x027:    irmovl $1,%eax      # Should not be executed
0x02d:    irmovl $2,%ecx      # Should not be executed
0x033:    irmovl $3,%edx      # Should not be executed
0x039:    irmovl $4,%ebx      # Should not be executed
0x100: .pos 0x100
0x100: Stack:                 # Stack: Stack pointer
```

- Previously executed three additional instructions

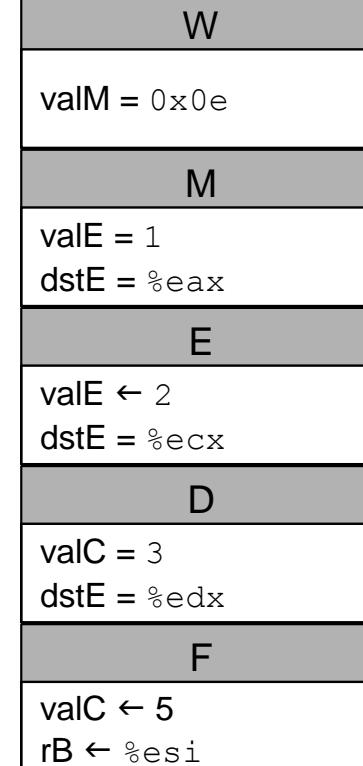
Incorrect Return Example

demo-ret

```
0x023:    ret
0x024:    irmovl $1,%eax # Oops!
0x02a:    irmovl $2,%ecx # Oops!
0x030:    irmovl $3,%edx # Oops!
0x00e:    irmovl $5,%esi # Return
```

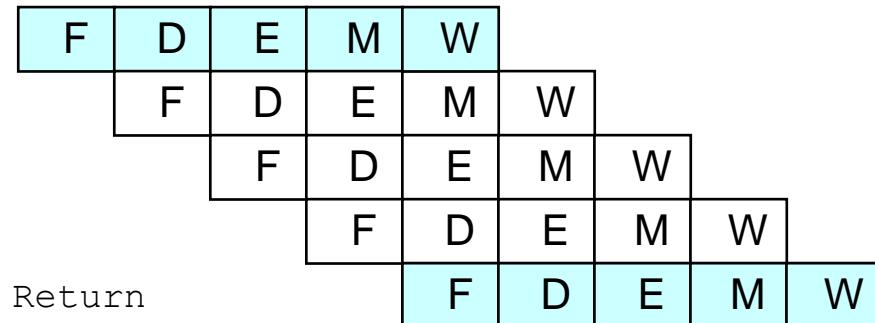


- Incorrectly execute 3 instructions following `ret`

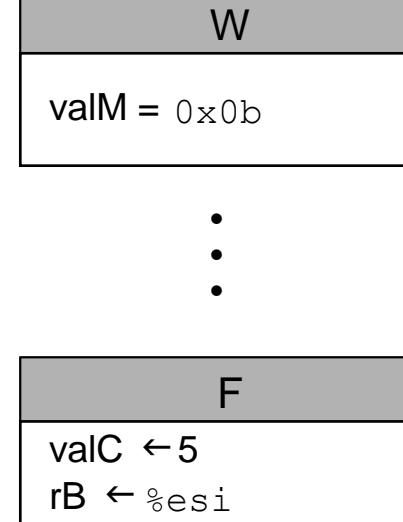


Correct Return Example

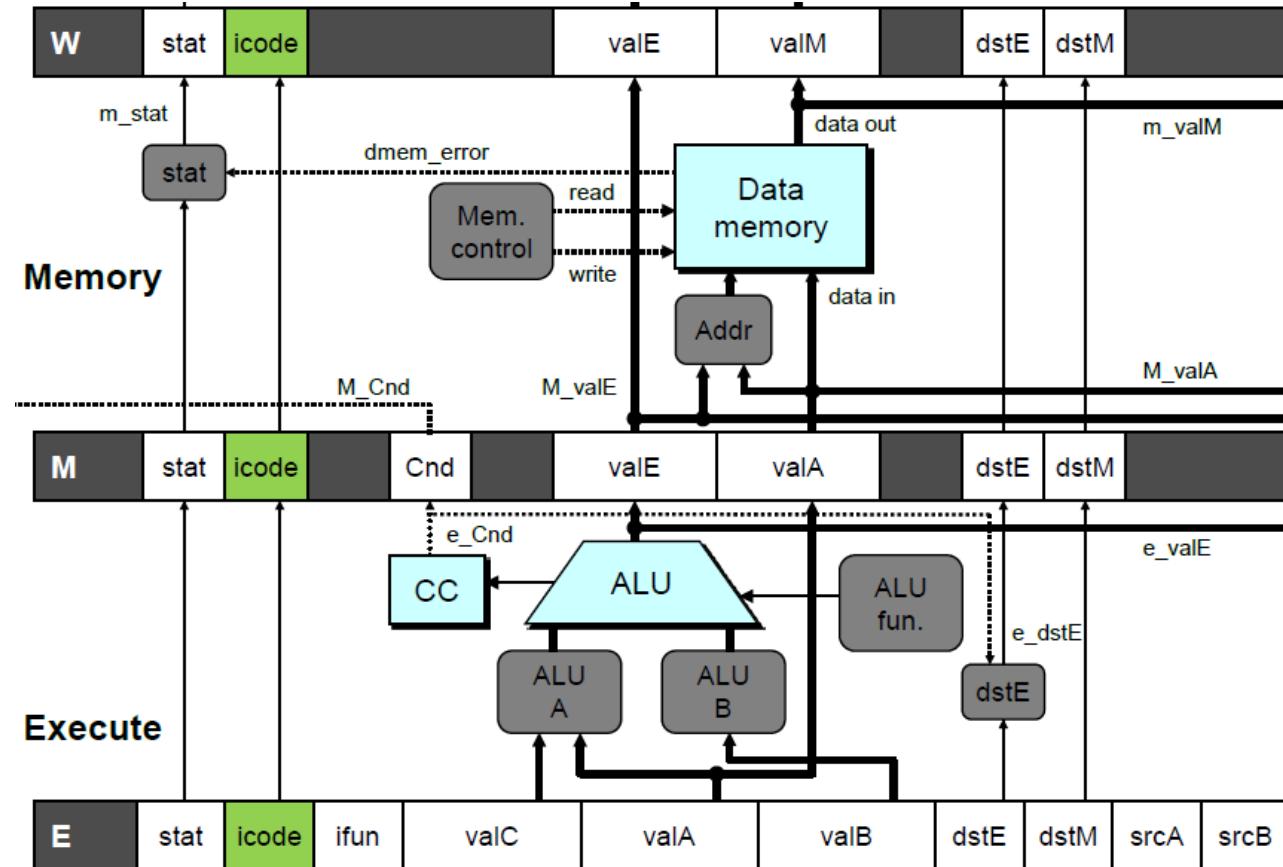
0x026: ret
bubble
bubble
bubble
0x00b: irmovl \$5,%esi # Return



- As `ret` passes through pipeline, stall at fetch stage
 - While in decode, execute, and memory stage
- Inject bubble into decode stage
- Release stall when reach write-back stage



Detecting Return

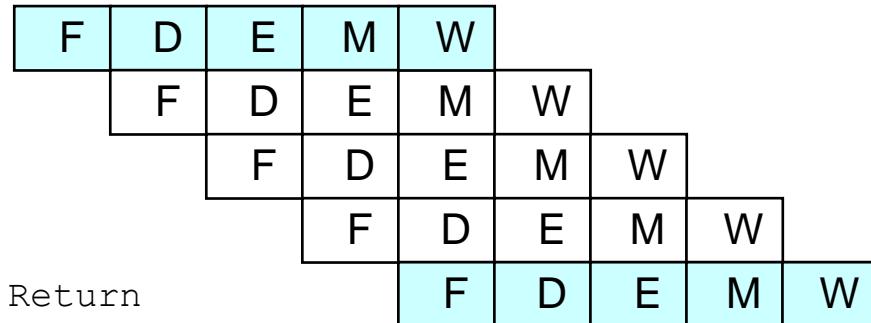


Condition	Trigger
Processing ret	IRET in { D_icode, E_icode, M_icode }

Control for Return

demo-retb

0x026: ret



Condition	F	D	E	M	W
Processing ret	stall	bubble	normal	normal	normal

Special Control Cases

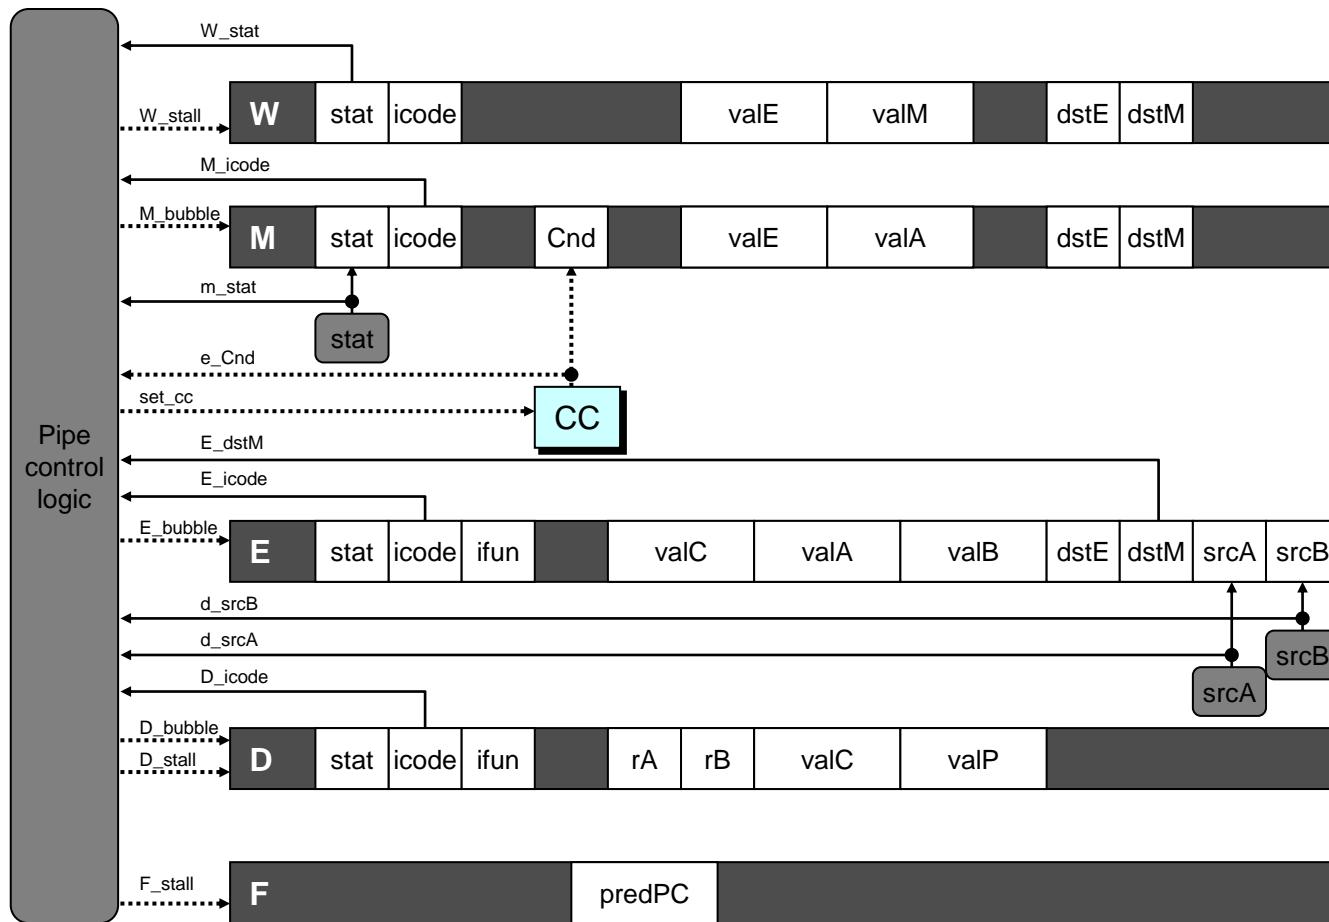
Detection

Condition	Trigger
Processing ret	IRET in { D_icode, E_icode, M_icode }
Load/Use Hazard	E_icode in { IMRMOVL, IPOPL } && E_dstM in { d_srcA, d_srcB }
Mispredicted Branch	E_icode = IJXX & !e_Cnd

Action (on next cycle)

Condition	F	D	E	M	W
Processing ret	stall	bubble	normal	normal	normal
Load/Use Hazard	stall	stall	bubble	normal	normal
Mispredicted Branch	normal	bubble	bubble	normal	normal

Implementing Pipeline Control



- Combinational logic generates pipeline control signals
- Action occurs at start of following cycle

Pipeline Control Logic

- A sequence of control instructions complicates the control logic
 - in particular, should stall in Decode stage (instead of bubble, as an initial inspection suggests)
- Load/use hazard should get priority
- `ret` instruction should be held in decode stage for additional cycle

Condition	F	D	E	M	W
Processing <code>ret</code>	stall	bubble	normal	normal	normal
Load/Use Hazard	stall	stall	bubble	normal	normal
Combination	<i>stall</i>	<u><i>stall</i></u>	<i>bubble</i>	<i>normal</i>	<i>normal</i>

Pipeline Summary

Concept

- Break instruction execution into 5 stages
- Run instructions through in pipelined mode

Limitations

- Can't handle dependencies between instructions when instructions follow too closely
- Data dependencies
 - One instruction writes register, later one reads it
- Control dependency
 - Instruction sets PC in way that pipeline did not predict correctly
 - Mispredicted branch and return

Pipeline Summary

Data Hazards

- Read-after-write dependencies handled by forwarding
 - No performance penalty
- Load/use hazard requires one cycle stall

Control Hazards

- Cancel instructions when detect mispredicted branch
 - Two clock cycles wasted
- Stall fetch stage while `ret` passes through pipeline
 - Three clock cycles wasted

Control Combinations

- Must analyze carefully
- First version had subtle bug
 - Only arises with unusual instruction combination