

# Machine-Level Programming I – Introduction

CSCI 2400: Computer Architecture

**Instructor:**

David Ferry

*Slides adapted from Bryant & O'Hallaron's slides  
via Jason Fritts*

# Turning a corner...

## Course theme:

- Low level (machine level) operation of a processor

## What we've seen so far:

- Bit-level representation of data
  - int, unsigned, char, float, double
  - Strings (arrays)
- Bit-level operations on data
  - Arithmetic (+, -, \*, /, %)
  - Bitwise (&, |, ^, ~, <<, >>)
- We've done data at a low level, what about programs?

# State of the Course

## Past:

- Machine organization
- Data representation
- Data operations
- Intro to C and Linux

## Future:

- Program representation (assembly language)
- Program execution
- Processor architecture / organization
- Memory and cache architecture / organization

# Machine Programming I – Basics

## ■ Instruction Set Architecture

- Software Architecture vs. Hardware Architecture
- Common Architecture Classifications

## ■ The Intel x86 ISA – History and Microarchitectures

## ■ Dive into C, Assembly, and Machine code

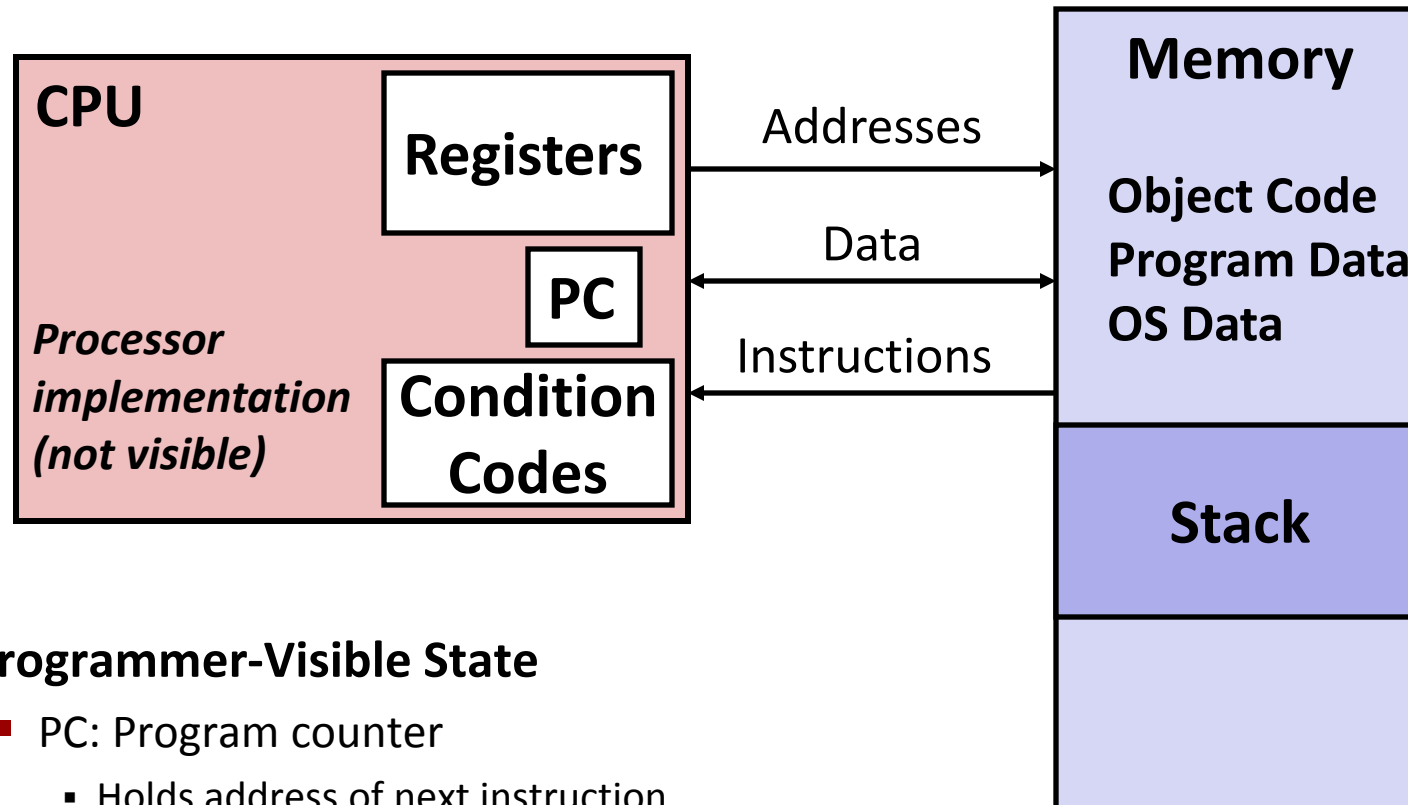
## ■ The Intel x86 Assembly Basics:

- Registers and Operands
- `mov` instruction

## ■ Intro to x86-64

- AMD was first!

# Assembly Programmer's View



## ■ Programmer-Visible State

- **PC: Program counter**
  - Holds address of next instruction
- **Register file**
  - Temp storage for program data
- **Condition codes**
  - Store status info about recent operation
  - Used for conditional branching

## ■ Memory

- Byte addressable array
- Code, user data, (some) OS data
- Includes stack used to support procedures

# Hardware vs. Software Architecture

- There are two parts to the computer architecture of a processor:
  - *Software architecture*
    - known as the **Architecture** or **Instruction Set Architecture (ISA)**
  - *Hardware architecture*
    - known as the **Microarchitecture**
    - a hardware architecture implements an ISA
- The **(software) architecture** includes all aspects of the design that are visible to programmers
- The **microarchitecture** refers to one specific implementation of a software architecture
  - e.g. number of cores, processor frequency, cache sizes, etc.
  - the set of all independent hardware architectures for a given software architecture is known as the *processor family*
    - e.g. the Intel x86 family

# Separation of hardware and software

- The reason for the separation of the (software) **architecture** from the **microarchitecture** (hardware) is backwards compatibility
- **Backwards compatibility ensures:**
  - software written on older processors will run on newer processors (of the same ISA)
  - processor families can always utilize the latest technology by creating new hardware architectures (for the same ISA)
- However, new **microarchitectures** often add to the (software) **architecture**, so software written on newer processors may not run on older processors

# Example Parts of the ISA

## ■ Register file

- Fast, on-processor data storage, very limited
- On Hopper:
  - 14 general purpose registers (rax, rbc, rcx, rdx, rsi, rdi, and r8-r15)
  - Two stack registers (rbp, rsp)
  - Instruction pointer (rip)
  - Flags register (eflags)

## ■ Instruction set

- The set of available instructions
- `movl` – moves 32-bit data (“*movl %edx, %eax*” moves %edx to %eax)
- `addl` – adds two 32-bit operands (“*addl %eax, %ebx*” adds %eax to %ebx)
- `call` – call a function

## ■ These instructions map directly to binary machine code!



# Parts of the Software Architecture

- **There are 4 parts to the (software) architecture**
  - **processor instruction set**
    - the set of available instructions and the rules for using them
  - **register file organization**
    - the number, size, and rules for using registers
  - **memory organization & addressing**
    - the organization of the memory and the rules for accessing data
  - **operating modes**
    - the various modes of execution for the processor
    - there are usually at least two modes:
      - *user* mode (for general use)
      - *system* mode (allows access to privileged instructions and memory)

# Software Architecture: Instruction Set

## ■ The *Instruction Set* defines

- the set of available instructions
- fundamental nature of the instructions
  - simple and fast (how many cycles?)
  - complex and concise
- instruction formats
  - define the rules for using the instructions
- the width (in bits) of the datapath
  - this defines the fundamental size of data in the CPU, including:
    - the size (number of bits) for the data buses in the CPU
    - the number of bits per register in the register file
    - the width of the processing units
    - the number of address bits for accessing memory

# Software Architecture: Instruction Set

## ■ There are 9 fundamental categories of instructions

- arithmetic
  - these instructions perform integer arithmetic, such as add, subtract, multiply, and negate
    - Note: integer division is commonly done in software
- logical
  - these instructions perform Boolean logic (AND, OR, NOT, etc.)
- relational
  - these instructions perform comparisons, including  
==, !=, <, >, <=, >=
  - some ISAs perform comparisons in the conditional branches
- control
  - these instructions enable changes in control flow, both for decision making and modularity
  - the set of control instructions includes:
    - conditional branches
    - unconditional jumps
    - procedure calls and returns

# Software Architecture: Instruction Set

- memory
  - these instructions allow data to be read from or written to memory
- floating-point
  - these instruction perform real-number operations, including add, subtract, multiply, division, comparisons, and conversions
- shifts
  - these instructions allow bits to be shifted or rotated left or right
- bit manipulation
  - these instructions allow data bits to be set or cleared
  - some ISAs do not provide these, since they can be done via logic instructions
- system instructions
  - specialized instructions for system control purposes, such as
    - STOP or HALT (stop execution)
    - cache hints
    - interrupt handling
  - some of these instructions are privileged, requiring system mode

# Software Architecture: Register File

- **The *Register File* is a small, fast temporary storage area in the processor's CPU**
  - it serves as the primary place for holding data values currently being operated upon by the CPU
- **The organization of the register file determines**
  - the number of registers
    - a large number of registers is desirable, but having too many will negatively impact processor speed
  - the number of bits per register
    - this is equivalent to the width of the datapath
  - the purpose of each register
    - ideally, most registers should be general-purpose
    - however, some registers serve specific purposes

# Purpose of Register File

- **Registers are much faster to access than memory**
  - Time to access a local register: ~1 CPU cycle
  - Time to access memory (RAM): hundreds to thousands of CPU cycles
  
- **Operating on memory data requires loads and stores**
  - More instructions to be executed
  
- **Compilers store values in registers whenever possible**
  - Only spill to memory for less frequently used variables
  - Register optimization is important!

# Software Architecture: Memory

## ■ The *Memory Organization & Addressing* defines

### ■ how memory is organized in the architecture

- where data and program memory are unified or separate
- the amount of addressable memory
  - usually determined by the datapath width
- the number of bytes per address
  - most processors are *byte-addressable*, so each byte has a unique addr
- whether it employs virtual memory, or just physical memory
  - virtual memory is usually required in complex computer systems, like desktops, laptops, servers, tablets, smart phones, etc.
  - simpler systems use embedded processors with only physical memory

### ■ rules identifying how instructions access data in memory

- what instructions may access memory (usually only loads, stores)
- what addressing modes are supported
- the ordering and alignment rules for multi-byte primitive data types

# Software Architecture: Operating Modes

- ***Operating Modes*** define the processor's modes of execution
- **The ISA typically supports at least two operating modes**
  - user mode
    - this is the mode of execution for typical use
  - system mode
    - allows access to privileged instructions and memory
    - aside from interrupt and exception handling, system mode is typically only available to system programmers and administrators
    - used to implement operating system privileges
- Processors also generally have hardware testing modes, but these are usually part of the microarchitecture, not the (software) architecture



# Machine Programming I – Basics

## ■ Instruction Set Architecture

- Software Architecture vs. Hardware Architecture
- Common Architecture Classifications

## ■ The Intel x86 ISA – History and Microarchitectures

## ■ Dive into C, Assembly, and Machine code

## ■ The Intel x86 Assembly Basics:

- Registers
- Operands
- `mov` instruction

## ■ Intro to x86-64

- AMD was first!

# Common Architecture (ISA) Classifications:

## Concise vs. Fast: CISC vs. RISC

- *CISC – Complex Instruction Set Computers*
  - complex instructions targeting efficient program representation
  - variable-length instructions
  - versatile addressing modes
  - specialized instructions and registers implement complex tasks
  - NOT optimized for speed – tend to be SLOW
- *RISC – Reduced Instruction Set Computers*
  - small set of simple instructions targeting high speed implementation
  - fixed-length instructions
  - simple addressing modes
  - many general-purpose registers
  - leads to FAST hardware implementations
  - but less memory efficient

# Is x86 CISC? How does it get speed?

- Hard to match RISC performance, but Intel has done just that!

....In terms of speed; less so for power

- CISC instruction set makes implementation difficult

- Hardware translates instructions to simpler *micro-operations*
  - simple instructions: 1-to-1
  - complex instructions: 1-to-many
- Micro-engine similar to RISC
- Market share makes this economically viable

- Comparable performance to RISC

- Compilers avoid CISC instructions

# Classifications: Unified vs. Separate Memory

## ■ von Neumann vs. Harvard architecture

- relates to whether program and data in unified or separate memory
- *von Neumann* architecture
  - program and data are stored in the same unified memory space
  - requires only one physical memory
  - allows self-modifying code
  - however, code and data must share the same memory bus
  - used by most general-purpose processors (e.g. Intel x86)
- *Harvard* architecture
  - program and data are stored in separate memory spaces
  - requires separate physical memory
  - code and data do not share same bus, giving higher bandwidths
  - often used by digital signal processors for data-intensive applications

# Classifications: Performance vs. Specificity

## ■ Microprocessor vs. Microcontroller

### ■ *Microprocessor*

- processors designed for high-performance and flexibility in personal computers and other general purpose applications
- architectures target high performance through a combination of high speed and parallelism
- processor chip contains only CPU(s) and cache
- no peripherals included on-chip

### ■ *Microcontroller*

- processors designed for specific purposes in embedded systems
- only need performance sufficient to needs of that application
- processor chip generally includes:
  - a simple CPU
  - modest amounts of RAM and (Flash) ROM
  - appropriate peripherals needed for specific application
- also often need to meet low power and/or real-time requirements

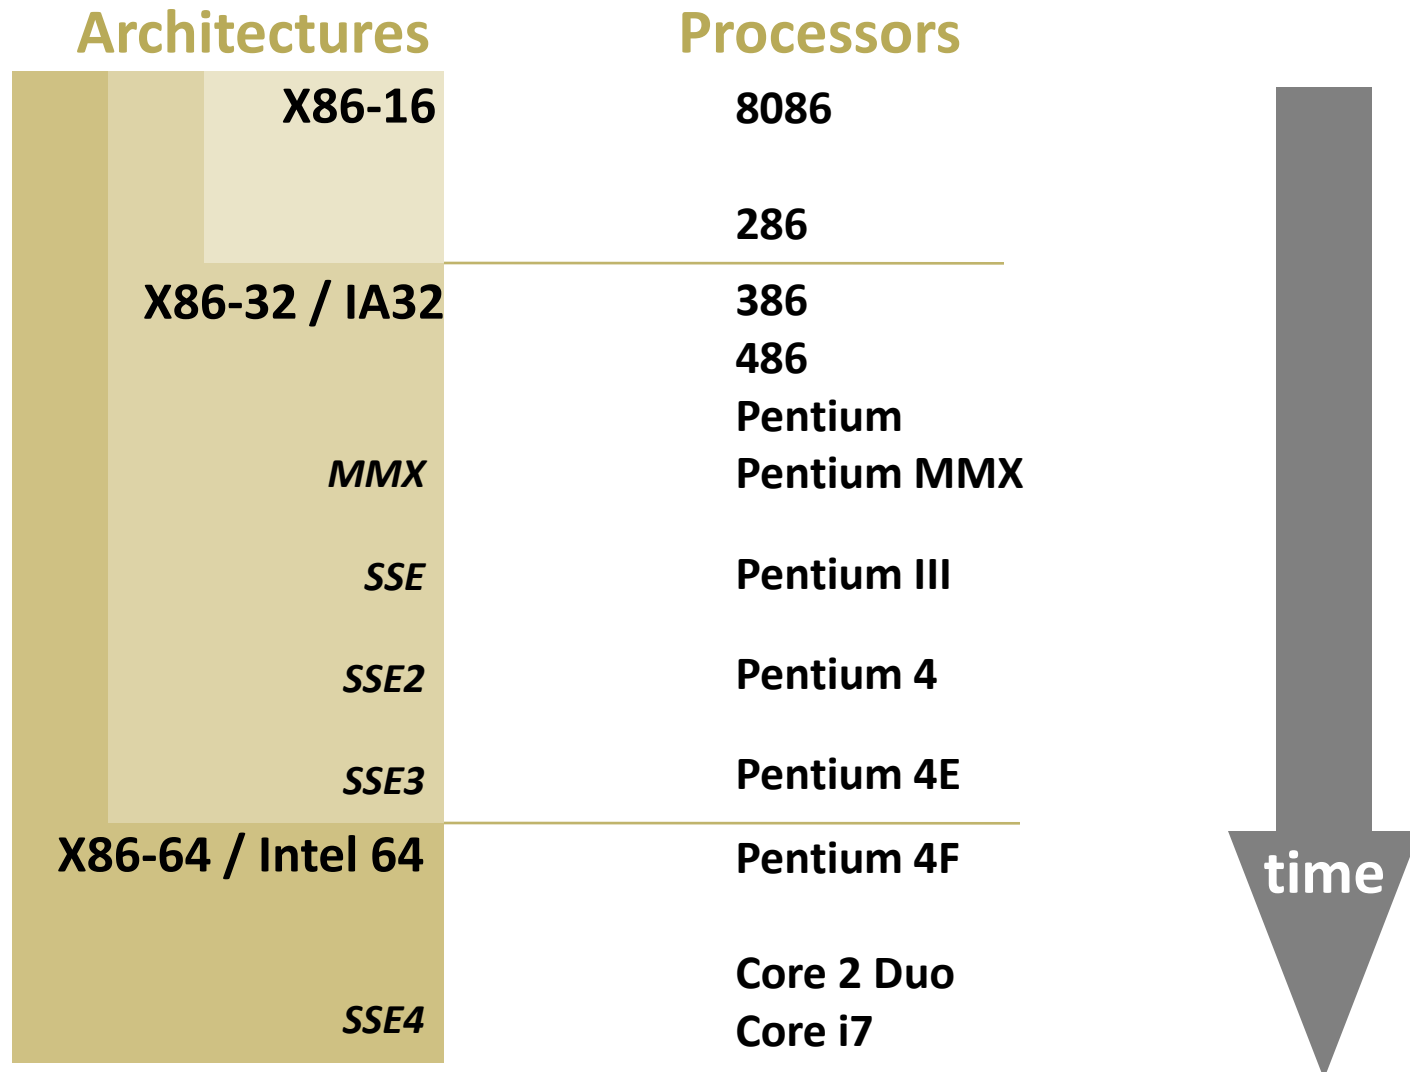
# Machine Programming I – Basics

- **Instruction Set Architecture**
  - Software Architecture vs. Hardware Architecture
  - Common Architecture Classifications
- **The Intel x86 ISA – History and Microarchitectures**
- **Dive into C, Assembly, and Machine code**
- **The Intel x86 Assembly Basics:**
  - Registers
  - Operands
  - `mov` instruction
- **Intro to x86-64**
  - AMD was first!

# Intel x86 Processors

- **The main software architecture for Intel is the x86 ISA**
  - also known as IA-32
  - for 64-bit processors, it is known as x86-64
- **Totally dominate laptop/desktop/server market**
- **Evolutionary design**
  - Backwards compatible back to 8086, introduced in 1978
  - Added more features as time goes on
- **Complex instruction set computer (CISC)**
  - Many different instructions with many different formats
    - *but, only small subset used in Linux programs*

# Intel x86 Family: Many Microarchitectures




IA: often redefined as latest Intel architecture



# Software architecture can grow

- **Backward compatibility does not mean instruction set is fixed**
  - new instructions and functionality can be added to the software architecture over time
- **Intel added additional features over time**
  - Instructions to support multimedia operations (*MMX, SSE*)
    - SIMD parallelism – same operation done across multiple data
  - Instructions enabling more efficient conditional operations



**x86 instruction set**


# Intel x86: Milestones & Trends

<u><i>Name</i></u>	<u><i>Date</i></u>	<u><i>Transistors</i></u>	<u><i>MHz</i></u>
■ <b>8086</b>	<b>1978</b>	<b>29K</b>	<b>5-10</b>
<ul style="list-style-type: none"> <li>■ First 16-bit processor. Basis for IBM PC &amp; DOS</li> <li>■ 1MB address space</li> </ul>			
■ <b>386</b>	<b>1985</b>	<b>275K</b>	<b>16-33</b>
<ul style="list-style-type: none"> <li>■ First 32 bit processor, referred to as IA32</li> <li>■ Added “flat addressing”</li> </ul>			
■ <b>Pentium</b>	<b>1993</b>	<b>3.1M</b>	<b>50-75</b>
■ <b>Pentium II</b>	<b>1996</b>	<b>7.5M</b>	<b>233-300</b>
■ <b>Pentium III</b>	<b>1999</b>	<b>9.5-21M</b>	<b>450-800</b>
■ <b>Pentium 4F</b>	<b>2004</b>	<b>169M</b>	<b>3200-3800</b>
<ul style="list-style-type: none"> <li>■ First 64-bit processor</li> <li>■ Got very hot (up to 115 watts!)</li> </ul>			
■ <b>Core i7</b>	<b>2008</b>	<b>731M</b>	<b>2667-3333</b>

# Intel's 64-Bit History

- **2001: Intel Attempts Radical Shift from IA32 to IA64**
  - Totally different architecture (Itanium)
  - Executes IA32 code only as legacy
  - Performance disappointing
- **2003: AMD Steps in with Evolutionary Solution**
  - x86-64 (now called “AMD64”)
- **Intel Felt Obligated to Focus on IA64**
  - Hard to admit mistake or that AMD is better
- **2004: Intel Announces EM64T extension to IA32**
  - Extended Memory 64-bit Technology
  - Almost identical to x86-64!
- **All but low-end x86 processors support x86-64**
  - But, lots of code still runs in 32-bit mode

# Processor Trends



- Number of transistors has continued to double every 2 years
- In 2004 – we hit the *Power Wall*
  - Processor clock speeds started to leveled off
- Recently – multi-cores have hit the *Memory Wall*

# 2015 State of the Art

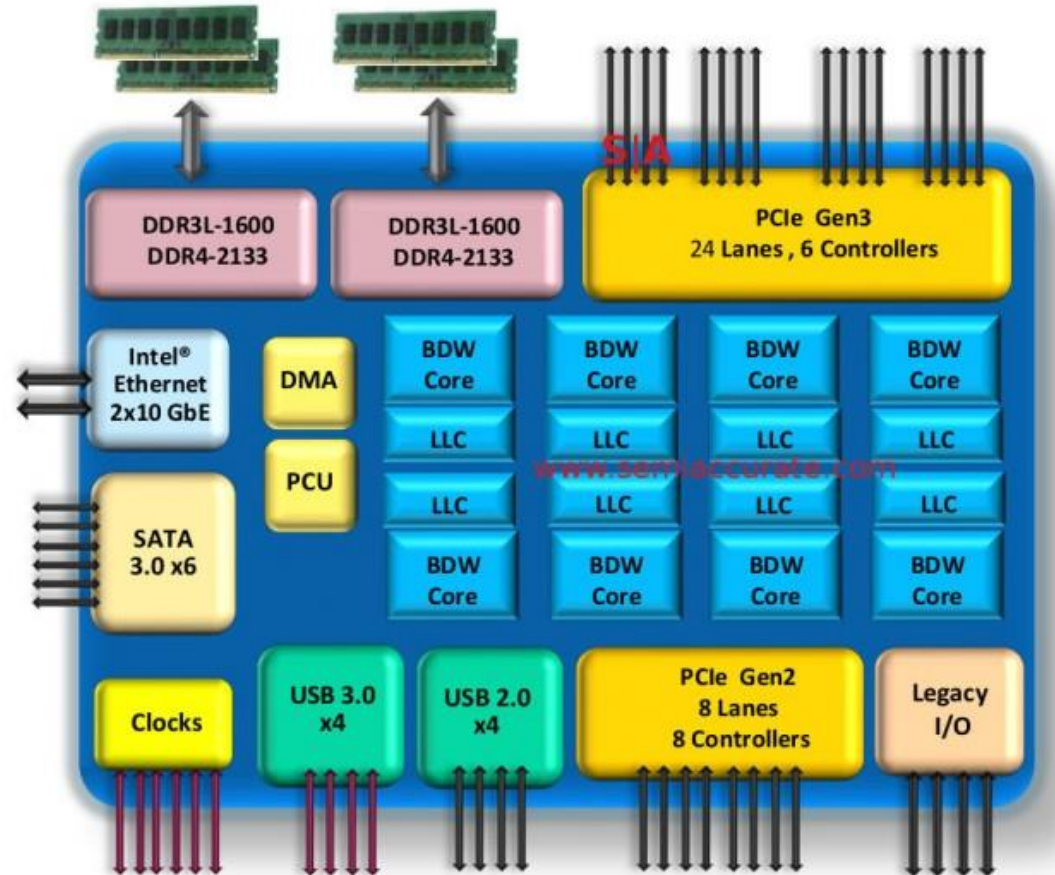
- Core i7 Broadwell 2015

## ■ Desktop Model

- 4 cores
- Integrated graphics
- 3.3-3.8 GHz
- 65W

## ■ Server Model

- 8 cores
- Integrated I/O
- 2-2.6 GHz
- 45W



# Machine Programming I – Basics

## ■ Instruction Set Architecture

Software Architecture (“Architecture” or “ISA”)

vs.

Hardware Architecture (“Microarchitecture”)

## ■ The Intel x86 ISA – History and Microarchitectures

## ■ Dive into C, Assembly, and Machine code

## ■ The Intel x86 Assembly Basics:


- Registers
- Operands
- `mov` instruction

## ■ Intro to x86-64

- AMD was first!

# Turning C into Object Code

- Code in separate *translation units*: `p1.c` `p2.c`
- Compile with command: `gcc -O1 -m32 p1.c p2.c -o p`
  - Use basic optimizations (`-O1`)
  - Put resulting binary in file `p`
  - On 64-bit machines, specify 32-bit x86 code (`-m32`)



# Compiling Into Assembly

## C Code

```
int sum(int x, int y)
{
    int t = x+y;
    return t;
}
```

Some compilers use  
instruction "leave"



## Generated IA32 Assembly

```
sum:
    pushl %ebp
    movl %esp,%ebp
    movl 12(%ebp),%eax
    addl 8(%ebp),%eax
    popl %ebp
    ret
```

Obtain with command:

```
gcc -O1 -S -m32 code.c
```

**-S** specifies compile to assembly (vs object) code, and  
produces file **code.s**



# Assembly Characteristics: Simple Types

- **Integer data of 1, 2, or 4 bytes**
  - Data values
  - Addresses (`void*` pointers)
- **Floating point data of 4, 8, or 10 bytes**
- **No concept of aggregate types such as arrays or structures**
  - Just contiguously allocated bytes in memory

# Assembly Characteristics: Operations

## ■ Perform some operation on register or memory data

- arithmetic
- logical
- bit shift or manipulation
- comparison (relational)

## ■ Transfer data between memory and register

- Load data from memory into register
- Store register data into memory

## ■ Transfer control

- Unconditional jumps to/from procedures
- Conditional branches

# Object Code

## Code for `sum`

`0x401040 <sum>:`

`0x55`

`0x89`

`0xe5`

`0x8b`

`0x45`

`0x0c`

`0x03`

`0x45`

`0x08`

`0x5d`

`0xc3`

- **Total of 11 bytes**
- **Each instruction 1, 2, or 3 bytes**
- **Starts at address 0x401040**

## ■ Assembler

- Translates `.s` into `.o`
- Binary encoding of each instruction
- Nearly-complete image of executable code
- Missing linkages between code in different files

## ■ Linker

- Resolves references between files
- Combines with static run-time libraries
  - E.g., code for `malloc`, `printf`
- Some libraries are *dynamically linked*
  - Linking occurs when program begins execution

# Machine Instruction Example

```
int t = x+y;
```

```
addl 8(%ebp), %eax
```

Similar to expression:

```
x += y
```

More precisely:

```
int eax;
```

```
int *ebp;
```

```
eax += ebp[2]
```

```
0x80483ca: 03 45 08
```

## ■ C Code

- Add two signed integers

## ■ Assembly

- Add 2 4-byte integers
  - “Long” words in GCC parlance
  - Same instruction whether signed or unsigned
- Operands:
 

<b>x:</b>	Register	<b>%eax</b>
<b>y:</b>	Memory	<b>M[%ebp+8]</b>
<b>t:</b>	Register	<b>%eax</b>

  - Return function value in **%eax**

## ■ Object Code

- 3-byte instruction
- Stored at address **0x80483ca**

# Disassembling Object Code

## Disassembled

```
080483c4 <sum>:  
80483c4: 55          push    %ebp  
80483c5: 89 e5       mov     %esp, %ebp  
80483c7: 8b 45 0c    mov     0xc(%ebp), %eax  
80483ca: 03 45 08    add     0x8(%ebp), %eax  
80483cd: 5d          pop     %ebp  
80483ce: c3          ret
```

## ■ Disassembler

`objdump -d p`

- Useful tool for examining object code
- Analyzes bit pattern of series of instructions
- Produces approximate rendition of assembly code
- Can be run on either a `.out` (complete executable) or `.o` file

# Alternate Disassembly

## Object

0x401040:

0x55

0x89

0xe5

0x8b

0x45

0x0c

0x03

0x45

0x08

0x5d

0xc3

## Disassembled

Dump of assembler code for function sum:

```
0x080483c4 <sum+0>:      push    %ebp
0x080483c5 <sum+1>:      mov     %esp, %ebp
0x080483c7 <sum+3>:      mov     0xc(%ebp), %eax
0x080483ca <sum+6>:      add     0x8(%ebp), %eax
0x080483cd <sum+9>:      pop     %ebp
0x080483ce <sum+10>:     ret
```

### ■ Within gdb Debugger

`gdb p`

`disassemble sum`

- Disassemble procedure

`x/11xb sum`

- Examine the 11 bytes starting at `sum`

# What Can be Disassembled?

```
% objdump -d WINWORD.EXE

WINWORD.EXE:      file format pei-i386

No symbols in "WINWORD.EXE".
Disassembly of section .text:

30001000 <.text>:
30001000:  55                      push    %ebp
30001001:  8b ec                  mov     %esp,%ebp
30001003:  6a ff                  push    $0xffffffff
30001005:  68 90 10 00 30 push    $0x30001090
3000100a:  68 91 dc 4c 30 push    $0x304cdc91
```

- Anything that can be interpreted as executable code
- Disassembler examines bytes and reconstructs assembly source

# Machine Programming I – Basics

- **Instruction Set Architecture**
  - Software Architecture vs. Hardware Architecture
  - Common Architecture Classifications
- **The Intel x86 ISA – History and Microarchitectures**
- **Dive into C, Assembly, and Machine code**
- **The Intel x86 Assembly Basics:**
  - Common instructions
  - Registers, Operands, and `mov` instruction
  - Addressing modes
- **Intro to x86-64**
  - AMD was first!



# World-wary aside: Instruction Syntax

## Two prevalent assembler syntaxes:

### ■ AT&T syntax

- Aka GNU Assembler syntax, aka GAS syntax
- Dominant in Unix/Linux world
- Subject of this class
- E.g.: `movl $5, %eax`
- E.g.: `movl 8(%ebp), %eax`

### ■ Intel Syntax

- Aka Microsoft Assembler syntax, aka MASM syntax
- Dominant in Microsoft world
- E.g.: `mov eax, 5`
- E.g.: `mov eax, [ebp + 8]`

# Typical Instructions in Intel x86

## ■ Arithmetic

- `add, sub, neg, imul, div, inc, dec, leal, ...`

## ■ Logical (bit-wise Boolean)

- `and, or, xor, not`

## ■ Relational

- `cmp, test, sete, ...`

## ■ Control

- `je, jle, jg, jb, jmp, call, ret, ...`

## ■ Moves & Memory Access

- `mov, push, pop, movswl, movzbl, cmov, ...`
- *nearly all x86 instructions can access memory*

## ■ Shifts


- `shr, sar, shl, sal` (same as `shl`)

## ■ Floating-point


- `fld, fadd, fsub, fxch, addsd, movss, cvt..., ucom...`
- *float-point change completely with x86-64*

# CISC Instructions: Variable-Length


a. JE EIP + displacement




b. CALL




c. MOV EBX, [EDI + 45]



d. PUSH ESI



e. ADD EAX, #6765



f. TEST EDX, #42



# Machine Programming I – Basics

## ■ Instruction Set Architecture

Software Architecture (“Architecture” or “ISA”)

vs.

Hardware Architecture (“Microarchitecture”)

## ■ The Intel x86 ISA – History and Microarchitectures

## ■ Dive into C, Assembly, and Machine code

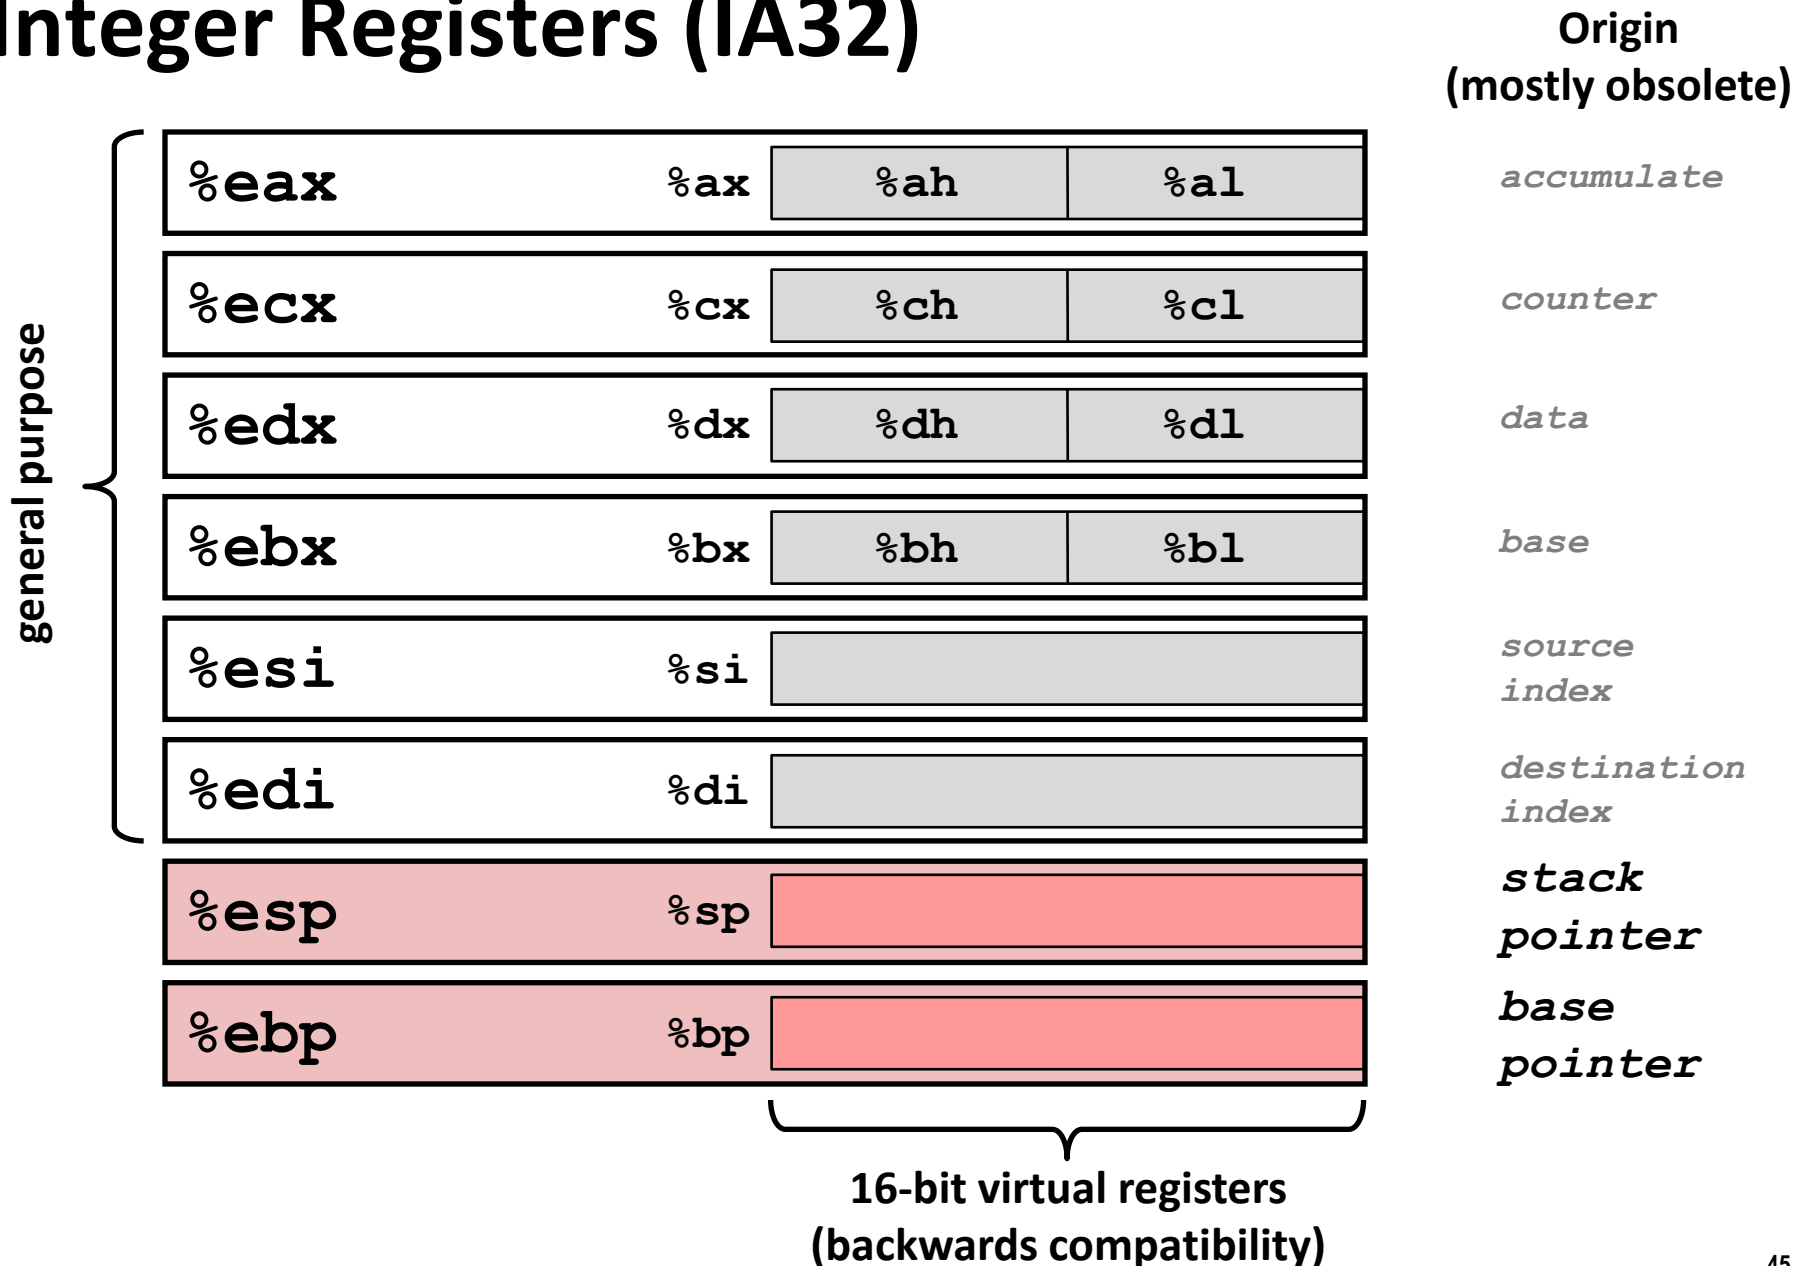
## ■ The Intel x86 Assembly Basics:

- Common instructions
- Registers, Operands, and **mov** instruction
- Addressing modes

## ■ Intro to x86-64

- AMD was first!

# Integer Registers (IA32)



# Moving Data: IA32

## ■ Moving Data

`movl Source, Dest`

## ■ Operand Types

- **Immediate:** Constant integer data
  - example: `$0x400`, `$-533`
  - like C constant, but prefixed with ``$'`
  - encoded with 1, 2, or 4 bytes
- **Register:** One of 8 integer registers
  - example: `%eax`, `%edx`
  - but `%esp` and `%ebp` reserved for special use
  - others have special uses in particular situations
- **Memory:** 4 consecutive bytes of memory at address given by register
  - simplest example: `(%eax)`
  - various other “address modes”

`%eax``%ecx``%edx``%ebx``%esi``%edi``%esp``%ebp`

# movl Operand Combinations

	Source	Dest	Src, Dest	C Analog
movl {	Imm {	Reg	movl \$0x4,%eax	temp = 0x4;
		Mem	movl \$-147, (%eax)	*p = -147;
	Reg {	Reg	movl %eax,%edx	temp2 = temp1;
		Mem	movl %eax, (%edx)	*p = temp;
	Mem	Reg	movl (%eax), %edx	temp = *p;

**Cannot do memory-memory transfer with a single instruction**

# Machine Programming I – Basics

## ■ Instruction Set Architecture

Software Architecture (“Architecture” or “ISA”)

vs.

Hardware Architecture (“Microarchitecture”)

## ■ The Intel x86 ISA – History and Microarchitectures

## ■ Dive into C, Assembly, and Machine code

## ■ The Intel x86 Assembly Basics:

- Common instructions
- Registers, Operands, and `mov` instruction
- Addressing modes

## ■ Intro to x86-64

- AMD was first!



# Simple Memory Addressing Modes

## ■ Normal:

**(R)                      Mem[Reg[R]]**

- Register R specifies memory address

```
movl (%ecx) , %eax
```

## ■ Displacement:

**D(R)                      Mem[Reg[R]+D]**

- Register R specifies start of memory region
- Constant displacement D specifies offset

```
movl 8 (%ebp) , %edx
```

# Using Simple Addressing Modes

```
void swap(int *xp, int *yp)
{
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

swap:

pushl %ebx

} Set Up

movl 8(%esp), %edx  
 movl 12(%esp), %eax  
 movl (%edx), %ecx  
 movl (%eax), %ebx  
 movl %ebx, (%edx)  
 movl %ecx, (%eax)

} Body

popl %ebx  
 ret

} Finish

# Using Simple Addressing Modes

```
void swap(int *xp, int *yp)
{
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

swap:

**pushl %ebx**

} **Set Up**

**movl 8(%esp), %edx**  
**movl 12(%esp), %eax**  
**movl (%edx), %ecx**  
**movl (%eax), %ebx**  
**movl %ebx, (%edx)**  
**movl %ecx, (%eax)**


} **Body**

**popl %ebx**  
**ret**

} **Finish**

# Understanding Swap

```
void swap(int *xp, int *yp)
{
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```



Register	Value
%edx	xp
%ecx	yp
%ebx	t0
%eax	t1

```
movl 8(%esp), %edx    # edx = xp
movl 12(%esp), %eax   # eax = yp
movl (%edx), %ecx     # ecx = *xp (t0)
movl (%eax), %ebx     # ebx = *yp (t1)
movl %ebx, (%edx)     # *xp = t1
movl %ecx, (%eax)     # *yp = t0
```

# Understanding Swap

%eax	
%edx	
%ecx	
%ebx	
%esi	
%edi	
%esp	
%ebp	0x104

		Address
		123
		0x124
		456
		0x120
		0x11c
		0x118
		0x114
yp	12	0x120
xp	8	0x124
	4	Rtn adr
%esp	0	0x104

```

movl 8(%esp), %edx    # edx = xp
movl 12(%esp), %eax   # eax = yp
movl (%edx), %ecx     # ecx = *xp (t0)
movl (%eax), %ebx     # ebx = *yp (t1)
movl %ebx, (%edx)     # *xp = t1
movl %ecx, (%eax)     # *yp = t0

```

# Understanding Swap

%eax	
%edx	0x124
%ecx	
%ebx	
%esi	
%edi	
%esp	
%ebp	0x104

		Address
		123
		0x124
		456
		0x120
		0x11c
		0x118
		0x114
yp	12	0x120
xp	8	0x10c
	4	Rtn adr
%esp	0	0x104

```

movl 8(%esp), %edx    # edx = xp
movl 12(%esp), %eax   # eax = yp
movl (%edx), %ecx     # ecx = *xp (t0)
movl (%eax), %ebx     # ebx = *yp (t1)
movl %ebx, (%edx)     # *xp = t1
movl %ecx, (%eax)     # *yp = t0

```

# Understanding Swap

%eax	0x120
%edx	0x124
%ecx	
%ebx	
%esi	
%edi	
%esp	
%ebp	0x104

		Address
		123
		0x124
		456
		0x120
		0x11c
		0x118
		0x114
yp	12	0x120
xp	8	0x124
	4	Rtn adr
%esp	0	
		0x104

```

movl 8(%esp), %edx    # edx = xp
movl 12(%esp), %eax   # eax = yp
movl (%edx), %ecx     # ecx = *xp (t0)
movl (%eax), %ebx     # ebx = *yp (t1)
movl %ebx, (%edx)     # *xp = t1
movl %ecx, (%eax)     # *yp = t0

```

# Understanding Swap

%eax	0x120
%edx	0x124
%ecx	123
%ebx	
%esi	
%edi	
%esp	
%ebp	0x104

		Address
	Offset	
		0x124
		0x120
		0x11c
		0x118
yp	12	0x114
		0x110
		0x10c
xp	8	0x108
	4	0x104
%esp → 0		

```

movl 8(%esp), %edx    # edx = xp
movl 12(%esp), %eax   # eax = yp
movl (%edx), %ecx    # ecx = *xp (t0)
movl (%eax), %ebx     # ebx = *yp (t1)
movl %ebx, (%edx)     # *xp = t1
movl %ecx, (%eax)     # *yp = t0

```



# Understanding Swap

%eax	0x120
%edx	0x124
%ecx	123
%ebx	456
%esi	
%edi	
%esp	
%ebp	0x104

		Offset	Address
			0x124
			0x120
			0x11c
			0x118
			0x114
yp	12	0x120	0x110
xp	8	0x124	0x10c
	4	Rtn adr	0x108
%esp	→ 0		0x104

```
movl    8(%esp), %edx    # edx = xp
movl    12(%esp), %eax   # eax = yp
movl    (%edx), %ecx     # ecx = *xp (t0)
movl    (%eax), %ebx     # ebx = *yp (t1)
movl    %ebx, (%edx)     # *xp = t1
movl    %ecx, (%eax)     # *yp = t0
```

# Understanding Swap

%eax	0x120
%edx	0x124
%ecx	123
%ebx	456
%esi	
%edi	
%esp	
%ebp	0x104

		Address
		0x124
		0x120
		0x11c
		0x118
		0x114
yp	12	0x110
xp	8	0x10c
	4	0x108
%esp	0	0x104

```

movl 8(%esp), %edx    # edx = xp
movl 12(%esp), %eax   # eax = yp
movl (%edx), %ecx     # ecx = *xp (t0)
movl (%eax), %ebx     # ebx = *yp (t1)
movl %ebx, (%edx)   # *xp = t1
movl %ecx, (%eax)     # *yp = t0

```

# Understanding Swap

%eax	0x120
%edx	0x124
%ecx	123
%ebx	456
%esi	
%edi	
%esp	
%ebp	0x104

		Address
		456
		0x124
		123
		0x120
		0x11c
		0x118
		0x114
yp	12	0x120
xp	8	0x124
	4	Rtn adr
%esp	0	0x104

```

movl 8(%esp), %edx    # edx = xp
movl 12(%esp), %eax   # eax = yp
movl (%edx), %ecx     # ecx = *xp (t0)
movl (%eax), %ebx     # ebx = *yp (t1)
movl %ebx, (%edx)     # *xp = t1
movl %ecx, (%eax)   # *yp = t0

```

# Complete Memory Addressing Modes

## ■ Most General Form

$$D(Rb, Ri, S) \qquad \text{Mem[ Reg[}Rb\text{] + } S * \text{ Reg[}Ri\text{] + } D\text{]}$$

- $D$ : Constant “displacement” 1, 2, or 4 bytes
- $Rb$ : Base register: Any of 8 integer registers
- $Ri$ : Index register: Any, except for `%esp` (likely not `%ebp` either)
- $S$ : Scale: 1, 2, 4, or 8 (*why these numbers?*)

## ■ Special Cases

$$(Rb, Ri) \qquad \text{Mem[ Reg[}Rb\text{] + Reg[}Ri\text{] ]}$$

$$D(Rb, Ri) \qquad \text{Mem[ Reg[}Rb\text{] + Reg[}Ri\text{] + } D\text{]}$$

$$(Rb, Ri, S) \qquad \text{Mem[ Reg[}Rb\text{] + } S * \text{ Reg[}Ri\text{] ]}$$

# Machine Programming I – Basics

- **Instruction Set Architecture**
  - Software Architecture vs. Hardware Architecture
  - Common Architecture Classifications
- **The Intel x86 ISA – History and Microarchitectures**
- **Dive into C, Assembly, and Machine code**
- **The Intel x86 Assembly Basics:**
  - Common instructions
  - Registers, Operands, and `mov` instruction
  - Addressing modes
- **Intro to x86-64**
  - AMD was first!

# AMD created first 64-bit version of x86

## ■ Historically

- AMD has followed just behind Intel
- A little bit slower, a lot cheaper

## ■ 2003, developed 64-bit version of x86: **x86-64**

- Recruited top circuit designers from DEC and other diminishing companies
- Built Opteron: tough competitor to Pentium 4

# Intel's 64-Bit

- **Intel Attempted Radical Shift from IA32 to IA64**
  - Totally different architecture (*Itanium*)
  - Executes IA32 code only as legacy
  - Performance disappointing
- **2003: AMD Stepped in with Evolutionary Solution**
  - Originally called **x86-64** (now called **AMD64**)
- **2004: Intel Announces their 64-bit extension to IA32**
  - Originally called **EMT64** (now called **Intel 64**)
  - Almost identical to **x86-64**!
- **Collectively known as **x86-64****
  - minor differences between the two

# Data Representations: IA32 vs. x86-64

## ■ Sizes of C Objects (in bytes)

<u><i>C Data Type</i></u>	<u><i>Intel IA32</i></u>	<u><i>x86-64</i></u>
▪ unsigned	4	4
▪ int	4	4
▪ long int	4	8
▪ char	1	1
▪ short	2	2
▪ float	4	4
▪ double	8	8
▪ long double	10/12	16
▪ pointer (e.g. char *)	4	8



# x86-64 Integer Registers

<b>%rax</b>	<b>%eax</b>	<b>%r8</b>	<b>%r8d</b>
<b>%rbx</b>	<b>%ebx</b>	<b>%r9</b>	<b>%r9d</b>
<b>%rcx</b>	<b>%ecx</b>	<b>%r10</b>	<b>%r10d</b>
<b>%rdx</b>	<b>%edx</b>	<b>%r11</b>	<b>%r11d</b>
<b>%rsi</b>	<b>%esi</b>	<b>%r12</b>	<b>%r12d</b>
<b>%rdi</b>	<b>%edi</b>	<b>%r13</b>	<b>%r13d</b>
<b>%rsp</b>	<b>%esp</b>	<b>%r14</b>	<b>%r14d</b>
<b>%rbp</b>	<b>%ebp</b>	<b>%r15</b>	<b>%r15d</b>

- Up to 6 function arguments are passed via registers
- Explicitly makes **%ebp/%rbp** general purpose

# New Instructions for 64-bit Operands

- Long word **l** (4 Bytes)  $\leftrightarrow$  Quad word **q** (8 Bytes)
  
- New instructions:
  - `movl`  $\Rightarrow$  `movq`
  - `addl`  $\Rightarrow$  `addq`
  - `sall`  $\Rightarrow$  `salq`
  - etc.
  
- 32-bit instructions that generate 32-bit results
  - Set higher order bits of destination register to 0
  - Example: `addl`

# 32-bit code for `int swap`

`swap:`

```
void swap(int *xp, int *yp)
{
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

`pushl %ebx`

**Set  
Up**

```
movl 8(%esp), %edx
movl 12(%esp), %eax
movl (%edx), %ecx
movl (%eax), %ebx
movl %ebx, (%edx)
movl %ecx, (%eax)
```

**Body**

```
popl %ebx
ret
```

**Finish**

# 64-bit code for `int` swap

swap:

```
void swap(int *xp, int *yp)
{
    int t0 = *xp;
    int t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

```
movl    (%rdi), %edx
movl    (%rsi), %eax
movl    %eax, (%rdi)
movl    %edx, (%rsi)
```

ret

} Set  
Up

} Body

} Finish

- **Operands passed in registers (why useful?)**
  - First input arg (**x**p) in %rdi, second input arg (**y**p) in %rsi
  - 64-bit pointers
- **No stack operations required**
- **32-bit ints held temporarily in %eax and %edx**

# 64-bit code for long int swap

swap\_1:

```
void swap(long *xp, long *yp)
{
    long t0 = *xp;
    long t1 = *yp;
    *xp = t1;
    *yp = t0;
}
```

```
movq    (%rdi), %rdx
movq    (%rsi), %rax
movq    %rax, (%rdi)
movq    %rdx, (%rsi)
```

ret

} Set  
Up

} Body

} Finish

## ■ 64-bit long ints

- Pass input arguments in registers **%rax** and **%rdx**
- **movq** operation
  - “q” stands for quad-word

# Machine Programming I – Basics

- **Instruction Set Architecture**
  - Software Architecture vs. Hardware Architecture
  - Common Architecture Classifications
- **The Intel x86 ISA – History and Microarchitectures**
- **Dive into C, Assembly, and Machine code**
- **The Intel x86 Assembly Basics:**
  - Common instructions
  - Registers, Operands, and `mov` instruction
  - Addressing modes
- **Intro to x86-64**
  - AMD was first!

# Machine Programming I – Summary

## ■ Instruction Set Architecture

- Many different varieties and features of processor architectures
- Separation of (software) Architecture and Microarchitecture is key for backwards compatibility

## ■ The Intel x86 ISA – History and Microarchitectures

- Evolutionary design leads to many quirks and artifacts

## ■ Dive into C, Assembly, and Machine code

- Compiler must transform statements, expressions, procedures into low-level instruction sequences

## ■ The Intel x86 Assembly Basics:

- The x86 move instructions cover wide range of data movement forms

## ■ Intro to x86-64

- A major departure from the style of code seen in IA32