

Course Overview

CSCI 2400/ ECE 3217: Computer Architecture

Instructors:

David Ferry

Slides adapted from Bryant & O'Hallaron's slides

Overview

- **Course theme**
- **Five realities**
- **Logistics**

Course Theme:

Abstraction Is Good But Don't Forget Reality

■ Most CS and CE courses emphasize abstraction

- Abstract data types
- Asymptotic analysis

■ These abstractions have limits

- Especially in the presence of bugs
- Need to understand details of underlying implementations

■ Useful outcomes


- Become more effective programmers
 - Able to find and eliminate bugs efficiently
 - Able to understand and tune for program performance
- Prepare for later “systems” classes in CS & ECE
 - Compilers, Operating Systems, Networks, Computer Architecture, Embedded Systems

Great Reality #1:

Ints are not Integers, Floats are not Reals

■ Example 1: Is $x^2 \geq 0$?

■ Float's: Yes!



■ Int's:

- $40000 * 40000 \Rightarrow 1600000000$
- $50000 * 50000 \Rightarrow ??$

■ Example 2: Is $(x + y) + z = x + (y + z)$?

■ Unsigned & Signed Int's: Yes!

■ Float's:

- $(1e20 + -1e20) + 3.14 \rightarrow 3.14$
- $1e20 + (-1e20 + 3.14) \rightarrow ??$

Consequences of Abstraction:

Code Security

```
/* Kernel memory region holding user-accessible data */
#define KSIZE 1024
char kbuf[KSIZE];

/* Copy at most maxlen bytes from kernel region to user buffer */
int copy_from_kernel(void *user_dest, int maxlen) {
    /* Byte count len is minimum of buffer size and maxlen */
    int len = KSIZE < maxlen ? KSIZE : maxlen;
    memcpy(user_dest, kbuf, len);
    return len;
}
```

- Similar to code found in FreeBSD's implementation of `getpeername`
- There are legions of smart people trying to find vulnerabilities in programs

Typical Usage

```
/* Kernel memory region holding user-accessible data */
#define KSIZE 1024
char kbuf[KSIZE];

/* Copy at most maxlen bytes from kernel region to user buffer */
int copy_from_kernel(void *user_dest, int maxlen) {
    /* Byte count len is minimum of buffer size and maxlen */
    int len = KSIZE < maxlen ? KSIZE : maxlen;
    memcpy(user_dest, kbuf, len);
    return len;
}
```

```
#define MSIZE 528

void getstuff() {
    char mybuf[MSIZE];
    copy_from_kernel(mybuf, MSIZE);
    printf("%s\n", mybuf);
}
```

Malicious Usage

```
/* Kernel memory region holding user-accessible data */
#define KSIZE 1024
char kbuf[KSIZE];

/* Copy at most maxlen bytes from kernel region to user buffer */
int copy_from_kernel(void *user_dest, int maxlen) {
    /* Byte count len is minimum of buffer size and maxlen */
    int len = KSIZE < maxlen ? KSIZE : maxlen;
    memcpy(user_dest, kbuf, len);
    return len;
}
```

```
#define MSIZE 528

void getstuff() {
    char mybuf[MSIZE];
    copy_from_kernel(mybuf, -MSIZE);
    . . .
}
```

- AKA “buffer over-read”
- This is essentially what happened with the Heartbleed bug (2014)

Computer Arithmetic

■ Does not generate random values

- Arithmetic operations have important mathematical properties

■ Cannot assume all “usual” mathematical properties

- Due to finiteness of representations
- Integer operations satisfy “ring” properties
 - Commutativity, associativity, distributivity
- Floating point operations satisfy “ordering” properties
 - Monotonicity, values of signs

■ Observation

- Need to understand which abstractions apply in which contexts
- Important issues for compiler writers and serious application programmers

Great Reality #2:

You've Got to Know Assembly

- **Chances are, you'll never write programs in assembly**
 - Compilers are much better & more patient than you are
- **But: Assembly is key to understanding machine-level execution**
 - Behavior of programs in presence of bugs
 - High-level language models break down
 - Tuning program performance
 - Understand optimizations done / not done by the compiler
 - Understanding sources of program inefficiency
 - Implementing system software
 - Compiler has machine code as target
 - Operating systems must manage process state
 - Creating / fighting malware
 - x86 assembly is the language of choice!

Assembly Code Example

■ Time Stamp Counter

- Special 64-bit register in Intel-compatible machines
- Incremented every clock cycle
- Read with rdtsc instruction

■ Application

- Measure time (in clock cycles) required by procedure

```
double t;  
start_counter();  
P();  
t = get_counter();  
printf("P required %f clock cycles\n", t);
```

Code to Read Counter

- Write small amount of assembly code using GCC's asm facility
- Inserts assembly code into machine code generated by compiler

```
static unsigned cyc_hi = 0;
static unsigned cyc_lo = 0;

/* Set *hi and *lo to the high and low order bits
   of the cycle counter.
*/
void access_counter(unsigned *hi, unsigned *lo)
{
    asm("rdtsc; movl %%edx,%0; movl %%eax,%1"
        : "=r" (*hi), "=r" (*lo)
        :
        : "%edx", "%eax");
}
```

Great Reality #3: Memory Matters

Random Access Memory Is an Unphysical Abstraction

■ Memory is not unbounded

- It must be allocated and managed
- Many applications are memory dominated
- Memory access is relatively slow

■ Memory referencing bugs especially pernicious

- Effects are distant in both time and space

■ Memory performance is not uniform

- Cache and virtual memory effects can greatly affect program performance
- Adapting program to characteristics of memory system can lead to major speed improvements

Memory Referencing Bug Example

```
double fun(int i)
{
    volatile double d[1] = {3.14};
    volatile long int a[2];
    a[i] = 1073741824; /* Possibly out of bounds */
    return d[0];
}
```

```
fun(0)    =>    3.14
fun(1)    =>    3.14
fun(2)    =>    3.1399998664856
fun(3)    =>    2.00000061035156
fun(4)    =>    3.14, then segmentation fault
```

■ **Result is architecture specific**

Memory Referencing Bug Example

```
double fun(int i)
{
    volatile double d[1] = {3.14};
    volatile long int a[2];
    a[i] = 1073741824; /* Possibly out of bounds */
    return d[0];
}
```

fun(0) => 3.14
 fun(1) => 3.14
 fun(2) => 3.13999998664856
 fun(3) => 2.000000061035156
 fun(4) => 3.14, then segmentation fault

Explanation:

Saved State	4	} Location accessed by fun(i)
d7 ... d4	3	
d3 ... d0	2	
a[1]	1	
a[0]	0	

Memory Referencing Errors

■ C and C++ do not provide any memory protection

- Out of bounds array references
- Invalid pointer values
- Abuses of malloc/free

■ Can lead to nasty bugs

- Whether or not bug has any effect depends on system and compiler
- Action at a distance
 - Corrupted object logically unrelated to one being accessed
 - Effect of bug may be first observed long after it is generated


■ How can I deal with this?

- Program in Java, Python, etc.
- Understand what possible interactions may occur
- Use or develop tools to detect referencing errors (e.g. Valgrind)

Memory System Performance Example

```
void copyij(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (i = 0; i < 2048; i++)
        for (j = 0; j < 2048; j++)
            dst[i][j] = src[i][j];
}
```

```
void copyji(int src[2048][2048],
            int dst[2048][2048])
{
    int i,j;
    for (j = 0; j < 2048; j++)
        for (i = 0; i < 2048; i++)
            dst[i][j] = src[i][j];
}
```



21 times slower
(Pentium 4)

■ Hierarchical memory organization


- Caches trade off speed and size

■ Performance depends on access patterns

- Including how step through multi-dimensional array

The Memory Mountain

Intel Core i7
2.67 GHz
32 KB L1 d-cache
256 KB L2 cache
8 MB L3 cache




Great Reality #4: There's more to performance than asymptotic complexity

- **Constant factors matter too!**
- **And even exact op count does not predict performance**
 - Easily see 10:1 performance range depending on how code written
 - Must optimize at multiple levels: algorithm, data representations, procedures, and loops
- **Must understand system to optimize performance**
 - How programs compiled and executed
 - How to measure program performance and identify bottlenecks
 - How to improve performance without destroying code modularity and generality

Example Matrix Multiplication

Matrix-Matrix Multiplication (MMM) on 2 x Core 2 Duo 3 GHz (double precision)
Gflop/s




- Standard desktop computer, vendor compiler, using optimization flags
- Both implementations have **exactly** the same operations count ($2n^3$)
- What is going on?

MMM Plot: Analysis

Matrix-Matrix Multiplication (MMM) on 2 x Core 2 Duo 3 GHz

Gflop/s



- Reason for 20x: Blocking or tiling, loop unrolling, array scalarization, instruction scheduling, search to find best choice
- **Effect: fewer register spills, L1/L2 cache misses, and TLB misses**

Great Reality #5:

Computers do more than execute programs

■ They need to get data in and out

- I/O system critical to program reliability and performance
- Hard drives, SSDs, USB, CD/DVD, keyboards, mice, etc.
- What's important? Response time, latency, or throughput?
- Large writes vs. Small (4Kb) writes

■ They communicate with each other over networks

- Many system-level issues arise in presence of network
 - Concurrent operations by autonomous processes
 - Coping with unreliable media
 - Cross platform compatibility
 - Complex performance issues

Course Perspective

■ Most Systems Courses are Builder-Centric

- Computer Architecture
 - Design pipelined processor in Verilog
- Operating Systems
 - Implement large portions of operating system
- Compilers
 - Write compiler for simple language
- Networking
 - Implement and simulate network protocols

Course Perspective (Cont.)

■ This Course is more Programmer-Centric

- Purpose is to show how by knowing more about the underlying system, one can be more effective as a programmer
- Enable you to
 - Understand how C/C++ executes in hardware
 - Write programs that are more reliable and efficient
 - Incorporate features that require hooks into OS
 - E.g., concurrency, signal handlers
- Not just a course for dedicated hackers
 - We bring out the hidden hacker in everyone
 - Imagine a zoologist that didn't know biology or chemistry
- Cover material in this course that you won't see elsewhere

Course Website

■ Class Website: <http://cs.slu.edu/~dferry/courses/csci2400>

- Detailed class information and policies
- Full Schedule, including:
 - lecture topics and code examples
 - assignments
 - exam dates
- All assignments posted on website
- Most lecture slides posted on website
- Class website is the official syllabus and takes precedence over anything that is said in slides or in class

■ SLU Blackboard

- Blackboard is not used for this course

Textbook

■ Randal E. Bryant and David R. O'Hallaron,

- “Computer Systems: A Programmer’s Perspective”, Third Edition (*CS:APP3e*), Prentice Hall, 2016
- Textbook’s website: <http://csapp.cs.cmu.edu>
- Recommend getting a hardcopy, since exams are often open book & notes
 - Laptops, tablets, etc. not allowed during exams

■ C reference textbook

- “C Programming”
- a free online reference text for C programming, that may prove beneficial for those who haven’t used C (or C++) before

Grading

■ Exams (50%)

- mid-semester exams: 15% each
- final 20%

■ Assignments (45%): approximately 7-9 assignments

■ Class Participation (5%)

- for participation in hands-on work during class

■ Curving policy on website

■ Late Policy:

- 10% penalty for up to 24 hours late
- 20% penalty for up to 48 hours late
- Penalty waived or accepted after 48 hours only at instructor's discretion

Attendance and Class Guidelines

- **Attendance is at students' discretion, but highly recommended**
- **Questions and Participation highly encouraged**
 - If you have a question or need clarification, it's very likely that other students will likewise benefit from your question
- **Laptops / computers may be used during class**
 - But NOT during exams

Policy for Collaborating on Assignments

- **Collaboration allowed, even encouraged, PROVIDED that:**
 - you only discuss the problem, not the solution
 - students may help guide each other in the process of solving the problem, BUT each student MUST turn in their own answer
 - students MUST indicate who they collaborated with on their cover sheet

Cheating

■ What is cheating?

- *Sharing code or answers:* copying, retyping, looking at, or supplying a file
- *Detailed coaching:* helping your friend to write code or an answer, line by line
- *Copying code* from previous course or from elsewhere on WWW
 - only allowed to use code supplied in class or on course website
 - must cite such code

■ What is NOT cheating?

- Explaining how to use systems or tools
- Helping others understand design issues or the process for solving a problem

■ Penalty for cheating:

- Ranges, based on severity, from zero on assignment to being sent before Academic Honesty Committee
- Records saved for all incidents of cheating

■ Detection of cheating:

- Instructor is (unfortunately) extremely experienced at detecting cheating

Topic: Programs and Data

■ Topics

- Assembly (and machine) language vs. High-level languages (HLLs)
- Instruction set architecture
 - CPU (core)
 - register file
 - processing units (ALU, FPU, etc.)
- Types of instructions
 - arithmetic
 - logical
 - shifts and bit manipulation
 - memory
 - compares
 - branches and jumps
 - procedure calls & returns
- Representation of variables, arrays and data structures

Topic: Computer Architecture

■ Topics

- Fundamentals of Logic Design (gates & circuits)
- Processor Organization
- CPU (core) Organization
- Fetch-Decode-Execute Cycle and Datapath Flow
- Sequential (single-cycle) Datapath
- Pipelined Datapath
 - purpose / benefit
 - data and control dependencies
 - hazards
 - bypassing / forwarding
 - branch prediction

Topic: Memory and the Memory Hierarchy

■ Topics

- Data representation
- Memory technology (disk vs. RAM vs. ROM vs. cache)
- Loads & Stores (reads & writes)
- Physical vs. Virtual memory
 - page tables, address translation, and TLB
 - how memory organized within a process
 - global vs. heap vs. stack memory
- Cache memory
 - purpose / benefit
 - locality
 - how it works

Topic: Performance and Optimization

■ Topics

- How simple modifications in assembly / machine code can dramatically affect execution time
- Co-optimization (control and data)
- Measuring time on a computer
- Related to architecture, compilers, and OS

*Welcome
and Enjoy!*